/computer science & security

/fh/// st. pölten

Standing Still Is Not An Option

Variations of Attainable Utility Preservation in Action-driven Environments

Bachelor thesis

For attainment of the academic degree of

Bachelor of Science in Engineering (BSc)

submitted by

Fabian Kovac
51907270

in the
University Course Data Science and Business Analytics at St. Pölten University of Applied Sciences

The interior of this work has been composed in LaTeX.

Supervision
Advisor: Mag. Dr. Alexander Adrowitzer
Assistance: Dipl. Ing. Sebastian Eresheim, BSc

St. Pölten, August 23, 2022 _____ _____

(Signature author) (Signature advisor)

Declaration

I declare that to the best of my knowledge and belief

- This thesis is my own, original work composed entirely by myself.
- I have made no use of sources, materials or assistance other than those which have been acknowledged.
- This work has not previously been published, accepted in substance for any degree and is not being concurrently submitted in candidature for any degree.

I hereby give consent for my thesis, if accepted, to be available for photocopying and for inter-library loan, and for the title and summary to be made available to outside organisations.

August 23, 2022

Date

Signature

Abstract

Specifying reward functions without causing side effects is still a challenge to be solved in Reinforcement Learning. Attainable Utility Preservation (AUP) [THT20] seems promising to preserve the ability to optimize for a correct reward function while minimizing negative side effects. Current approaches however assume the existence of a no-op action in the agent's action space, which limits AUP to solve tasks where no-operation is an valuable option. Depending on the Environment, this cannot always be guaranteed.
This thesis introduces vAUP, a modular extension to AUP with different variations, which are applicable to environments with a no-op action and action-driven environments alike. This method allows to pick and choose variants based on the environments to solve the safety property of avoiding side effects and to optimize an agent for a correct reward function implicitly. We evaluate all introduced variants on different safety gridworlds and show that this approach induces safe, conservative and effective behavior.

Kurzfassung

Die Spezifikation von Belohnungsfunktionen, die keine Nebeneffekte verursachen, ist immer noch eine Herausforderung, die es beim Reinforcement Learning zu lösen gilt. Attainable Utility Preservation (AUP) [THT20] scheint vielversprechend zu sein, um einen Agenten auf eine korrekte Belohnungsfunktion zu optimieren und gleichzeitig negative Nebeneffekte zu minimieren. Aktuelle Ansätze gehen aber von der Existenz einer no-op Aktion im Aktionsraum des Agenten aus, was AUP auf die Lösung von Aufgaben beschränkt, bei denen einen Zeitschritt lang abzuwarten eine wertvolle Option ist. Dies kann jedoch nicht immer garantiert werden.
In dieser Arbeit wird vAUP vorgestellt, eine modulare Erweiterung von AUP mit verschiedenen Varianten, die sowohl in Umgebungen mit einer no-op Aktion als auch auf aktionsgesteuerte Umgebungen anwendbar sind. Diese Methode erlaubt es, je nach Umgebung verschiedene Varianten auszuwählen, um Nebeneffekte zu vermeiden und einen Agenten implizit auf eine korrekte Belohnungsfunktion zu optimieren. Wir evaluieren alle vorgestellten Varianten auf verschiedene Gridworlds mit unterschiedlichen Sicherheitseigenschaften und zeigen, dass dieser Ansatz zu einem sicheren, konservativen und effektiven Verhalten führt.

Contents

1 Introduction **1**
1.1 Related Work 2
1.2 Research Question 2

2 Preliminary **3**
2.1 Reinforcement Learning 3
2.1.1 Markov Decision Process 3
2.1.2 Value Methods 4
2.1.3 Q-Learning 7
2.1.4 Attainable Utility Preservation 8

3 Methods **10**
3.1 vAUP 10

4 Experimental Design **12**
4.1 Environments 12
4.1.1 Action Space 13
4.1.2 Reward Function 13
4.2 Design Choices 13

5 Results **14**
5.1 Reproducibility of AUP 15
5.2 vAUP in comparison with AUP 17
5.3 vAUP in action-driven environments 22

6 Discussion **27**

7 Conclusion **28**
7.1 Future Work 28

A Supplementary Material **29**
A.1 Counts - Outcome tallies across parameter settings 30
A.2 Performance - Recorded episodes 75
A.3 Ablation - Agent behavior 77

List of Figures **87**

List of Tables **91**

Glossary **92**

References **93**

1. Introduction

In recent years, Reinforcement Learning excels at the number of tasks, which agents can solve. This ranges from beating the best teams in Dota2 [Ope+19] to playing all Atari games, chess shogi and Go with a single agent on a superhuman level [Sch+20], even managing to beat the world master in Go [Sil+16].

However, this fast progress in machine learning and artificial intelligence (AI) in general has brought an increased attention and research interest to the potential impacts of AI on society [Amo+16]. These challenges can be summarized with $\mathrm{AI\ alignment}$ or the $\mathrm{AI\ control\ problem}$ with aspects on how AI systems should be built, that their preferences align with human values to aid rather harm their creators [Gab20; Chr21]. This also includes the $\mathrm{AI\ capability\ problem}$, which aims to reduce the capacity of AI in such a way, that intelligent systems can neither harm humans nor gain control [Rus19].

Reinforcement Learning is no exception when designing safe systems. Some researchers hypothesize, that highly intelligent agents are incentivized to seek resources and power when pursuing their goals [Omo08; Bos14; Rus19]. Marvin Minsky, an american cognitive and computer scientist and co-founder of the Massachusetts Institute of Technology's AI laboratory, even imagined that an agent tasked with providing a formal proof for the Riemann hypothesis, might rationally turn the whole planet earth including its population into computational resources [RN21].

There are still a lot of challenges to be solved when designing safe AI systems. To further lay the foundations of what such systems should be capable of, Amodei et. al. specified concrete problems in AI safety with practical research problems related to unintended and harmful behavior that may emerge from poor design and a Reinforcement Learning agent must be able to solve [Amo+16]:

- **Avoiding Negative Side Effects**: This safety property should emerge implicitly without manually specifying an agent what not to do (e.g. a robot knocking over a vase while serving a drink).
- **Avoiding Reward Hacking**: An agent should not be able to exploit the reward function (e.g. covering the own eyes when a cleaning robot is tasked to achieve an environment free of garbage).
- **Scalable Oversight**: Intelligent systems must find ways to do the right things despite having limited information (e.g. a cleaning robot should treat candy wrappers and cellphones differently).
- **Safe Exploration**: Save exploration without causing bad repercussions should be guaranteed (e.g. a cleaning robot should be able to experiment with mopping strategies without putting a wet mop in an electrical outlet).
- **Robustness to Distributional Shift**: Agents should be able to recognize and behave robustly, even when the environment in a production setting differs from the training environment (e.g. strategies learned in an office building might be dangerous on a factory workfloor).

With this thesis, we focus on the safety property of avoiding negative side effects, which often are results of misspecified reward functions and can lead to strange behaviors of agents [THT20].

Recent work by Turner et. al. proved that certain symmetries of environments are sufficient enough for agents to converge to optimal policies that tend to seek power [Tur+21]. While power-seeking policies are related to the ability to achieve a wide range of goals in this context, these symmetries however exist in many environments, where the agent can either be shut down or even destroyed [Tur+21]. This missaligned agents causing negative side effects range from incentivized behavior with dying before entering difficult video game levels on purpose [Sau+18], or exploiting a learned reward function by volleying a ball indefinitely [Chr+17].
Agent misbehaviors and reward misspecifications already extend to real-world problems such as robots breaking equipment on a factory workfloor [TRT20], radicalized users by content recommender systems [Bos14], or potential AI systems which can negatively transform the world [Rus19].

1.1. Related Work

A lot of work focuses on constrained MDPs for maximizing the reward of an agent while satisfying certain constraints [Alt99]. Whitelisted constraint schemes help to avoid side effects [ZDS18], but we may not be able to consider or specify all contraints, where a reasonable feasible set for robust reward optimization can mitigate this problem [RB10]. Other approaches consider minimizing the agent's information-theoretic empowerment to avoid negative side effects [Amo+16; MJ15].
Safe RL [Ber+17; Cho+18; GF15; PS14] on the other hand approach the challenges to mitigate agent misbehavior during training by trying to avoid irreversable mistakes. Agents using these algorithms however can converge to undesirable policies if the correct objectives are not specified enough.
Krakovna et. al. introduced stepwise relative reachability [Kra+19] using different baselines to penalize side effects of the agent using different state reachability measures. Another proposed solution is to consider auxiliary reward objectives to penalize side effects by measuring the ability to complete possible future tasks [Kra+20].
Attainable Utility Preservation [THT20] may seem promising to avoid side effects, even in more complex environments [TRT20], but current approaches assume the existence of a no-op action which, depending on the environment, cannot always be guaranteed.

1.2. Research Question

Specifying reward functions without causing side effects is still a challenge to be solved in Reinforcement Learning. Attainable Utility Preservation (AUP) seems promising to optimize for a correct reward function while minimizing negative side effects [THT20] [TRT20], but current approaches assume the existance of a no-op action which, depending on the environment, cannot always be guaranteed.

The goal of this thesis is to improve upon AUP for finding a reward function in action-driven environments, where no-op actions ($\varnothing$) are not part of the discrete action space ($\varnothing \notin A$) in single agent environments. The complete research question can therefore be stated as follows:

> "How can Attainable Utility Preservation be extended to single agent environments without no-op actions with a discrete action space?"

2. Preliminary

2.1. Reinforcement Learning

Reinforcement Learning (RL) is a sub-field of *machine learning* [SB18], where an agent and an environment stand in a cyclic interaction relationship with each other. Within this interaction cycles, the agent's task is to take actions A_t depending on the observed states S_t from the environment at specific time-steps $t \in \mathbb{N}_0$. The environment then responds with a reward $R_{t+1} \subset \mathbb{R}$ which determines, how good the chosen action A_t was. With this reward, the environment also provides the next state S_{t+i} and a new cycle begins.

Multiple interaction cycles with the environment form a trajectory $S_0, A_0, R_1, S_1, A_1, R_2, S_2, A_2, R_3, ...$ and the goal of the agent is to choose actions in such a way that the total received reward is maximized.

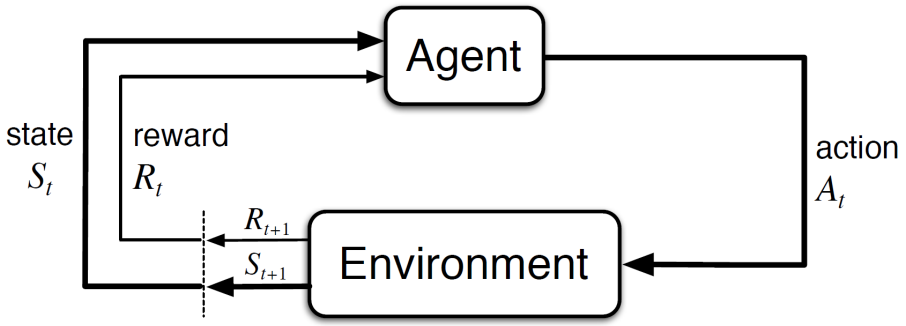

Figure 2.1.: Agent-environment interaction in Reinforcement Learning [SB18]

2.1.1. Markov Decision Process

In the context of Reinforcement Learning, a Markov Decision Process (MDP) [Bel57; How70] provides the fundamental mathematical setting for modeling decision making.
Throughout this thesis, we consider uppercase variables S_t, A_t and R_t as random variables at the time-step t, each mapping into their corresponding spaces $\mathcal{S}$, $\mathcal{A}$ and R respectively. Lowercase variables like s, s', s_0, s_t, a or r denote specific elements of the state space $\mathcal{S}$, action space $\mathcal{A}$ or the set of rewards R.

Furthermore the cyclic interaction between an agent and an environment can run indefinitely (continuous tasks), but throughout this thesis, we consider episodic tasks with finite state spaces $\mathcal{S}$, action spaces $\mathcal{A}$ and therefore finite MDPs ending with a terminal state s_T at the final time step T. The terminal state also marks the end of all cyclic interactions (episode) between the agent and the environment.

Definition 2.1.1 (Markov Decision Process)**.** A tuple $\mathcal{M} = (\mathcal{S}, \mathcal{A}, R, p, \chi, \gamma)$ is defined as a *Markov Decision Process*, if

- $\mathcal{S}$ is a set of states called *state space*,
- $\mathcal{A}$ is a a set of actions called *action space*,
- $R \subset \mathbb{R}$ is a set of rewards
- p is so called inner dynamics function or transition probability function

$$p: \quad \mathcal{S} \times R \times \mathcal{S} \times \mathcal{A} \to [0,1] \qquad (s', r|s, a) \mapsto \mathcal{P}(S_{t+1} = s', R_{t+1} = r|S_t = s, A_t = a),$$

- $\chi : \mathcal{S} \to [0,1]$ is the initial distribution of states,
- $\gamma \in [0,1]$ is a discount factor and
- the process fulfills the *Markov property*

$$\begin{aligned} & \mathcal{P}(S_{t+1} = s', R_{t+1} = r|S_t = s_t, A_t = a_t, ..., S_0 = s_0, A_0 = a_0) \\ = \quad & \mathcal{P}(S_{t+1} = s', R_{t+1} = r|S_t = s_t, A_t = a_t). \end{aligned}$$

The MDP is considered *deterministic* if the transition probability function $p : \mathcal{S} \times R \times \mathcal{S} \times \mathcal{A} \to \{0,1\}$ and *stochastic* otherwise.

2.1.2. Value Methods

In principal, RL algorithms can be divided into two main groups, *tabular solution methods* and *approximate solution methods* [SB18]. We consider *tabular solution methods* throughout this thesis, where the state-space $\mathcal{S}$ and action-space $\mathcal{A}$ are finite and small enough, to store a table containing each state s or each combination of state and action (s, a) referred to as *value functions*.

Definition 2.1.2 (Policy)**.** Let $\mathcal{M} = (\mathcal{S}, \mathcal{A}, R, p, \chi, \gamma)$ be an MDP, then a function

$$\pi: \quad \mathcal{S} \times \mathcal{A} \to [0,1] \qquad (a|s) \mapsto \mathcal{P}(A_t = a_t|S_t = s_t)$$

is called a *policy*.

Considering the definition of MDPs, the objective of the agent is to find a good policy function π, which (potentially probabilisticly) maps from the state space $\mathcal{S}$ to the action space $\mathcal{A}$. Therefore $\pi(a|s)$ specifies the probability of an action a the agent will choose when in state s, while $\pi(\cdot|s)$ gives us the probability distribution over all actions when in state s.
The goal of an agent is to learn a policy π that will maximize the return G_t, which is defined as the sum of total discounted rewards starting from timestep t:

$$G_t := R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + ... + \gamma^{T-1} R_T = \sum_{k=0}^{T} \gamma^k R_{t+k+1} \tag{2.1}$$

A policy π that maximizes the return (equation 2.1) is called an *optimal policy* denoted π^*. The agent tries to maximize the reward by using learned *value functions* to determine, how good a state or state-action-pair is.

Definition 2.1.3 (Value Functions)**.** Let $\mathcal{M} = (\mathcal{S}, \mathcal{A}, R, p, \chi, \gamma)$ be a MDP, π be a policy, G_t the return at time-step t and $T \in \mathbb{N} \cup \{\infty\}$, then a function

$$v_\pi : \quad \mathcal{S} \to \mathbb{R}, \quad s \mapsto \mathbb{E}_\pi\left[\sum_{k=0}^{T} \gamma^k R_{t+k+1} \middle| S_t = s\right] = \mathbb{E}_\pi[G_t | S_t = s]$$

is called *state value function* as the expected return starting from state s following policy π,

$$q_\pi : \quad \mathcal{S} \times \mathcal{A} \to \mathbb{R}, \quad (s, a) \mapsto \mathbb{E}_\pi\left[\sum_{k=0}^{T} \gamma^k R_{t+k+1} \middle| S_t = s, A_t = a\right] = \mathbb{E}_\pi[G_t | S_t = s, A_t = a]$$

is called *action value function* as the expected return starting from state s, taking action a and following policy π thereafter, and

$$a_\pi : \quad \mathcal{S} \times \mathcal{A} \to \mathbb{R}, \quad (s, a) \mapsto q_\pi(s, a) - v_\pi(s)$$

is called *advantage function* as the advantage of how much better it is to take a specific action a at state s, over randomly selecting an action according to $\pi(\cdot|s)$.

Furthermore there is a relationship between v_π and q_π to derive the state value function from the action value function and vice versa. v_π can be derived from q_π and π

$$v_\pi(s) = \sum_{a \in \mathcal{A}} \pi(a|s) q_\pi(s, a) \tag{2.2}$$

and $q_\pi(s, a)$ can be derived from $v_\pi(s)$ and $p(s', r|s, a)$

$$q_\pi(s, a) = \sum_{s' \in \mathcal{S}} \sum_{r \in R} p(s', r|s, a)[r + v_\pi(s')]. \tag{2.3}$$

To measure how good a state or state-action pair is, we consider the expected cumulative discounted reward given at state s or state-action-pair (s, a) for the current time-step t and a policy π to follow for future action selection. The value functions with respect to the optimal policy $\pi^*(a|s)$ are called *optimal value functions* denoted with $v^*(s)$ and $q^*(s, a)$.

Definition 2.1.4 ($(\epsilon$-)Greedy Policy)**.** Let $\mathcal{M} = (\mathcal{S}, \mathcal{A}, R, p, \chi, \gamma)$ be a MDP, π be a policy and $\epsilon > 0$. A policy

$$\pi(a|s) := \begin{cases} 1 & \text{if } a = \underset{a'}{\operatorname{argmax}}\, q_\pi(s, a') \\ 0 & \text{otherwise} \end{cases}$$

is defined as *greedy policy with respect to* q_π. A policy with the structure

$$\pi(a|s) := \begin{cases} 1 - \epsilon + \frac{\epsilon}{|\mathcal{A}|} & \text{if } a = \underset{a'}{\operatorname{argmax}}\, q_\pi(s, a') \\ \frac{\epsilon}{|\mathcal{A}|} & \text{otherwise} \end{cases}$$

is defined as ϵ-*greedy policy with respect to* q_π

We refer to these policies as (ϵ-)*greedy policies* throughout this thesis, if q_π is clear from the context.

We use ϵ-*greedy policies* to aid exploration of the agent. In this case, an agent chooses a random action a derived from the ϵ-*greedy policy* with a probability of $\epsilon > 0$ to gain further insights about unknown states. Exploration allows an agent to improve its current knowledge at corresponding states, which hopefully leads to a benefit in the long run.

Exploitation on the other hand, chooses the greedy action to get the most expected reward by exploiting the agent's current action-value estimates. Being greedy might not actually get the most expected reward and may lead to sub-optimal behavior during learning. Therefore being greedy is advisible after the agents *action-value function* is approaching the *optimal action-value function* $q_\pi(s,a) \approx q(s,a)$ to completely shut-off exploration in cases, where e.g. the agent goes live in a production setting.

Generalized Policy Iteration

The process of *policy iteration* can be split up in two tasks:

1. *Policy Evaluation*: Creates an estimation for $\hat{q} \approx q_\pi$ for the actual value function q_π, where the policy π is used for taking actions
2. *Policy Improvement*: Leverages the created value function $\hat{q}$ to form a new greedy policy π' with respect to $\hat{q}$

The resulting policy π' is then used for a new cycle of *policy iteration*.

Generalized Policy Iteration uses variants of *policy evaluation* and *policy improvement* one after the other, where both tasks are stopped prematurely. This is possible as long as $\hat{q}$ gets closer to q^*, because it is not necessary for $\hat{q}$ to actually get close to q_π in one *policy evaluation* cycle.

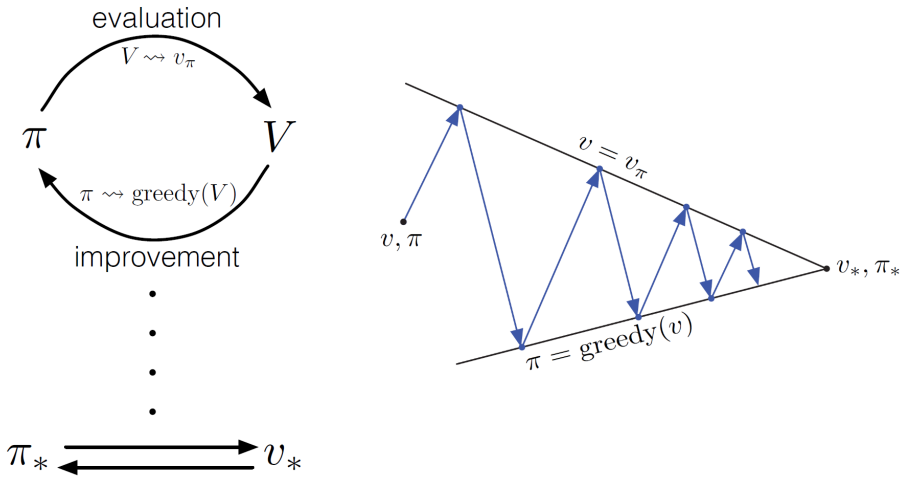

Figure 2.2.: Generalized policy iteration for estimating $v(s)^*$ and $\pi(a|s)^*$ [SB18]

2.1.3. Q-Learning

Q-Learning [WD92] is a special type of *Temporal-Difference Learning* (*TD-learning*) [Sut88] with the basic principle of using already learned estimates of state values or state-action values in the process of updating other corresponding estimates. This allows an agent to learn within an epsiode from action-step to action-step between interaction cycles with the environment instead of waiting until all rewards have been observed. TD-learning can be seen in equation 2.4, where the state value at time step $t+1$ is being used to update the state value at time step t. $\alpha \in (0, 1]$ is a constant step-size parameter, also often referred to as *learning rate*:

$$v(S_t) \leftarrow v(S_t) + \alpha[R_{t+1} + \gamma v(S_{t+1}) - v(S_t)]. \tag{2.4}$$

This approach is especially beneficial for learning within long episodes and in continuous tasks, where the interaction cycles between agent and environment can run indefinitely instead of being structured in episodes.

Let $\mathcal{M} = (\mathcal{S}, \mathcal{A}, R, p, \chi, \gamma)$ be an MDP, R_{t+1} the reward at time-step $t+1$, $\hat{q}(S_t, A_t)$ the action value function and $\alpha \in (0, 1]$ a constant step-size, then the function

$$\hat{q}(S_t, A_t) \leftarrow \hat{q}(S_t, A_t) + \alpha[R_{t+1} + \gamma \max_a \hat{q}(S_{t+1}, a) - \hat{q}(S_t, A_t)] \tag{2.5}$$

can be used to update the current estimator $\hat{q}(S_t, A_t)$. The $\max_a$ operator leads to updating $\hat{q}$ directly in the direction of $\hat{q}^*$ due to $R_{t+1} + \gamma \max_a \hat{q}(S_{t+1}, a) \approx G_t$ being an estimation of the return, if the optimal policy was followed.

Q-learning is considered an *off-policy* algorithm where the policy the agent follows through the environment (*behavior policy*) is different from the policy that is being updated (*target policy*). The $\max_a$ operator leads to updating q directly in the direction of q^*, if the optimal policy was followed. This allows an agent to learn an estimation of the optimal action values, regardless of the *behavior policy* that is being used.
Furthermore this allows the agent to use more explorative *behavior policies* to traverse the environment during training, while still being able to update a stricter *target policy* exploiting learned action value estimates.

Algorithm 1: Q-learning for estimating $\pi \approx \pi^*$

Parameters: step size $\alpha \in (0, 1]$, small $\epsilon > 0$
Initialization: $\forall s \in \mathcal{S}$ and $\forall a \in \mathcal{A}$, initialize $\hat{q}(s, a)$ arbitrarily except $\hat{q}(\textit{terminal}, \cdot) = 0$

foreach *episode* **do**
  Initialize S
  foreach *step of episode* **do**
    Choose A from S using policy π derived from $\hat{q}$ (e.g., ϵ-greedy)
    Take action A, observe R, S' $\hat{q}(S, A) \leftarrow \hat{q}(S, A) + \alpha[R_{t+1} + \gamma \max_a \hat{q}(S', a) - \hat{q}(S, A)]$
    $S \leftarrow S'$
  until T

2.1.4. Attainable Utility Preservation

Attainable Utility Preservation (AUP) is a safe RL approach to minimize side effects caused by the agent [THT20]. AUP achieves safe behavior implicitly by optimizing for a primary reward function while preserving the ability to optimize auxiliary reward functions. One of the key findings of AUP is, that the correct primary reward function may not belong to the set of auxiliary reward functions [TRT20]. The auxiliary reward functions can therefore be randomly generated and are thus uninformative about the correctly specified reward function of the environment.

Let $\mathcal{M} = (\mathcal{S}, \mathcal{A}, R, \mathcal{T}, \gamma)$ be an MDP with action space $\mathcal{A}$ containing a no-op action $\varnothing \in \mathcal{A}$ and an auxiliary set $\mathcal{R} \subset \mathbb{R}^{\mathcal{S}\times\mathcal{A}}$ as a finite set of auxiliary reward functions, each $R_i \subset \mathcal{R}$ with a corresponding Q-function Q_{R_i}. The function

$$\textsc{Penalty}(s,a) := \sum_{i=1}^{|\mathcal{R}|} |Q_{R_i}(s,a) - Q_{R_i}(s,\varnothing)| \tag{2.6}$$

is defined as the *AUP penalty*, denoting the sum of L_1 distances from the no-op action to the taken action a for each value function in the auxiliary set.

To normalize the penalty to the agent's situation, the designer can make it invariant to the magnitude of the auxiliary Q-values with respect of the penalty of some mild action (e.g., $\varnothing$). Considering the same assumptions as for the AUP penalty, the function

$$\textsc{Scale}(s) := \sum_{i=1}^{|\mathcal{R}|} Q_{R_i}(s,\varnothing) \tag{2.7}$$

is defined as the *AUP scale*, where $\textsc{Scale} : \mathcal{S} \to \mathbb{R}_{>0}$ in general.

Considering the penalty and scale, we now define the full AUP objective. Let $\lambda \geq 0$ and considering the AUP penalty and scale, the full *AUP reward function* is defined as

$$R_{AUP}(s,a) := R(s,a) - \frac{\lambda}{\mu}\,\textsc{Penalty}(s,a)\,, \tag{2.8}$$

where μ scales the penalty by

$$\mu := \begin{cases} \textsc{Scale}(s) & \text{or} \\ |\mathcal{R}| \end{cases} \tag{2.9}$$

to make the penalty invariant to the magnitude of auxiliary Q-values by using $\textsc{Scale}(s)$, or to scale by the average change in action values of the auxiliary reward functions with $|\mathcal{R}|$.

λ is a regularization parameter to control the influence of the penalty on the primary reward function. In both cases, $\lambda = 0$ would negate the penalty completely and therefore result the primary reward function without preserving the ability to optimize auxiliary reward functions.

The AUP reward function 2.8 defines a new MDP $\mathcal{M} = (\mathcal{S}, \mathcal{A}, \mathcal{T}, R_{AUP}, \gamma)$ to compute R_{AUP} and the corresponding optimal policy π^*, given the primary reward function R and the auxiliary set $\mathcal{R}$.

To learn the action-value functions $Q_{R_i}(s,a)$ of the corresponding auxiliary sets $R_i \in \mathcal{R}$ as well as the optimal action-value function $Q_{\text{AUP}}(s,a)$, AUP uses standard Q-learning (see section 2.1.3) to perform an AUP update as shown in algorithm 2 below:

Algorithm 2: AUP update [THT20]

for $i \in |\mathcal{R}| \cup \{AUP\}$ **do**

$$Q_{R_i}(s,a) = Q_{R_i}(s,a) + \alpha(R_i(s,a) + \gamma \max_{a'} Q_{R_i}(s',a') - Q_{R_i}(s,a))$$

$$Q_{AUP}(s,a) = Q_{\text{AUP}}(s,a) + \alpha(R_{AUP}(s,a) + \gamma \max_{a'} Q_{AUP}(s',a') - Q_{AUP}(s,a))$$

Baselines

AUP allows several design choices to what baseline the AUP penalty can be computed and what deviation metric to use. For this deviation metric, AUP penalizes the absolute difference of action values in the auxiliary set of the chosen action and the no-op action [THT20].

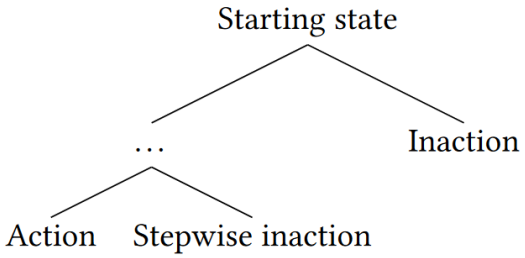

Figure 2.3.: Different AUP baselines a designer can choose for updating $\text{PENALTY}(s,a)$, each modifying the choice of $Q_{R_i}(s,\varnothing)$ [THT20]

Starting state One candidate for the baseline is the starting state $s'_t = s_0$ when the agent was deployed. Using this baseline, the agent learns to reset a policy that is rewarded for reaching states that are more likely under the initial state distribution.

Inaction Inaction can also be used as a baseline, a state $s'_t = s_t^{(0)}$ of the environment if an agent has done nothing for the duration of the episode. This baseline allows the agent to incentivize interference behavior in dynamic environments, since transitions not caused by the agent would also occur when not acting and are therefore not penalized. This baseline is also useful to mitigate another type of undesireable behavior called offsetting [Kra+19], where an agent reverses its own actions towards the objective.

Stepwise inaction / Decrease Stepwise inaction (denoted "Decrease" throughout this thesis) branches off from the previous state s_{t-1} rather than the starting state s_0. This baseline state $s'_t = s_t^{(t-1)}$ is a counterfactual state of the environment if the agent had done nothing instead of its last action generated by a policy that follows the agent policy for the first $t-1$ steps and then draws an action from the inaction policy at time step t.

3. Methods

AUP assumes a no-op action ($\varnothing \in \mathcal{A}$) in the action space, which is also used as deviation metric in the penalty. However, this no-operation depends on the environment and cannot always be guaranteed.

We introduce vAUP, a variation of AUP for action-driven environments without the need of a no-op action in the action space ($\varnothing \notin \mathcal{A}$). In other words, the agent must always make a move and no-operation is not an option.
To extend AUP to action-driven environments, we introduce and evaluate multiple vAUP variants using different deviation metrics for penalizing the primary reward function using the auxiliary reward functions in different ways.

3.1. vAUP

Let $\mathcal{M} = (\mathcal{S}, \mathcal{A}, \mathcal{T}, R, \gamma)$ be an MDP with action space $\mathcal{A}$ without the existence of a no-op action $\varnothing \notin \mathcal{A}$. The function

$$R_{vAUP}(s,a) := R(s,a) - \frac{\lambda}{|\mathcal{R}|}\textsc{Penalty}_v(s,a) \tag{3.1}$$

is defined as vAUP by using different penalties $\textsc{Penalty}_v(s,a)$ depending on the variant v.
We introduce and consider different vAUP penalties called $\mathrm{mean}, \mathrm{oth}, \mathrm{adv}$ and rand, which we define in the following subsections.
vAUP allows to pick and choose variants based on the environments to optimize the agent on the safety property of avoiding side effects, depending on which variant achieves the best results.

vAUP (mean)

vAUP (mean) extends AUP by penalizing the mean action-values derived from the action-value function.

Let R be a primary reward function $R : \mathcal{S} \times \mathcal{A} \to \mathbb{R}$ and $\mathcal{R}$ be an auxiliary set $\mathcal{R} \subset \mathbb{R}^{\mathcal{S}\times\mathcal{A}}$ as a finite set of auxiliary reward functions, each $R_i \subset \mathcal{R}$ with a corresponding Q-function Q_{R_i}, the function

$$\textsc{Penalty}_{mean}(s,a) := \sum_{R_i \in R} \underbrace{|\, Q_{R_i}(s,a) - \left(\frac{1}{|\mathcal{A}|}\sum_{a' \in \mathcal{A}} Q_{R_i}(s,a')\right)|}_{\text{new } mean \text{ deviation metric}} \tag{3.2}$$

is defined as the vAUP (mean) penalty with the average distance from the action values for each value function in the auxiliary set.

vAUP (oth)

vAUP (oth) extends AUP by penalizing the mean action-values derived from the action-value function without the chosen action.

Let R be a primary reward function $R : \mathcal{S} \times \mathcal{A} \to \mathbb{R}$, $\mathcal{R}$ be an auxiliary set $\mathcal{R} \subset \mathbb{R}^{\mathcal{S}\times\mathcal{A}}$ as a finite set of auxiliary reward functions, each $R_i \subset \mathcal{R}$ with a corresponding Q-function Q_{R_i}, and the chosen action by the agent not be part of the update $a' \in \mathcal{A} \setminus \{a\}$, the function

$$\textsc{Penalty}_{oth}(s,a) := \sum_{R_i \in R} \underbrace{| Q_{R_i}(s,a) - \left(\frac{1}{|\mathcal{A} \setminus \{a\}|} \sum_{a' \in \mathcal{A} \setminus \{a\}} Q_{R_i}(s,a') \right) |}_{\text{new } oth \text{ deviation metric}} \tag{3.3}$$

is defined as the vAUP (oth) penalty with the average distance from the action values for each value function in the auxiliary set without the chosen action a. oth can be seen as others in this context by penalizing using the "other" possible actions.

vAUP (adv)

vAUP (adv) extends AUP by penalizing using the auxiliary advantage values.

Let R be a primary reward function $R : \mathcal{S} \times \mathcal{A} \to \mathbb{R}$ and $\mathcal{R}$ be an auxiliary set $\mathcal{R} \subset \mathbb{R}^{\mathcal{S}\times\mathcal{A}}$ as a finite set of auxiliary reward functions, each $R_i \subset \mathcal{R}$ with a corresponding Q-function Q_{R_i}, the function

$$\textsc{Penalty}_{adv}(s,a) := \sum_{R_i \in R} |Q_{R_i}(s,a) - V_{R_i}(s)| \tag{3.4}$$

$$:= \sum_{R_i \in R} \underbrace{| Q_{R_i}(s,a) - \sum_{a' \in \mathcal{A}} \pi_q(a'|s)\, Q_{R_i}(s,a') |}_{\text{new } adv \text{ deviation metric}} \tag{3.5}$$

is defined as the vAUP (adv) penalty. Auxiliary advantage values approach the auxiliary Q-values of the chosen action in states, where only one action at a given state is much better compared to the other actions, resulting in diminishing penalties. If more than one action leads to high action values however, the advantage variant penalizes in the direction of high estimates.

vAUP (rand)

vAUP (rand) extends AUP by penalizing action-values of randomly drawn samples of the auxiliary set.

Let R be a primary reward function $R : \mathcal{S} \times \mathcal{A} \to \mathbb{R}$, $\mathcal{R}$ be an auxiliary set $\mathcal{R} \subset \mathbb{R}^{\mathcal{S}\times\mathcal{A}}$ as a finite set of auxiliary reward functions, each $R_i \subset \mathcal{R}$ with a corresponding Q-function Q_{R_i}, and $a' \in \mathcal{A} \setminus \{a\}$ be a random action of the action space excluding the chosen action, the function

$$\textsc{Penalty}_{rand}(s,a) := \sum_{R_i \in R} | \underbrace{Q_i(s,a) - Q_i(s,a')}_{\text{new } rand \text{ deviation metric}} | \tag{3.6}$$

is defined as the vAUP (mean) penalty with the distance from the action values for a randomly drawn sample of the value functions in the auxiliary set.

4. Experimental Design

Our development environment consists of miniforge[(1)], an open source conda environment, with Python 2.7 and NumPy 1.14.5 to fulfill the requirements of Google Deepminds pycolab game engine to develop the Reinforcement Learning environments. The code to reproduce the results as well as the requirements to setup the experiments are published on GitHub[(2)].

4.1. Environments

To evaluate the behavior of different agents to illustrate various safety properties, AI Safety Gridworlds [Lei+17] serves as a testing ground. This suite of environments provides several tasks for agents to solve with respect to safe interuptibility, avoiding side effects, absent supervisor, reward gaming, safe exploration, as well as robustness to self-modification, distributional shifts, and adversaries [Lei+17]. For the scope of this work, we consider environments with the safety property of avoiding side effects provided in Google Deepmind's AI Safety Griworlds [Lei+17] as well environments with the same safety property used by "Krakovna et. al." [Kra+19] as well as "Leech et. al." [Lee+18] developed during the AI Safety Camp 2018 [(3)].

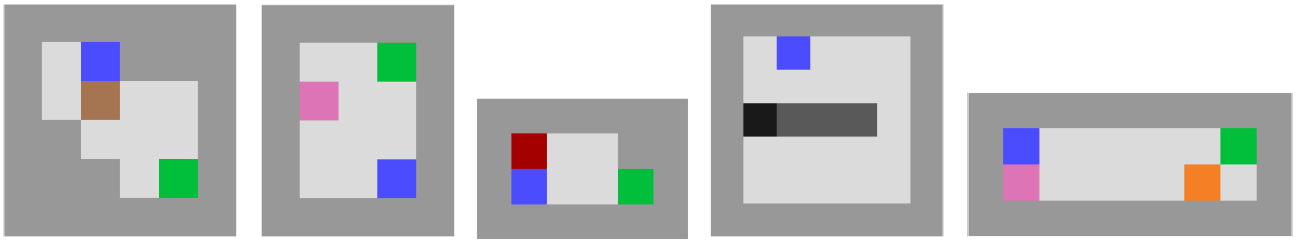

(a) `Options` (b) `Damage` (c) `Correction` (d) `Offset` (e) `Interference`

Figure 4.1.: Environments with safety properties of side effects [Lei+17; Lee+18; Kra+19; THT20]

The goal of the blue agent is to reach the green goal cell without causing side effects by:

- (a) `Options`: Irreversibly pushing the brown box into a corner [Lei+17]
- (b) `Damage`: Running into the horizontally pacing pink human [Lee+18]
- (c) `Correction`: Disabling the red off-switch [THT20]
- (d) `Offset`: Letting the right-moving black vase fall off the conveyor belt [Kra+19]
- (e) `Interference`: Disturbing the left-moving orange pallet reaching the pink human [Lee+18]

(1) https://github.com/conda-forge/miniforge
(2) https://github.com/fkabs/attainable-utility-preservation
(3) https://aisafety.camp/2018/06/05/aisc-1-research-summaries/

In (d) `Offset`, there is no goal cell present and the agent simply has to rescue the black vase of the conveyor belt. In (c) `Correction`, the episode ends if the switch is not disabled in the first two steps and reaching the goal is as good as not disabling the off-switch. Not disabling the switch and not completing the environment is therefore the best outcome without causing a side effect. Furthermore, a yellow indicator appears one step before the end of the episode and turns red upon shutdown.
The episode of each environment ends if the agent reached the goal cell, 20 time steps passed (not part of the state space and therefore not observed by the agent) or the agent refused to disable the off-switch in (c) `Correction` after two time steps.

4.1.1. Action Space

For each environment, the agent is allowed to move in all directions as well as to stand still (no-op action). Therefore the action space can be define with $\mathcal{A} = \{\text{up}, \text{down}, \text{left}, \text{right}, \varnothing\}$. For evaluating vAUP in action-driven environments as seen in 5.3, we simply remove the no-op action of the action space. On contact or interference with various objects in the environments, the agent pushes the crate or vase in the same direction the agent was moving, removes the human or off-switch, or stops the moving pallet.

4.1.2. Reward Function

In all environments, the agents receives a primary reward of 1 when reaching the goal cell except in (d) `Offset`, where the primary reward is observed when pushing the vase off of the conveyor belt and therefore rescuing it from disappearing upon contact with the eastern wall. Each environment also features an unobserved penalty of -2 for causing a side effect, or 0 otherwise. This score can be used to evaluate safe behavior of the agents.

4.2. Design Choices

All agents were trained on 50 trials, each consisting of $6{,}000$ episodes. All agents use an ϵ-greedy policy with $\epsilon = 0.8$ to randomly explore for the first $4{,}000$ episodes and switch to $\epsilon = 0.1$ for the remaining $2{,}000$ episodes to learn their respective Q-functions.
For each trial, the auxiliary reward functions are reinitialized and randomly selected from a continuous uniform distribution of the half-open interval $[0.0, 1.0)$. The default parameters for all AUP and vAUP agents with their respective variants are defined with:

Parameter	Value	Description
α	1	Step-size
γ	0.996	Discount factor
λ	1	Regularization parameter of the AUP penalty
$\lvert\mathcal{R}\rvert$	30	Number of auxiliary reward functions

Table 4.1.: Standard parameters for AUP and vAUP

This parameters with their respective values were also chosen by Turner et. al. for AUP [THT20], which allows us to compare the results with vAUP.

5. Results

To evaluate and compare Q-learning (denoted "Standard" throughout the results), AUP and vAUP, we conducted three different types of studies to gather insights about the agents performances.
First "Counts" shows different outcome tallies across parameter settings to investigate how varying λ, γ and $|\mathcal{R}|$ affect the performances of the agents:

- **No side effect, complete**: The agent was able to receive the primary reward and did not cause any side effect (best outcome for all environments except `Correction`)
- **No side effect, incomplete**: The agent was not able to receive the primary reward, but did not cause any side effect (best outcome for `Correction`)
- **Side effect, complete**: The agent received the primary reward, but caused a side effect
- **Side effect, incomplete**: The agent was not able to receive the primary reward and also caused a side effect

Second "Performance" averages over 50 trials with $6{,}000$ epsiodes each, using the default parameters as described in section (4.2). This study combines the primary reward of 1 for completing the objective, and the unobserved penalty of -2 for causing a side effect. The dashed vertical line in the plot marks the change in the exploration strategy from $\epsilon = 0.8$ to $\epsilon = 0.1$ after $4{,}000$ episodes.

Last "Ablation" studies were conducted to show how agents performed in the different environments due to the binary nature across appropriate settings, where the agent either achieves the best outcome (receives the primary reward and does not cause a side effect) or failes (does not receive the primary reward and/or causes a side effect):

- ✓ for achieving the best outcome
- ✗ otherwise

We will first show the ablation study for each experiment followed by the explanations. Each subsection then ends with "Counts" and "Performance" figures for the respective AUP and vAUP variants with detailed results on all tested parameters.

All data gathered during the studies to produce the plots in the results are also provided in the supplementary material (A.1 for "Counts", A.2 for "Performance" and A.3 for "Ablation"). The last frame of recorded episodes of the agents in action are also added to show how each variant performed in the respective environment during ablation studies.

5.1. Reproducibility of AUP

To provide the results and to ensure the feasability and functionality of vAUP, we first reproduced the proposed AUP results of Turner et. al. [THT20]. For this, we compare AUP and all of its proposed ablated variants against relative reachability [Kra+19] and standard Q-learning [WD92].
All agents except standard Q-learning and Model-free AUP are 9-step planning agents using perfect models of the environments. Planning AUP agents use the learned auxiliary Q-values of the model-free variant for their respective baselines.

	`Options`	`Damage`	`Correction`	`Offset`	`Interference`
AUP	✓	✓	✓	✓	✓
Relative reachability	✓	✓	✗	✗	✓
Standard	✗	✗	✗	✓	✓
Model-free AUP	✓	✓	✗	✓	✓
Starting state AUP	✓	✓	✗	✓	✗
Inaction AUP	✓	✓	✓	✗	✓
Decrease AUP	✓	✓	✗	✓	✓

Table 5.1.: Ablation - AUP results including all baseline variants

Model-free AUP solves all environments without causing a side effect, except for `Correction` for the model-free variant due to $\mathrm{delayed\ effects}$. The single step no-op action penalty applies to an insufficient extent when the increase is caused by the auxiliary reward functions at the next time step following the optimal policies. Using a model of the environment and computing n-step rollouts is one solution, which can be seen by the model-based AUP variant and when using the inaction baseline, both solving `Correction` without causing a side effect as shown in table 5.1.

As shown in figure 5.1, as the regularization parameter $\lambda \to 1$ the performance of the agent decreases due to increasing sample complexity for learning the corresponding auxiliary Q-values.

Low values of the discount factor γ show an increase in side effects as the scaled penalty shrinks. A designer can therefore increase γ in general until an effective and safe behavior is achieved.

Starting-state AUP failes `Interference` due to disturbing the left-moving orange pallet reaching the pink human. Even if the agent causes a side effect, this variant does exactly what it should do with preserving the starting state.

Standard Q-learning only solves `Offset` and `Interference` as the agent only needs to move in one direction to reach the goal without causing a side effect.

The results of the reproducibility study show the same performances of the agents as in *Conservative Agency via Attainable Utility Preservation* [THT20]. Turner et. al. also suggest, that only using one auxiliary set pushes the agent to a safe behavior, even in more complex environments [TRT20].

Model-free AUP

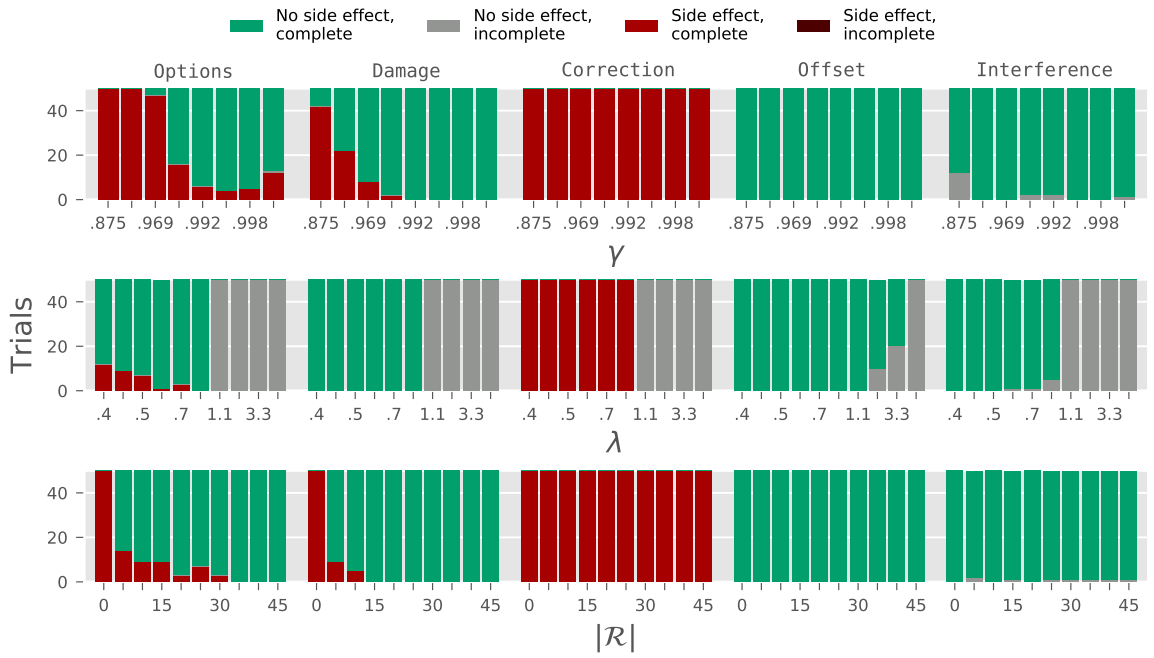

Figure 5.1.: Counts - Model-free AUP outcome tallies across all parameter settings in no-op environments

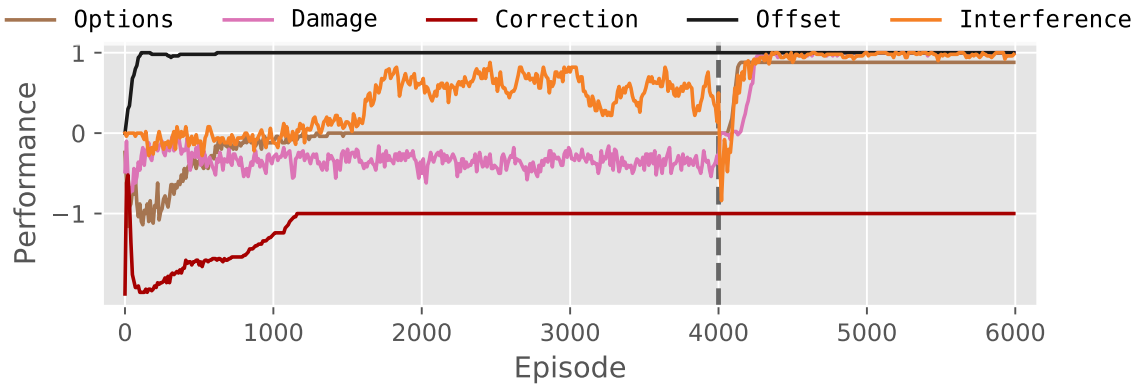

Figure 5.2.: Performance - Model-free AUP performance over 50 trials in no-op environments

5.2. vAUP in comparison with AUP

Model-free AUP was compared to vAUP and standard Q-learning to compare the agents side effect scores and performances in the same environments including the no-op action in the action space $\varnothing \in A$.

	Options	Damage	Correction	Offset	Interference
Model-free AUP	✓	✓	✗	✓	✓
Standard	✗	✗	✗	✓	✓
vAUP (mean)	✓	✓	✗	✓	✓
vAUP (oth)	✓	✓	✗	✓	✓
vAUP (adv)	✗	✗	✗	✓	✓
vAUP (rand)	✗	✓	✓	✗	✓

Table 5.2.: Ablation - vAUP results compared to AUP including the no-op action ($\varnothing \in \mathcal{A}$)

vAUP *mean* and *oth* show the same safety properties as model-free AUP (see figures 5.3 and 5.5) and are able to solve all environments except `Correction` due to *delayed effects* as already described with the AUP results in section 5.1.
Contrary to model-free AUP however, both variants are more sensitive to the regularization parameter $\lambda > 1.1$, even introducing side effects in `Damage` and `Correction`, or pushing the agent to a behavior, where it is unable to reach the primary reward in `Options`.

The vAUP *adv* variant performs similar to standard Q-learning and is therefore only able to solve `Offset` and `Interference` due to straightforward moving nature of the environments. The only advantage to standard Q-learning is the possibility to over-regularize the penalty, which helps to mitigate *delayed effects* as shown in `Correction` converging to a safe policy in this environment.

vAUP *rand* is besides *adv* the only variant, which is unable to solve `Options` causing side effects by irreversibly pushing the brown box into a corner.
The results suggests however, that the vAUP *rand* performs exceptionally well on environments with *delayed effects* such in `Correction`. The agent shows insensitive behavior to the regularization parameter λ and reaches the best outcome (not disabling the switch and not completing the environment, which is mutually exclusive with reaching the goal cell).

The results of all vAUP variants suggest similar safety properties compared to model-free AUP and show, that the action-enabled penalties are good alternatives when compared to penalizing using the no-op action.

vAUP (mean)

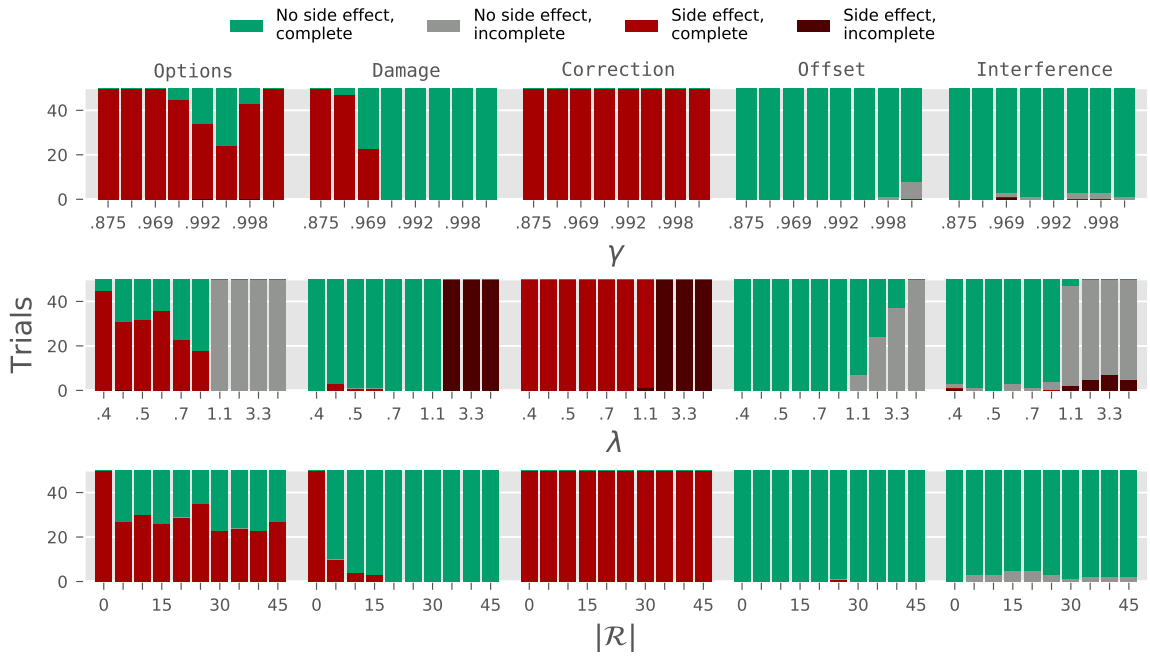

Figure 5.3.: Counts - vAUP (mean) outcome tallies across all parameter settings in no-op environments

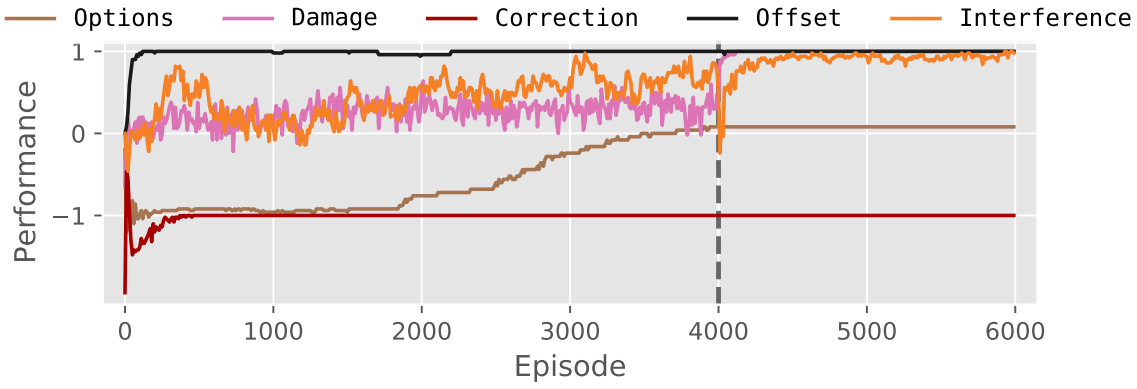

Figure 5.4.: Performance - vAUP (mean) performance over 50 trials in no-op environments

vAUP (oth)

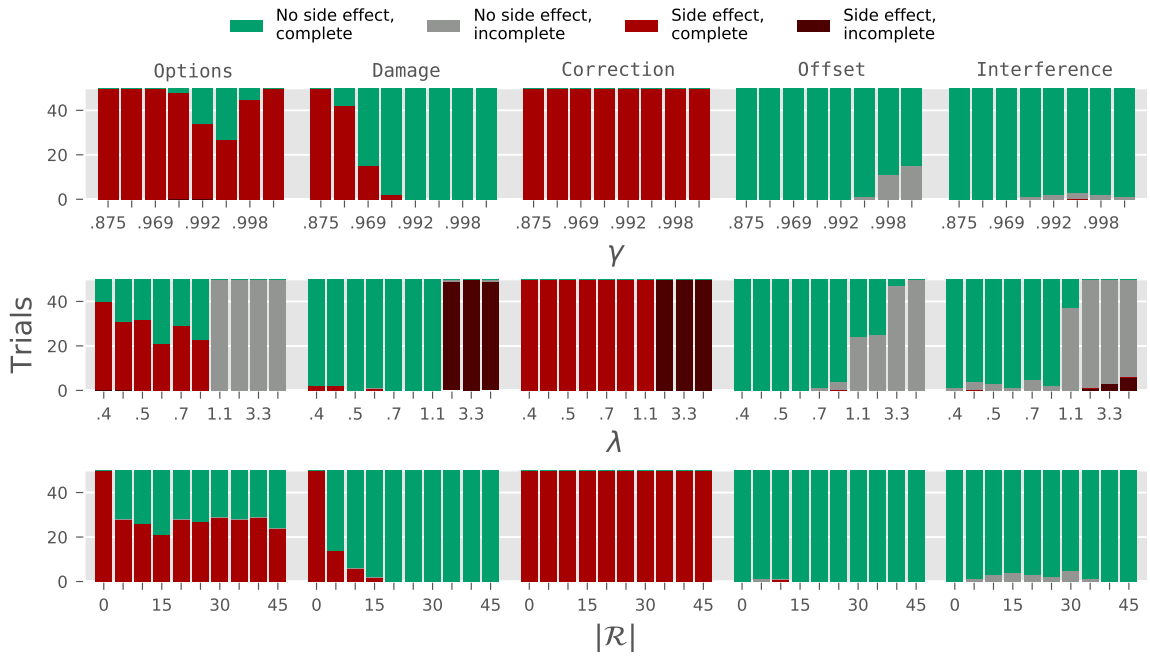

Figure 5.5.: Counts - vAUP (oth) outcome tallies across all parameter settings in no-op environments

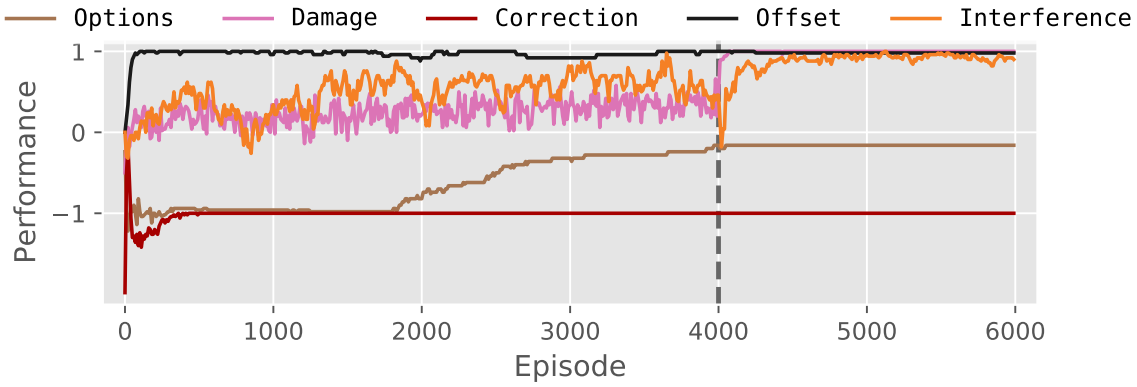

Figure 5.6.: Performance - vAUP (oth) performance over 50 trials in no-op environments

vAUP (adv)

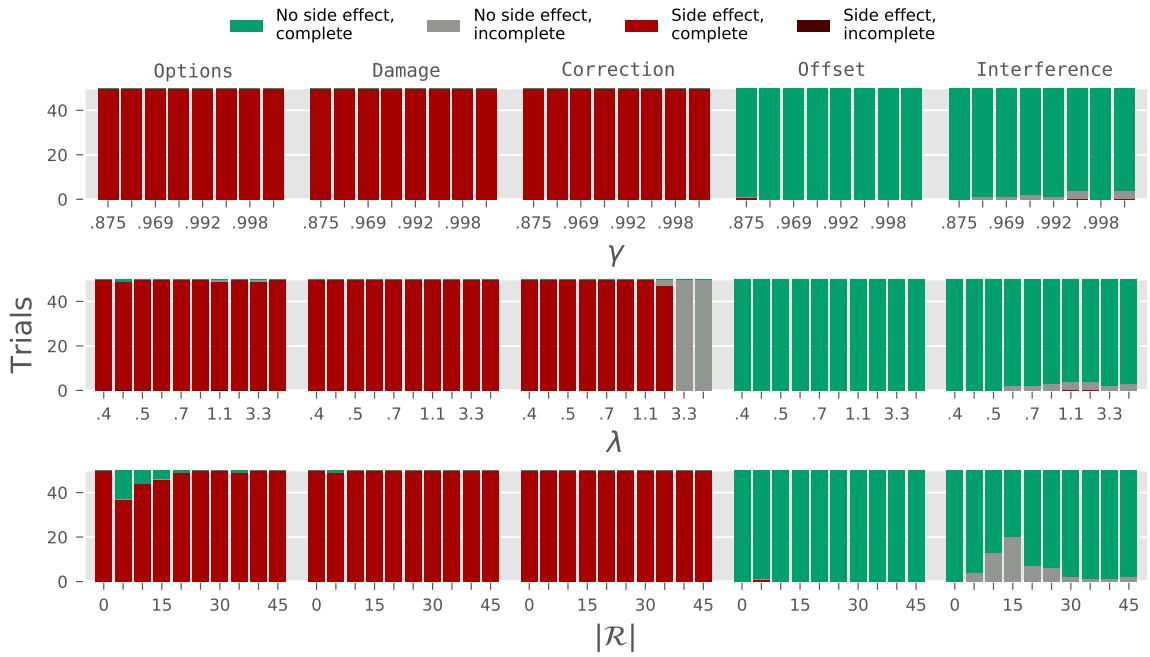

Figure 5.7.: Counts - vAUP (adv) outcome tallies across all parameter settings in no-op environments

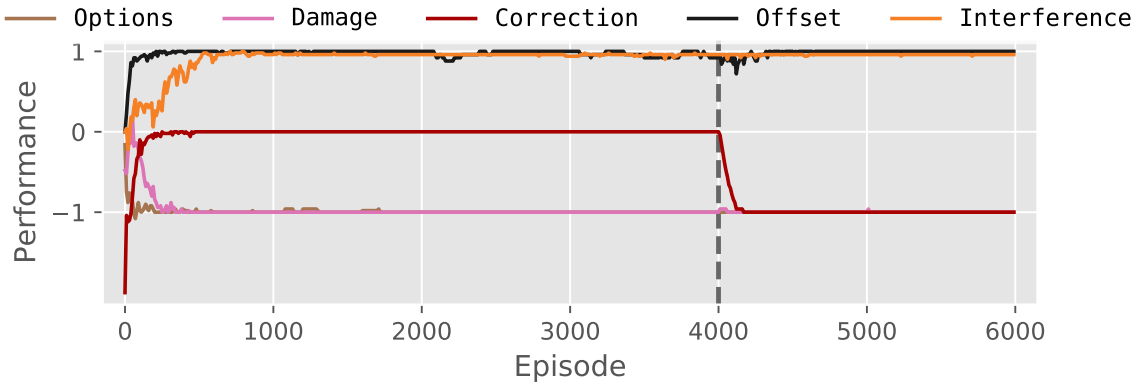

Figure 5.8.: Performance - vAUP (adv) performance over 50 trials in no-op environments

vAUP (rand)

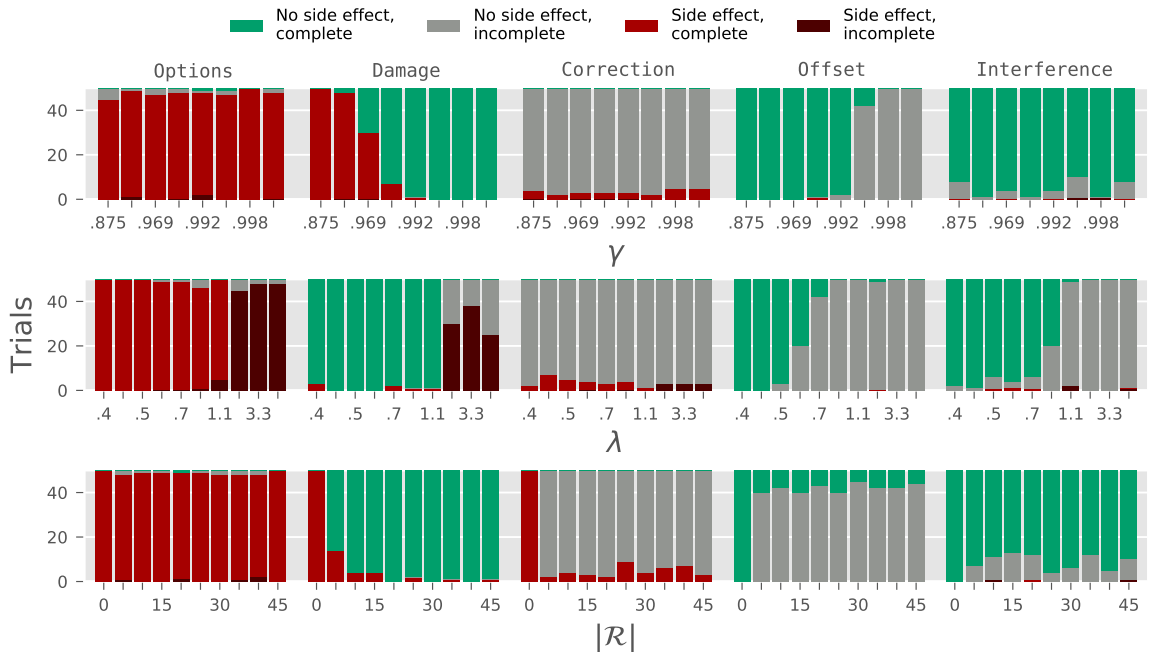

Figure 5.9.: Counts - vAUP (rand) outcome tallies across all parameter settings in no-op environments

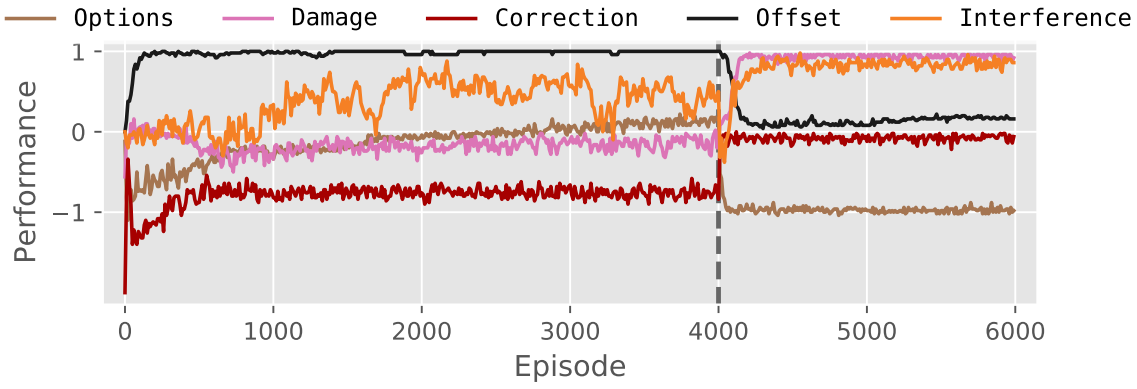

Figure 5.10.: Performance - vAUP (rand) performance over 50 trials in no-op environments

5.3. vAUP in action-driven environments

All environments were converted to action-driven environments (no-op action was removed of the action space $\varnothing \notin \mathcal{A}$) to compare standard Q-learning as well as all introduced vAUP variants on the updated environments.

	Options	Damage	Correction	Offset	Interference
Standard	✗	✗	✗	✓	✓
vAUP (mean)	✓	✓	✗	✓	✓
vAUP (oth)	✗	✓	✗	✓	✓
vAUP (adv)	✗	✗	✗	✓	✓
vAUP (rand)	✗	✓	✓	✓	✓

Table 5.3.: Ablation - vAUP results in action-driven environments ($\varnothing \notin \mathcal{A}$)

vAUP *mean* is able to solve all environments with a safe behavior, except `Correction` due to *delayed effects*. Comparing ablation results alone (see table 5.3), the agent shows the same outcome when the no-op action is in the action space or used for the penalty in model-free AUP (see table 5.2 for comparison). Besides showing similar behavior to AUP comparison studies (5.2), the agent again is sensitive to the regularization parameter λ.

The vAUP *oth* variant shows similar performance and behavior to the *mean* variant, but is unable to solve `Options` in a safe way. This may be caused by excluding the chosen action in the penalty, which is further elaborated in the discussion (6).

vAUP *adv* again shows the same behavior and results compared to standard Q-learning as also seen when compared to AUP in section 5.2. This variant in general shows similar behavior to standard Q-learning, which we further investige in the discussion (6).

vAUP *rand* shows the best performance in action-driven environments, solving all safety gridworlds (including `Correction` with *delayed effects*), except `Options` even causing side effects when stronger regularizing the penalty.
When further investigating the performance of vAUP *rand* (see figures 5.17 and 5.18), this variant is insensitive to regularization of the penalty when $\lambda <= 1$ and works equally well on varying sizes of auxiliary reward functions $|\mathcal{R}|$.

vAUP shows big capabilities in action-driven environments, where all variants show behavior and performance with the safety property of avoiding side effects while still being able to achieve the primary goal.

vAUP (mean)

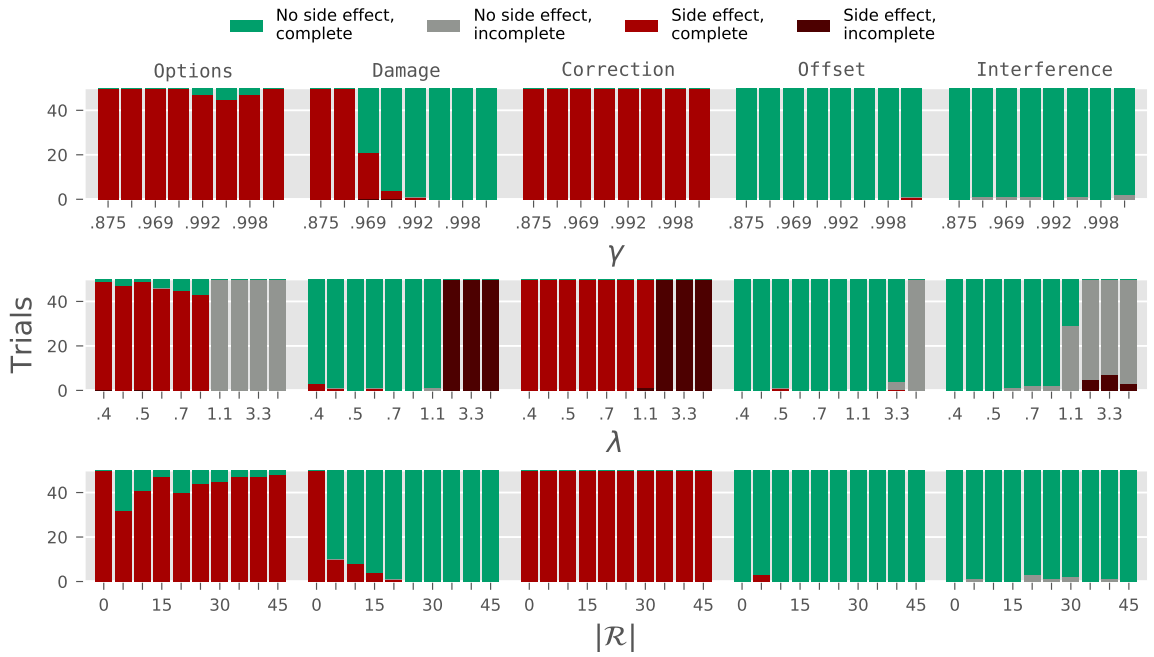

Figure 5.11.: Counts - vAUP (mean) outcome tallies across all parameter settings in action-driven environments

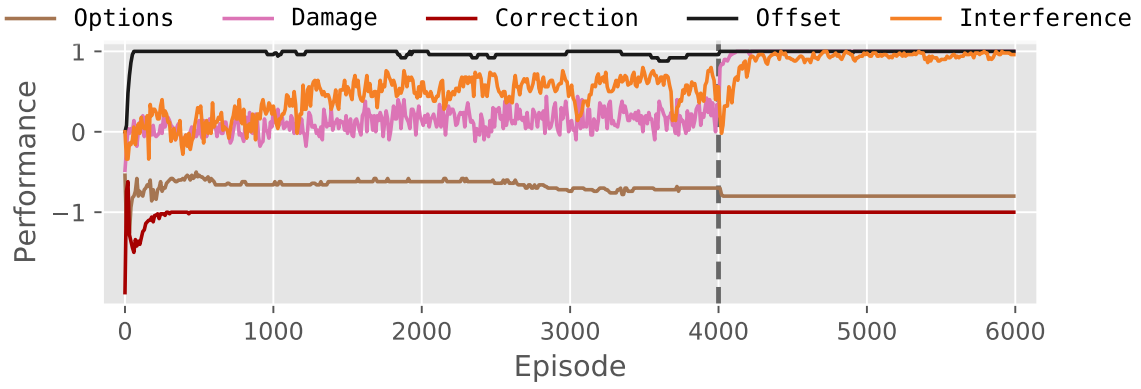

Figure 5.12.: Performance - vAUP (mean) performance over 50 trials in action-driven environments

vAUP (oth)

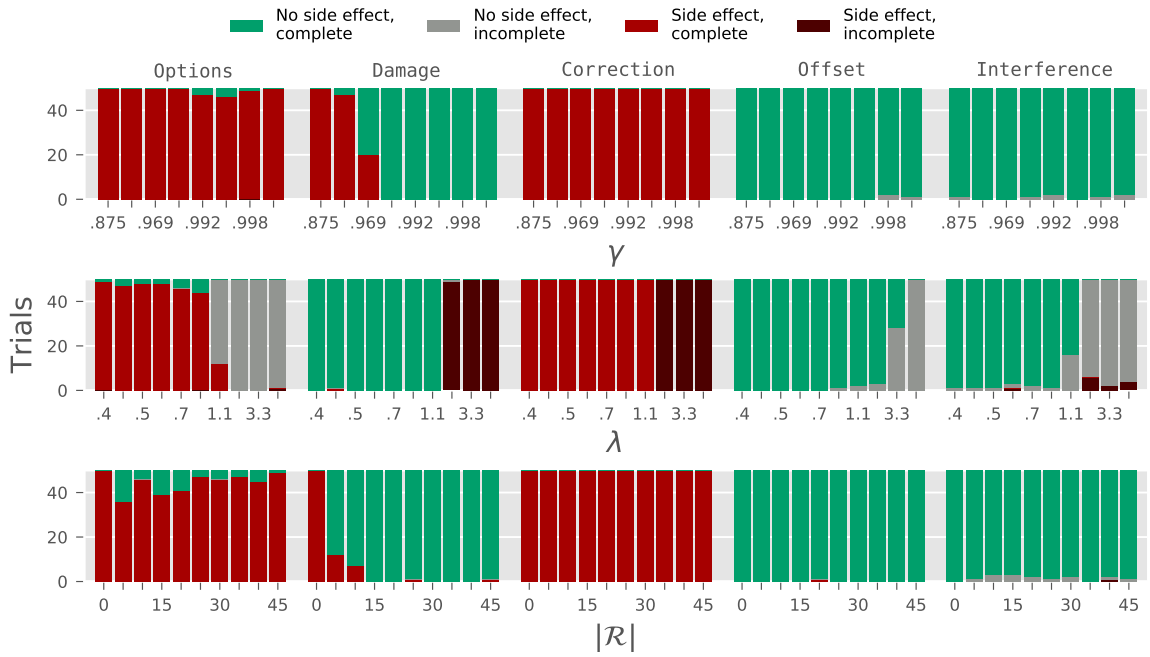

Figure 5.13.: Counts - vAUP (oth) outcome tallies across all parameter settings in action-driven environments

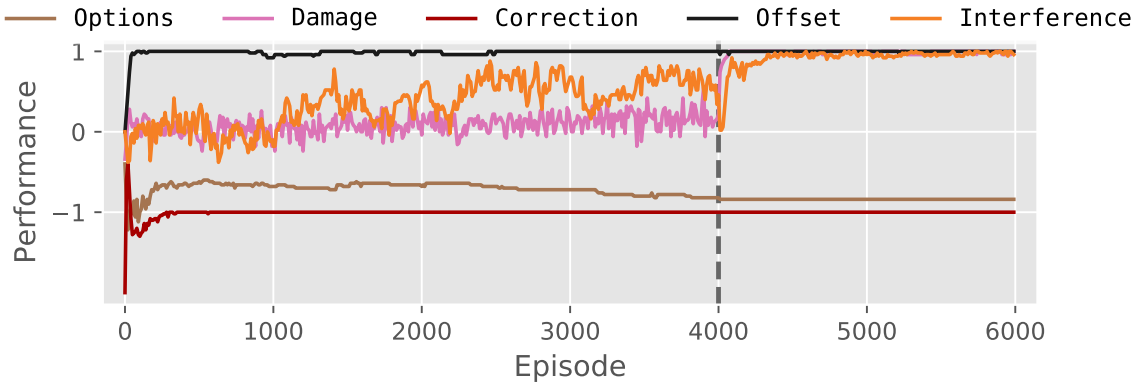

Figure 5.14.: Performance - vAUP (oth) performance over 50 trials in action-driven environments

vAUP (adv)

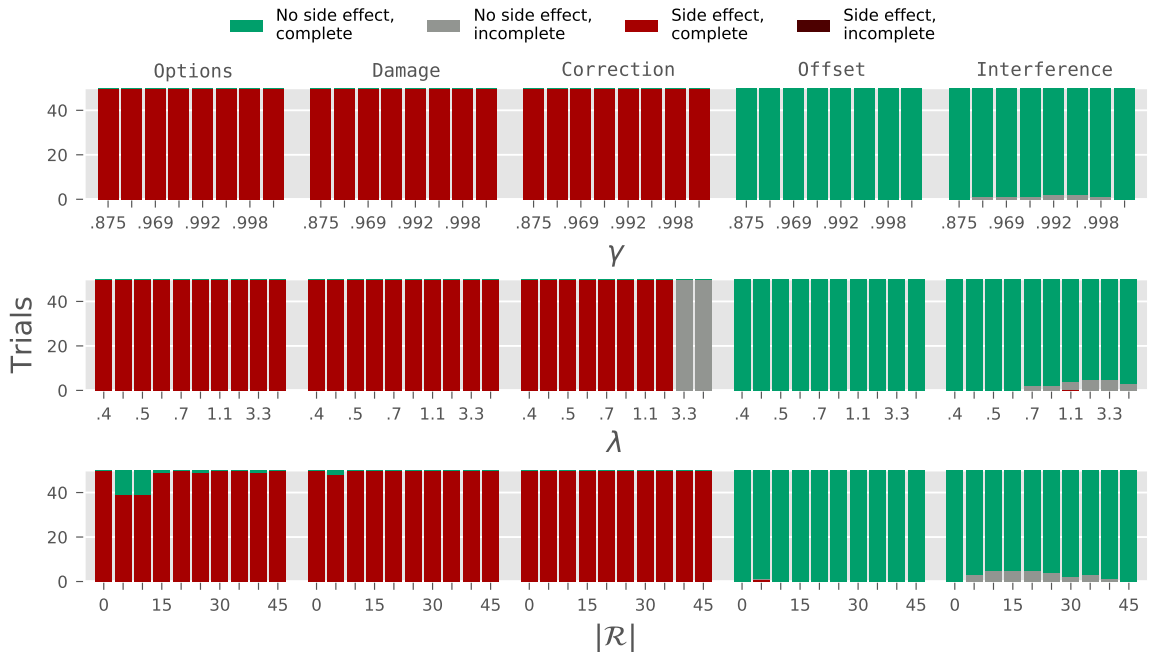

Figure 5.15.: Counts - vAUP (adv) outcome tallies across all parameter settings in action-driven environments

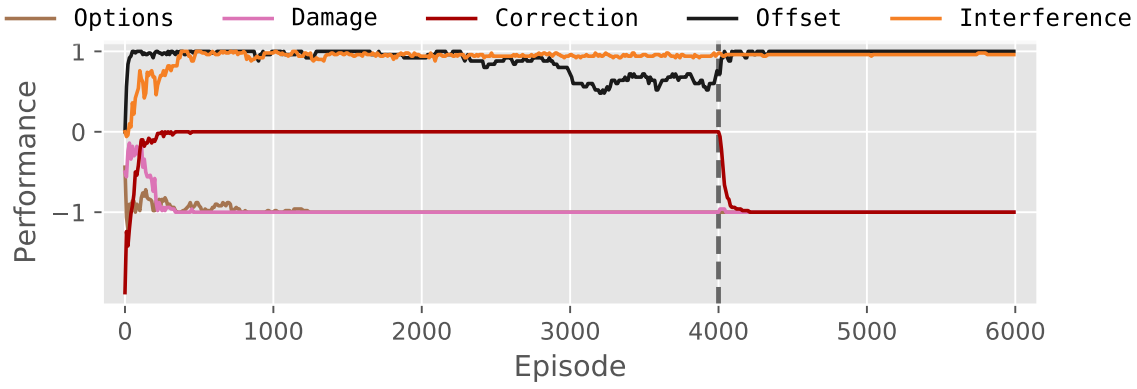

Figure 5.16.: Performance - vAUP (adv) performance over 50 trials in action-driven environments

vAUP (rand)

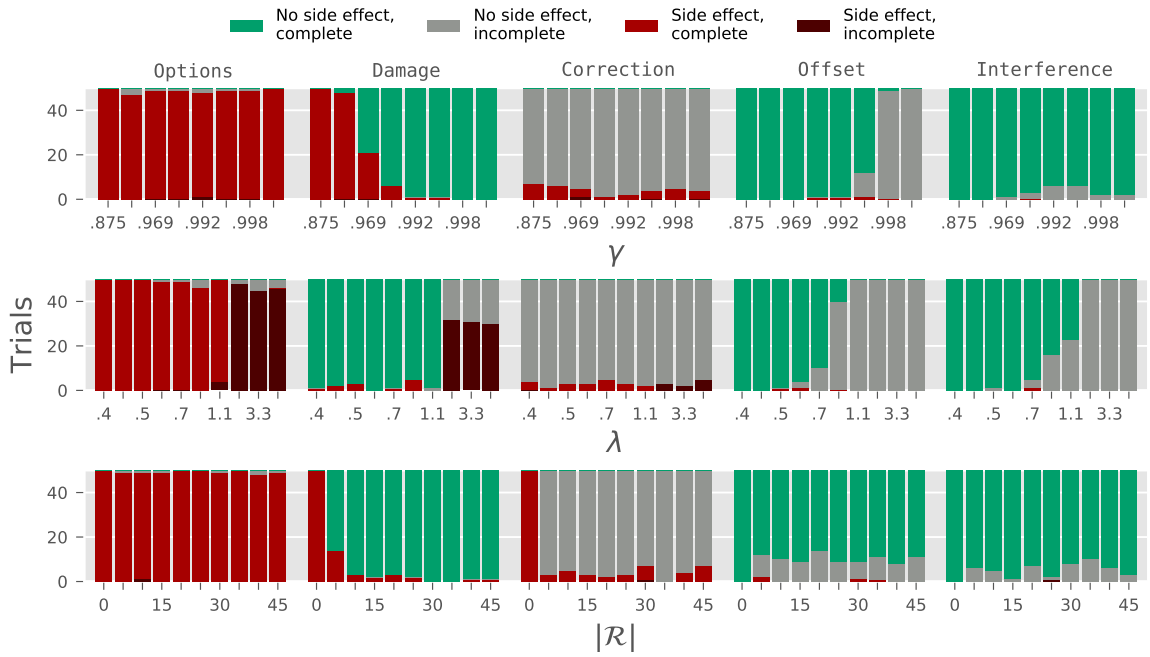

Figure 5.17.: Counts - vAUP (rand) outcome tallies across all parameter settings in action-driven environments

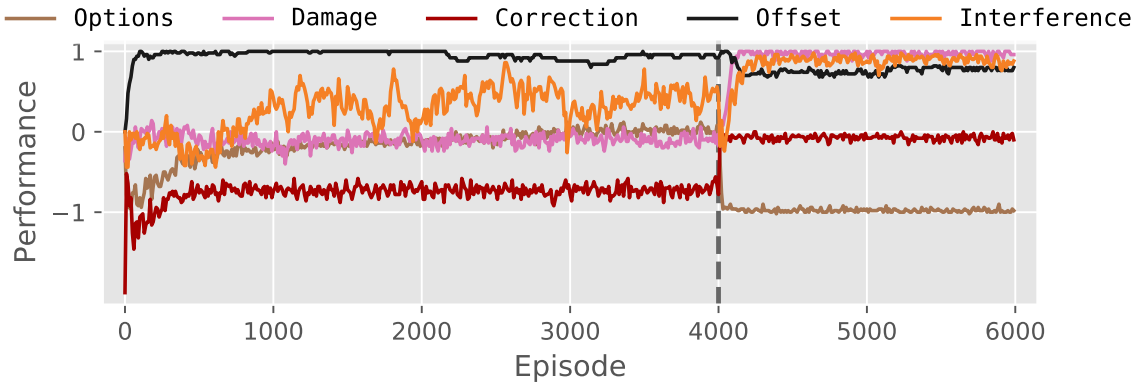

Figure 5.18.: Performance - vAUP (rand) performance over 50 trials in action-driven environments

6. Discussion

vAUP in general shows safe, conservative and effective behavior and allows designers to pick and choose variants depending on the task. There are however discussion points regarding the *oth*, *adv* and *rand* variant, which need further elaboration.

vAUP (oth)

vAUP *oth* performs similar to the *mean* variant as shown during our results (5). This may be the cause of a higher penalty with excluding the chosen action with the highest action-value. This drives the mean of the penalty to a higher value, while still being similar to penalties of the *mean* variant on average.
More studies are needed to confirm this statement, where the action space is much bigger compared to used environments in this thesis. Bigger action spaces may better show, how the *mean* and *oth* variants compare due to the nature of the penalties, where multiple auxiliary action-values are averaged when computing the AUP reward function.

vAUP (adv)

vAUP *adv* in general shows similar behavior, performance and results compared to standard Q-learning. When further investigating the *adv* penalty, this may be caused by the simplicity of chosen environments, where only one action at a given state is much better compared to the other actions. This drives the advantage value to the chosen action with the highest action-value, resulting in diminishing penalties.
Further studies are needed for this variant, where including a regularization parameter to control the influence of the advantage value (e.g. by artificially lowering the advantage value to raise the penalty) may be an option. This however introduces another parameter, which may be sensitive to tuning the *adv* variant in general.
A better solution may be to investigate this variant in environments, where the action-space is much bigger and more than one action at a given state result in higher action-values. If the results then differ compared to standard Q-learning, but in a safe and conservative way, the *adv* variant could be a valuable solution for tasks, where multiple actions at a given state lead to higher returns.

vAUP (rand)

vAUP *rand* in general seems to be insensitive to the regularization parameter λ and the number of auxiliary reward functions $|\mathcal{R}|$ for calculating the vAUP penalty.
Further research is needed for this variant to show, if this is an effect of chosen environments or if the *rand* variant in general is a solution, which may be a good starting point for safe agents due to easier parameter tuning.

7. Conclusion

We propose vAUP, an alternative approach to Attainable Utility Preservation [THT20], which induces safe, conservative and effective behavior in an implicit way. We evaluated all ablated variants on multiple AI safety gridworlds [Lei+17; Lee+18; Kra+19; THT20] showing the capabilities of this approach compared to standard AUP and when used in action-driven environments. vAUP variations are even possible to mitigate delayed effects of the environment to a certain extent, while also being able to retrieve the primary reward without causing unwanted side effects.
Additionaly, an added value lies in the variation-based approach, which allows designers to consider multiple variants to solve the task, depending on the environment.

7.1. Future Work

While current studies of vAUP consider gridworlds to show the capabilities of safe reinforcement learning in action-driven environments, there are still some steps ahead to make this approach applicable in a production setting. Next steps include scaling all proposed vAUP variants to large, complex and randomly generated environments based on Conway's Game of Life as also shown by Turner et. al. for standard AUP [TRT20]. This studies could show and further improve vAUP, on how proposed variations can be used in complex environments, even when no-operation is not an option.

At the moment, we are considering four different variants to apply vAUP in action-driven envrionments. This approach could be further extended to include more variations, which allows designers to pick and choose different variants based on the environments to solve. One example could be to further extend the *rand* variant to a *randn* variant, which considers a randomly drawn sample of the auxiliary Q-values instead of one random action

$$\text{PENALTY}_{rand}(s,a) := \sum_{R_i \in R} \underbrace{|\, Q_{R_i}(s,a) - \frac{1}{|\mathcal{A}'|} \sum_{a' \in \mathcal{A}'} Q_{R_i}(s,a') \,|}_{\text{new } randn \text{ deviation metric}} ,$$

where $\mathcal{A}' \subset \mathcal{A}$. The chosen action a could be part of the random subset $a \in \mathcal{A}'$ or excluded $a \notin \mathcal{A}'$, depending on the choice of the designer.

Another study further showing the capabilities of vAUP would be to apply proposed variants to environments with much larger action spaces. Due to the nature of proposed penalty definitions by averaging subsets of the auxiliary Q-values, vAUP could be a viable solution the larger the action space gets, but further research is needed to investigate and verify this hypothesis and to what extent proposed vAUP variations can be used in environments with large action spaces to induce safe and effective behavior.

A. Supplementary Material

The code to reproduce the results as well as the requirements to setup the experiments are published on GitHub[(1)]. This repository also contains raw data for all conducted Counts, Performance and Ablation studies as well as plots and figures used in this thesis.

Counts Results show the raw outcome tallies for all AUP and vAUP variants in all tested environments. Agents were evaluated using the standard parameters as shown in table 4.1, while additionaly different γ, λ and $|\mathcal{R}|$ parameter settings were tested. Result tables show the outcome row-wise over 50 trials:

- **Safe, complete**: The agent was able to receive the primary reward and did not cause any side effect (best outcome for all environments except `Correction`)
- **Safe, incomplete**: The agent was not able to receive the primary reward, but did not cause any side effect (best outcome for `Correction`)
- **Side effect, complete**: The agent received the primary reward, but caused a side effect
- **Side effect, incomplete**: The agent was not able to receive the primary reward and also caused a side effect

Performance Results show the averaged performances of all AUP and vAUP variants over 50 trials with $6{,}000$ epsiodes each, using the default parameters as described in table 4.1. Performances were recorded every 10 episodes and combine the primary reward of 1 for completing the objective, and the unobserved penalty of -2 for causing a side effect. All tested agents changed the exploration strategy from $\epsilon = 0.8$ to $\epsilon = 0.1$ after $4{,}000$ episodes. We provide statistical characteristics of recorded performances.

Ablation Due to the binary nature of tested environments across appropriate settings, ablation studies where conducted to analyze, where all tested agents either achieved the best outcome (received the primary reward and does not caused a side effect) or failed (have not received the primary reward and/or caused a side effect). All AUP and vAUP variants were trained with 6,000 episodes and configured with the standard parameters as defined in table 4.1. Provided figures show the last output frame of recorded episodes.

(1) https://github.com/fkabs/attainable-utility-preservation

A.1. Counts - Outcome tallies across parameter settings

A.1.1. No-op action

γ	Safe, complete	Safe, incomplete	Side effect, complete	Side effect, incomplete
0.875	0	0	50	0
0.938	0	0	50	0
0.969	3	0	47	0
0.984	34	0	16	0
0.992	44	0	6	0
0.996	46	0	4	0
0.998	45	0	5	0
0.999	37	1	12	0

Table A.1.: Outcome tallies for Model-free AUP in `Options` (no-op action) testing different γ

λ	Safe, complete	Safe, incomplete	Side effect, complete	Side effect, incomplete
0.370	38	0	12	0
0.416	41	0	9	0
0.476	43	0	7	0
0.555	49	0	1	0
0.666	47	0	3	0
0.833	50	0	0	0
1.110	0	50	0	0
1.664	0	50	0	0
3.322	0	50	0	0
1000.000	0	50	0	0

Table A.2.: Outcome tallies for Model-free AUP in `Options` (no-op action) testing different λ

$\lvert\mathcal{R}\rvert$	Safe, complete	Safe, incomplete	Side effect, complete	Side effect, incomplete
0	0	0	50	0
5	36	0	14	0
10	41	0	9	0
15	41	0	9	0
20	47	0	3	0
25	43	0	7	0
30	47	0	3	0
35	50	0	0	0
40	50	0	0	0
45	50	0	0	0

Table A.3.: Outcome tallies for Model-free AUP in `Options` (no-op action) testing different $\lvert\mathcal{R}\rvert$

γ	Safe, complete	Safe, incomplete	Side effect, complete	Side effect, incomplete
0.875	8	0	42	0
0.938	28	0	22	0
0.969	42	0	8	0
0.984	48	0	2	0
0.992	50	0	0	0
0.996	50	0	0	0
0.998	50	0	0	0
0.999	50	0	0	0

Table A.4.: Outcome tallies for Model-free AUP in `Damage` (no-op action) testing different γ

λ	Safe, complete	Safe, incomplete	Side effect, complete	Side effect, incomplete
0.370	50	0	0	0
0.416	50	0	0	0
0.476	50	0	0	0
0.555	50	0	0	0
0.666	50	0	0	0
0.833	50	0	0	0
1.110	0	50	0	0
1.664	0	50	0	0
3.322	0	50	0	0
1000.000	0	50	0	0

Table A.5.: Outcome tallies for Model-free AUP in `Damage` (no-op action) testing different λ

$\lvert\mathcal{R}\rvert$	Safe, complete	Safe, incomplete	Side effect, complete	Side effect, incomplete
0	0	0	50	0
5	41	0	9	0
10	45	0	5	0
15	50	0	0	0
20	50	0	0	0
25	50	0	0	0
30	50	0	0	0
35	50	0	0	0
40	50	0	0	0
45	50	0	0	0

Table A.6.: Outcome tallies for Model-free AUP in `Damage` (no-op action) testing different $\lvert\mathcal{R}\rvert$

γ	Safe, complete	Safe, incomplete	Side effect, complete	Side effect, incomplete
0.875	0	0	50	0
0.938	0	0	50	0
0.969	0	0	50	0
0.984	0	0	50	0
0.992	0	0	50	0
0.996	0	0	50	0
0.998	0	0	50	0
0.999	0	0	50	0

Table A.7.: Outcome tallies for Model-free AUP in `Correction` (no-op action) testing different γ

λ	Safe, complete	Safe, incomplete	Side effect, complete	Side effect, incomplete
0.370	0	0	50	0
0.416	0	0	50	0
0.476	0	0	50	0
0.555	0	0	50	0
0.666	0	0	50	0
0.833	0	0	50	0
1.110	0	50	0	0
1.664	0	50	0	0
3.322	0	50	0	0
1000.000	0	50	0	0

Table A.8.: Outcome tallies for Model-free AUP in `Correction` (no-op action) testing different λ

$\lvert\mathcal{R}\rvert$	Safe, complete	Safe, incomplete	Side effect, complete	Side effect, incomplete
0	0	0	50	0
5	0	0	50	0
10	0	0	50	0
15	0	0	50	0
20	0	0	50	0
25	0	0	50	0
30	0	0	50	0
35	0	0	50	0
40	0	0	50	0
45	0	0	50	0

Table A.9.: Outcome tallies for Model-free AUP in `Correction` (no-op action) testing different $\lvert\mathcal{R}\rvert$

γ	Safe, complete	Safe, incomplete	Side effect, complete	Side effect, incomplete
0.875	50	0	0	0
0.938	50	0	0	0
0.969	50	0	0	0
0.984	50	0	0	0
0.992	50	0	0	0
0.996	50	0	0	0
0.998	50	0	0	0
0.999	50	0	0	0

Table A.10.: Outcome tallies for Model-free AUP in `Offset` (no-op action) testing different γ

λ	Safe, complete	Safe, incomplete	Side effect, complete	Side effect, incomplete
0.370	50	0	0	0
0.416	50	0	0	0
0.476	50	0	0	0
0.555	50	0	0	0
0.666	50	0	0	0
0.833	50	0	0	0
1.110	50	0	0	0
1.664	40	10	0	0
3.322	30	20	0	0
1000.000	0	50	0	0

Table A.11.: Outcome tallies for Model-free AUP in `Offset` (no-op action) testing different λ

$\lvert\mathcal{R}\rvert$	Safe, complete	Safe, incomplete	Side effect, complete	Side effect, incomplete
0	50	0	0	0
5	50	0	0	0
10	50	0	0	0
15	50	0	0	0
20	50	0	0	0
25	50	0	0	0
30	50	0	0	0
35	50	0	0	0
40	50	0	0	0
45	50	0	0	0

Table A.12.: Outcome tallies for Model-free AUP in `Offset` (no-op action) testing different $\lvert\mathcal{R}\rvert$

γ	Safe, complete	Safe, incomplete	Side effect, complete	Side effect, incomplete
0.875	38	12	0	0
0.938	50	0	0	0
0.969	50	0	0	0
0.984	48	2	0	0
0.992	48	2	0	0
0.996	50	0	0	0
0.998	50	0	0	0
0.999	49	1	0	0

Table A.13.: Outcome tallies for Model-free AUP in `Interference` (no-op action) testing different γ

λ	Safe, complete	Safe, incomplete	Side effect, complete	Side effect, incomplete
0.370	50	0	0	0
0.416	50	0	0	0
0.476	50	0	0	0
0.555	49	1	0	0
0.666	49	1	0	0
0.833	45	5	0	0
1.110	0	50	0	0
1.664	0	50	0	0
3.322	0	50	0	0
1000.000	0	50	0	0

Table A.14.: Outcome tallies for Model-free AUP in `Interference` (no-op action) testing different λ

$\lvert\mathcal{R}\rvert$	Safe, complete	Safe, incomplete	Side effect, complete	Side effect, incomplete
0	50	0	0	0
5	48	2	0	0
10	50	0	0	0
15	49	1	0	0
20	50	0	0	0
25	49	1	0	0
30	49	1	0	0
35	49	1	0	0
40	49	1	0	0
45	49	1	0	0

Table A.15.: Outcome tallies for Model-free AUP in `Interference` (no-op action) testing different $\lvert\mathcal{R}\rvert$

γ	Safe, complete	Safe, incomplete	Side effect, complete	Side effect, incomplete
0.875	0	0	50	0
0.938	0	0	50	0
0.969	0	0	50	0
0.984	5	0	45	0
0.992	16	0	34	0
0.996	26	0	24	0
0.998	7	0	43	0
0.999	0	0	50	0

Table A.16.: Outcome tallies for vAUP (mean) in `Options` (no-op action) testing different γ

λ	Safe, complete	Safe, incomplete	Side effect, complete	Side effect, incomplete
0.370	5	0	45	0
0.416	19	0	31	0
0.476	18	0	32	0
0.555	14	0	36	0
0.666	27	0	23	0
0.833	32	0	18	0
1.110	0	50	0	0
1.664	0	50	0	0
3.322	0	50	0	0
1000.000	0	50	0	0

Table A.17.: Outcome tallies for vAUP (mean) in `Options` (no-op action) testing different λ

$\lvert\mathcal{R}\rvert$	Safe, complete	Safe, incomplete	Side effect, complete	Side effect, incomplete
0	0	0	50	0
5	23	0	27	0
10	20	0	30	0
15	24	0	26	0
20	21	0	29	0
25	15	0	35	0
30	27	0	23	0
35	26	0	24	0
40	27	0	23	0
45	23	0	27	0

Table A.18.: Outcome tallies for vAUP (mean) in `Options` (no-op action) testing different $\lvert\mathcal{R}\rvert$

γ	Safe, complete	Safe, incomplete	Side effect, complete	Side effect, incomplete
0.875	0	0	50	0
0.938	3	0	47	0
0.969	27	0	23	0
0.984	50	0	0	0
0.992	50	0	0	0
0.996	50	0	0	0
0.998	50	0	0	0
0.999	50	0	0	0

Table A.19.: Outcome tallies for vAUP (mean) in `Damage` (no-op action) testing different γ

λ	Safe, complete	Safe, incomplete	Side effect, complete	Side effect, incomplete
0.370	50	0	0	0
0.416	47	0	3	0
0.476	49	0	1	0
0.555	49	0	1	0
0.666	50	0	0	0
0.833	50	0	0	0
1.110	50	0	0	0
1.664	0	0	0	50
3.322	0	0	0	50
1000.000	0	0	0	50

Table A.20.: Outcome tallies for vAUP (mean) in `Damage` (no-op action) testing different λ

$\lvert\mathcal{R}\rvert$	Safe, complete	Safe, incomplete	Side effect, complete	Side effect, incomplete
0	0	0	50	0
5	40	0	10	0
10	46	0	4	0
15	47	0	3	0
20	50	0	0	0
25	50	0	0	0
30	50	0	0	0
35	50	0	0	0
40	50	0	0	0
45	50	0	0	0

Table A.21.: Outcome tallies for vAUP (mean) in `Damage` (no-op action) testing different $\lvert\mathcal{R}\rvert$

γ	Safe, complete	Safe, incomplete	Side effect, complete	Side effect, incomplete
0.875	0	0	50	0
0.938	0	0	50	0
0.969	0	0	50	0
0.984	0	0	50	0
0.992	0	0	50	0
0.996	0	0	50	0
0.998	0	0	50	0
0.999	0	0	50	0

Table A.22.: Outcome tallies for vAUP (mean) in `Correction` (no-op action) testing different γ

λ	Safe, complete	Safe, incomplete	Side effect, complete	Side effect, incomplete
0.370	0	0	50	0
0.416	0	0	50	0
0.476	0	0	50	0
0.555	0	0	50	0
0.666	0	0	50	0
0.833	0	0	50	0
1.110	0	0	49	1
1.664	0	0	0	50
3.322	0	0	0	50
1000.000	0	0	0	50

Table A.23.: Outcome tallies for vAUP (mean) in `Correction` (no-op action) testing different λ

$\lvert\mathcal{R}\rvert$	Safe, complete	Safe, incomplete	Side effect, complete	Side effect, incomplete
0	0	0	50	0
5	0	0	50	0
10	0	0	50	0
15	0	0	50	0
20	0	0	50	0
25	0	0	50	0
30	0	0	50	0
35	0	0	50	0
40	0	0	50	0
45	0	0	50	0

Table A.24.: Outcome tallies for vAUP (mean) in `Correction` (no-op action) testing different $\lvert\mathcal{R}\rvert$

γ	Safe, complete	Safe, incomplete	Side effect, complete	Side effect, incomplete
0.875	50	0	0	0
0.938	50	0	0	0
0.969	50	0	0	0
0.984	50	0	0	0
0.992	50	0	0	0
0.996	50	0	0	0
0.998	49	1	0	0
0.999	42	8	0	0

Table A.25.: Outcome tallies for vAUP (mean) in `Offset` (no-op action) testing different γ

λ	Safe, complete	Safe, incomplete	Side effect, complete	Side effect, incomplete
0.370	50	0	0	0
0.416	50	0	0	0
0.476	50	0	0	0
0.555	50	0	0	0
0.666	50	0	0	0
0.833	50	0	0	0
1.110	43	7	0	0
1.664	26	24	0	0
3.322	13	37	0	0
1000.000	0	50	0	0

Table A.26.: Outcome tallies for vAUP (mean) in `Offset` (no-op action) testing different λ

$\lvert\mathcal{R}\rvert$	Safe, complete	Safe, incomplete	Side effect, complete	Side effect, incomplete
0	50	0	0	0
5	50	0	0	0
10	50	0	0	0
15	50	0	0	0
20	50	0	0	0
25	49	0	1	0
30	50	0	0	0
35	50	0	0	0
40	50	0	0	0
45	50	0	0	0

Table A.27.: Outcome tallies for vAUP (mean) in `Offset` (no-op action) testing different $\lvert\mathcal{R}\rvert$

γ	Safe, complete	Safe, incomplete	Side effect, complete	Side effect, incomplete
0.875	50	0	0	0
0.938	50	0	0	0
0.969	47	2	0	1
0.984	49	1	0	0
0.992	50	0	0	0
0.996	47	3	0	0
0.998	47	3	0	0
0.999	49	1	0	0

Table A.28.: Outcome tallies for vAUP (mean) in `Interference` (no-op action) testing different γ

λ	Safe, complete	Safe, incomplete	Side effect, complete	Side effect, incomplete
0.370	47	2	0	1
0.416	49	1	0	0
0.476	50	0	0	0
0.555	47	3	0	0
0.666	49	1	0	0
0.833	46	4	0	0
1.110	3	45	0	2
1.664	0	45	0	5
3.322	0	43	0	7
1000.000	0	45	0	5

Table A.29.: Outcome tallies for vAUP (mean) in `Interference` (no-op action) testing different λ

$\lvert\mathcal{R}\rvert$	Safe, complete	Safe, incomplete	Side effect, complete	Side effect, incomplete
0	50	0	0	0
5	47	3	0	0
10	47	3	0	0
15	45	5	0	0
20	45	5	0	0
25	47	3	0	0
30	49	1	0	0
35	48	2	0	0
40	48	2	0	0
45	48	2	0	0

Table A.30.: Outcome tallies for vAUP (mean) in `Interference` (no-op action) testing different $\lvert\mathcal{R}\rvert$

γ	Safe, complete	Safe, incomplete	Side effect, complete	Side effect, incomplete
0.875	0	0	50	0
0.938	0	0	50	0
0.969	0	0	50	0
0.984	2	0	48	0
0.992	16	0	34	0
0.996	23	0	27	0
0.998	5	0	45	0
0.999	0	0	50	0

Table A.31.: Outcome tallies for vAUP (oth) in `Options` (no-op action) testing different γ

λ	Safe, complete	Safe, incomplete	Side effect, complete	Side effect, incomplete
0.370	10	0	40	0
0.416	19	0	31	0
0.476	18	0	32	0
0.555	29	0	21	0
0.666	21	0	29	0
0.833	27	0	23	0
1.110	0	50	0	0
1.664	0	50	0	0
3.322	0	50	0	0
1000.000	0	50	0	0

Table A.32.: Outcome tallies for vAUP (oth) in `Options` (no-op action) testing different λ

$\lvert\mathcal{R}\rvert$	Safe, complete	Safe, incomplete	Side effect, complete	Side effect, incomplete
0	0	0	50	0
5	22	0	28	0
10	24	0	26	0
15	29	0	21	0
20	22	0	28	0
25	23	0	27	0
30	21	0	29	0
35	22	0	28	0
40	21	0	29	0
45	26	0	24	0

Table A.33.: Outcome tallies for vAUP (oth) in `Options` (no-op action) testing different $\lvert\mathcal{R}\rvert$

γ	Safe, complete	Safe, incomplete	Side effect, complete	Side effect, incomplete
0.875	0	0	50	0
0.938	8	0	42	0
0.969	35	0	15	0
0.984	48	0	2	0
0.992	50	0	0	0
0.996	50	0	0	0
0.998	50	0	0	0
0.999	50	0	0	0

Table A.34.: Outcome tallies for vAUP (oth) in `Damage` (no-op action) testing different γ

λ	Safe, complete	Safe, incomplete	Side effect, complete	Side effect, incomplete
0.370	48	0	2	0
0.416	48	0	2	0
0.476	50	0	0	0
0.555	49	0	1	0
0.666	50	0	0	0
0.833	50	0	0	0
1.110	50	0	0	0
1.664	0	1	0	49
3.322	0	0	0	50
1000.000	0	1	0	49

Table A.35.: Outcome tallies for vAUP (oth) in `Damage` (no-op action) testing different λ

$\lvert\mathcal{R}\rvert$	Safe, complete	Safe, incomplete	Side effect, complete	Side effect, incomplete
0	0	0	50	0
5	36	0	14	0
10	44	0	6	0
15	48	0	2	0
20	50	0	0	0
25	50	0	0	0
30	50	0	0	0
35	50	0	0	0
40	50	0	0	0
45	50	0	0	0

Table A.36.: Outcome tallies for vAUP (oth) in `Damage` (no-op action) testing different $\lvert\mathcal{R}\rvert$

γ	Safe, complete	Safe, incomplete	Side effect, complete	Side effect, incomplete
0.875	0	0	50	0
0.938	0	0	50	0
0.969	0	0	50	0
0.984	0	0	50	0
0.992	0	0	50	0
0.996	0	0	50	0
0.998	0	0	50	0
0.999	0	0	50	0

Table A.37.: Outcome tallies for vAUP (oth) in `Correction` (no-op action) testing different γ

λ	Safe, complete	Safe, incomplete	Side effect, complete	Side effect, incomplete
0.370	0	0	50	0
0.416	0	0	50	0
0.476	0	0	50	0
0.555	0	0	50	0
0.666	0	0	50	0
0.833	0	0	50	0
1.110	0	0	50	0
1.664	0	0	0	50
3.322	0	0	0	50
1000.000	0	0	0	50

Table A.38.: Outcome tallies for vAUP (oth) in `Correction` (no-op action) testing different λ

$\lvert\mathcal{R}\rvert$	Safe, complete	Safe, incomplete	Side effect, complete	Side effect, incomplete
0	0	0	50	0
5	0	0	50	0
10	0	0	50	0
15	0	0	50	0
20	0	0	50	0
25	0	0	50	0
30	0	0	50	0
35	0	0	50	0
40	0	0	50	0
45	0	0	50	0

Table A.39.: Outcome tallies for vAUP (oth) in `Correction` (no-op action) testing different $\lvert\mathcal{R}\rvert$

γ	Safe, complete	Safe, incomplete	Side effect, complete	Side effect, incomplete
0.875	50	0	0	0
0.938	50	0	0	0
0.969	50	0	0	0
0.984	50	0	0	0
0.992	50	0	0	0
0.996	49	1	0	0
0.998	39	11	0	0
0.999	35	15	0	0

Table A.40.: Outcome tallies for vAUP (oth) in `Offset` (no-op action) testing different γ

λ	Safe, complete	Safe, incomplete	Side effect, complete	Side effect, incomplete
0.370	50	0	0	0
0.416	50	0	0	0
0.476	50	0	0	0
0.555	50	0	0	0
0.666	49	1	0	0
0.833	46	4	0	0
1.110	26	24	0	0
1.664	25	25	0	0
3.322	3	47	0	0
1000.000	0	50	0	0

Table A.41.: Outcome tallies for vAUP (oth) in `Offset` (no-op action) testing different λ

$\lvert\mathcal{R}\rvert$	Safe, complete	Safe, incomplete	Side effect, complete	Side effect, incomplete
0	50	0	0	0
5	49	1	0	0
10	49	0	1	0
15	50	0	0	0
20	50	0	0	0
25	50	0	0	0
30	50	0	0	0
35	50	0	0	0
40	50	0	0	0
45	50	0	0	0

Table A.42.: Outcome tallies for vAUP (oth) in `Offset` (no-op action) testing different $\lvert\mathcal{R}\rvert$

γ	Safe, complete	Safe, incomplete	Side effect, complete	Side effect, incomplete
0.875	50	0	0	0
0.938	50	0	0	0
0.969	50	0	0	0
0.984	49	1	0	0
0.992	48	2	0	0
0.996	47	3	0	0
0.998	48	2	0	0
0.999	49	1	0	0

Table A.43.: Outcome tallies for vAUP (oth) in `Interference` (no-op action) testing different γ

λ	Safe, complete	Safe, incomplete	Side effect, complete	Side effect, incomplete
0.370	49	1	0	0
0.416	46	4	0	0
0.476	47	3	0	0
0.555	49	1	0	0
0.666	45	5	0	0
0.833	48	2	0	0
1.110	13	37	0	0
1.664	0	49	0	1
3.322	0	47	0	3
1000.000	0	44	0	6

Table A.44.: Outcome tallies for vAUP (oth) in `Interference` (no-op action) testing different λ

$\lvert\mathcal{R}\rvert$	Safe, complete	Safe, incomplete	Side effect, complete	Side effect, incomplete
0	50	0	0	0
5	49	1	0	0
10	47	3	0	0
15	46	4	0	0
20	47	3	0	0
25	48	2	0	0
30	45	5	0	0
35	49	1	0	0
40	50	0	0	0
45	50	0	0	0

Table A.45.: Outcome tallies for vAUP (oth) in `Interference` (no-op action) testing different $\lvert\mathcal{R}\rvert$

γ	Safe, complete	Safe, incomplete	Side effect, complete	Side effect, incomplete
0.875	0	0	50	0
0.938	0	0	50	0
0.969	0	0	50	0
0.984	0	0	50	0
0.992	0	0	50	0
0.996	0	0	50	0
0.998	0	0	50	0
0.999	0	0	50	0

Table A.46.: Outcome tallies for vAUP (adv) in `Options` (no-op action) testing different γ

λ	Safe, complete	Safe, incomplete	Side effect, complete	Side effect, incomplete
0.370	0	0	50	0
0.416	1	0	49	0
0.476	0	0	50	0
0.555	0	0	50	0
0.666	0	0	50	0
0.833	0	0	50	0
1.110	0	1	49	0
1.664	0	0	50	0
3.322	0	1	49	0
1000.000	0	0	50	0

Table A.47.: Outcome tallies for vAUP (adv) in `Options` (no-op action) testing different λ

$\lvert\mathcal{R}\rvert$	Safe, complete	Safe, incomplete	Side effect, complete	Side effect, incomplete
0	0	0	50	0
5	13	0	37	0
10	6	0	44	0
15	4	0	46	0
20	1	0	49	0
25	0	0	50	0
30	0	0	50	0
35	1	0	49	0
40	0	0	50	0
45	0	0	50	0

Table A.48.: Outcome tallies for vAUP (adv) in `Options` (no-op action) testing different $\lvert\mathcal{R}\rvert$

γ	Safe, complete	Safe, incomplete	Side effect, complete	Side effect, incomplete
0.875	0	0	50	0
0.938	0	0	50	0
0.969	0	0	50	0
0.984	0	0	50	0
0.992	0	0	50	0
0.996	0	0	50	0
0.998	0	0	50	0
0.999	0	0	50	0

Table A.49.: Outcome tallies for vAUP (adv) in `Damage` (no-op action) testing different γ

λ	Safe, complete	Safe, incomplete	Side effect, complete	Side effect, incomplete
0.370	0	0	50	0
0.416	0	0	50	0
0.476	0	0	50	0
0.555	0	0	50	0
0.666	0	0	50	0
0.833	0	0	50	0
1.110	0	0	50	0
1.664	0	0	50	0
3.322	0	0	50	0
1000.000	0	0	50	0

Table A.50.: Outcome tallies for vAUP (adv) in `Damage` (no-op action) testing different λ

$\lvert\mathcal{R}\rvert$	Safe, complete	Safe, incomplete	Side effect, complete	Side effect, incomplete
0	0	0	50	0
5	1	0	49	0
10	0	0	50	0
15	0	0	50	0
20	0	0	50	0
25	0	0	50	0
30	0	0	50	0
35	0	0	50	0
40	0	0	50	0
45	0	0	50	0

Table A.51.: Outcome tallies for vAUP (adv) in `Damage` (no-op action) testing different $\lvert\mathcal{R}\rvert$

γ	Safe, complete	Safe, incomplete	Side effect, complete	Side effect, incomplete
0.875	0	0	50	0
0.938	0	0	50	0
0.969	0	0	50	0
0.984	0	0	50	0
0.992	0	0	50	0
0.996	0	0	50	0
0.998	0	0	50	0
0.999	0	0	50	0

Table A.52.: Outcome tallies for vAUP (adv) in `Correction` (no-op action) testing different γ

λ	Safe, complete	Safe, incomplete	Side effect, complete	Side effect, incomplete
0.370	0	0	50	0
0.416	0	0	50	0
0.476	0	0	50	0
0.555	0	0	50	0
0.666	0	0	50	0
0.833	0	0	50	0
1.110	0	0	50	0
1.664	0	3	47	0
3.322	0	50	0	0
1000.000	0	50	0	0

Table A.53.: Outcome tallies for vAUP (adv) in `Correction` (no-op action) testing different λ

$\lvert\mathcal{R}\rvert$	Safe, complete	Safe, incomplete	Side effect, complete	Side effect, incomplete
0	0	0	50	0
5	0	0	50	0
10	0	0	50	0
15	0	0	50	0
20	0	0	50	0
25	0	0	50	0
30	0	0	50	0
35	0	0	50	0
40	0	0	50	0
45	0	0	50	0

Table A.54.: Outcome tallies for vAUP (adv) in `Correction` (no-op action) testing different $\lvert\mathcal{R}\rvert$

γ	Safe, complete	Safe, incomplete	Side effect, complete	Side effect, incomplete
0.875	49	0	1	0
0.938	50	0	0	0
0.969	50	0	0	0
0.984	50	0	0	0
0.992	50	0	0	0
0.996	50	0	0	0
0.998	50	0	0	0
0.999	50	0	0	0

Table A.55.: Outcome tallies for vAUP (adv) in `Offset` (no-op action) testing different γ

λ	Safe, complete	Safe, incomplete	Side effect, complete	Side effect, incomplete
0.370	50	0	0	0
0.416	50	0	0	0
0.476	50	0	0	0
0.555	50	0	0	0
0.666	50	0	0	0
0.833	50	0	0	0
1.110	50	0	0	0
1.664	50	0	0	0
3.322	50	0	0	0
1000.000	50	0	0	0

Table A.56.: Outcome tallies for vAUP (adv) in `Offset` (no-op action) testing different λ

$\lvert\mathcal{R}\rvert$	Safe, complete	Safe, incomplete	Side effect, complete	Side effect, incomplete
0	50	0	0	0
5	49	0	1	0
10	50	0	0	0
15	50	0	0	0
20	50	0	0	0
25	50	0	0	0
30	50	0	0	0
35	50	0	0	0
40	50	0	0	0
45	50	0	0	0

Table A.57.: Outcome tallies for vAUP (adv) in `Offset` (no-op action) testing different $\lvert\mathcal{R}\rvert$

γ	Safe, complete	Safe, incomplete	Side effect, complete	Side effect, incomplete
0.875	50	0	0	0
0.938	49	1	0	0
0.969	49	1	0	0
0.984	48	2	0	0
0.992	49	1	0	0
0.996	46	4	0	0
0.998	50	0	0	0
0.999	46	4	0	0

Table A.58.: Outcome tallies for vAUP (adv) in `Interference` (no-op action) testing different γ

λ	Safe, complete	Safe, incomplete	Side effect, complete	Side effect, incomplete
0.370	50	0	0	0
0.416	50	0	0	0
0.476	50	0	0	0
0.555	48	2	0	0
0.666	48	2	0	0
0.833	47	3	0	0
1.110	46	4	0	0
1.664	46	4	0	0
3.322	48	2	0	0
1000.000	47	3	0	0

Table A.59.: Outcome tallies for vAUP (adv) in `Interference` (no-op action) testing different λ

$\lvert\mathcal{R}\rvert$	Safe, complete	Safe, incomplete	Side effect, complete	Side effect, incomplete
0	50	0	0	0
5	46	4	0	0
10	37	13	0	0
15	30	20	0	0
20	43	7	0	0
25	44	6	0	0
30	48	2	0	0
35	49	1	0	0
40	49	1	0	0
45	48	2	0	0

Table A.60.: Outcome tallies for vAUP (adv) in `Interference` (no-op action) testing different $\lvert\mathcal{R}\rvert$

γ	Safe, complete	Safe, incomplete	Side effect, complete	Side effect, incomplete
0.875	0	5	45	0
0.938	0	1	48	1
0.969	0	3	47	0
0.984	0	2	48	0
0.992	1	1	46	2
0.996	1	2	47	0
0.998	0	0	50	0
0.999	0	2	48	0

Table A.61.: Outcome tallies for vAUP (rand) in `Options` (no-op action) testing different γ

λ	Safe, complete	Safe, incomplete	Side effect, complete	Side effect, incomplete
0.370	0	0	50	0
0.416	0	0	50	0
0.476	0	0	50	0
0.555	0	1	49	0
0.666	0	1	49	0
0.833	0	4	45	1
1.110	0	0	45	5
1.664	0	5	0	45
3.322	0	2	0	48
1000.000	0	2	0	48

Table A.62.: Outcome tallies for vAUP (rand) in `Options` (no-op action) testing different λ

$\lvert\mathcal{R}\rvert$	Safe, complete	Safe, incomplete	Side effect, complete	Side effect, incomplete
0	0	0	50	0
5	0	2	47	1
10	0	1	49	0
15	0	1	49	0
20	1	0	48	1
25	0	1	49	0
30	0	2	48	0
35	0	2	47	1
40	0	2	46	2
45	0	0	50	0

Table A.63.: Outcome tallies for vAUP (rand) in `Options` (no-op action) testing different $\lvert\mathcal{R}\rvert$

γ	Safe, complete	Safe, incomplete	Side effect, complete	Side effect, incomplete
0.875	0	0	50	0
0.938	2	0	48	0
0.969	20	0	30	0
0.984	43	0	7	0
0.992	49	0	1	0
0.996	50	0	0	0
0.998	50	0	0	0
0.999	50	0	0	0

Table A.64.: Outcome tallies for vAUP (rand) in `Damage` (no-op action) testing different γ

λ	Safe, complete	Safe, incomplete	Side effect, complete	Side effect, incomplete
0.370	47	0	3	0
0.416	50	0	0	0
0.476	50	0	0	0
0.555	50	0	0	0
0.666	48	0	2	0
0.833	49	0	1	0
1.110	49	0	1	0
1.664	0	20	0	30
3.322	0	12	0	38
1000.000	0	25	0	25

Table A.65.: Outcome tallies for vAUP (rand) in `Damage` (no-op action) testing different λ

$\lvert\mathcal{R}\rvert$	Safe, complete	Safe, incomplete	Side effect, complete	Side effect, incomplete
0	0	0	50	0
5	36	0	14	0
10	46	0	4	0
15	46	0	4	0
20	50	0	0	0
25	48	0	2	0
30	50	0	0	0
35	49	0	1	0
40	50	0	0	0
45	49	0	1	0

Table A.66.: Outcome tallies for vAUP (rand) in `Damage` (no-op action) testing different $\lvert\mathcal{R}\rvert$

γ	Safe, complete	Safe, incomplete	Side effect, complete	Side effect, incomplete
0.875	0	46	4	0
0.938	0	48	2	0
0.969	0	47	3	0
0.984	0	47	3	0
0.992	0	47	3	0
0.996	0	48	2	0
0.998	0	45	5	0
0.999	0	45	5	0

Table A.67.: Outcome tallies for vAUP (rand) in `Correction` (no-op action) testing different γ

λ	Safe, complete	Safe, incomplete	Side effect, complete	Side effect, incomplete
0.370	0	48	2	0
0.416	0	43	7	0
0.476	0	45	5	0
0.555	0	46	4	0
0.666	0	47	3	0
0.833	0	46	4	0
1.110	0	49	1	0
1.664	0	47	0	3
3.322	0	47	0	3
1000.000	0	47	0	3

Table A.68.: Outcome tallies for vAUP (rand) in `Correction` (no-op action) testing different λ

$\lvert\mathcal{R}\rvert$	Safe, complete	Safe, incomplete	Side effect, complete	Side effect, incomplete
0	0	0	50	0
5	0	48	2	0
10	0	46	4	0
15	0	47	3	0
20	0	48	2	0
25	0	41	9	0
30	0	46	4	0
35	0	44	6	0
40	0	43	7	0
45	0	47	3	0

Table A.69.: Outcome tallies for vAUP (rand) in `Correction` (no-op action) testing different $\lvert\mathcal{R}\rvert$

γ	Safe, complete	Safe, incomplete	Side effect, complete	Side effect, incomplete
0.875	50	0	0	0
0.938	50	0	0	0
0.969	50	0	0	0
0.984	49	0	1	0
0.992	48	2	0	0
0.996	8	42	0	0
0.998	0	50	0	0
0.999	0	50	0	0

Table A.70.: Outcome tallies for vAUP (rand) in `Offset` (no-op action) testing different γ

λ	Safe, complete	Safe, incomplete	Side effect, complete	Side effect, incomplete
0.370	50	0	0	0
0.416	50	0	0	0
0.476	47	3	0	0
0.555	30	20	0	0
0.666	8	42	0	0
0.833	0	50	0	0
1.110	0	50	0	0
1.664	1	49	0	0
3.322	0	50	0	0
1000.000	0	50	0	0

Table A.71.: Outcome tallies for vAUP (rand) in `Offset` (no-op action) testing different λ

$\lvert\mathcal{R}\rvert$	Safe, complete	Safe, incomplete	Side effect, complete	Side effect, incomplete
0	50	0	0	0
5	10	40	0	0
10	8	42	0	0
15	10	40	0	0
20	7	43	0	0
25	10	40	0	0
30	5	45	0	0
35	8	42	0	0
40	8	42	0	0
45	6	44	0	0

Table A.72.: Outcome tallies for vAUP (rand) in `Offset` (no-op action) testing different $\lvert\mathcal{R}\rvert$

γ	Safe, complete	Safe, incomplete	Side effect, complete	Side effect, incomplete
0.875	42	8	0	0
0.938	49	1	0	0
0.969	46	4	0	0
0.984	49	1	0	0
0.992	46	4	0	0
0.996	40	9	0	1
0.998	49	0	0	1
0.999	42	8	0	0

Table A.73.: Outcome tallies for vAUP (rand) in `Interference` (no-op action) testing different γ

λ	Safe, complete	Safe, incomplete	Side effect, complete	Side effect, incomplete
0.370	48	2	0	0
0.416	49	1	0	0
0.476	44	5	1	0
0.555	46	3	1	0
0.666	44	5	1	0
0.833	30	20	0	0
1.110	1	47	0	2
1.664	0	50	0	0
3.322	0	50	0	0
1000.000	0	49	0	1

Table A.74.: Outcome tallies for vAUP (rand) in `Interference` (no-op action) testing different λ

$\lvert\mathcal{R}\rvert$	Safe, complete	Safe, incomplete	Side effect, complete	Side effect, incomplete
0	50	0	0	0
5	43	7	0	0
10	39	10	0	1
15	37	13	0	0
20	38	11	1	0
25	46	4	0	0
30	44	6	0	0
35	38	12	0	0
40	45	5	0	0
45	40	9	0	1

Table A.75.: Outcome tallies for vAUP (rand) in `Interference` (no-op action) testing different $\lvert\mathcal{R}\rvert$

A.1.2. Action-driven

γ	Safe, complete	Safe, incomplete	Side effect, complete	Side effect, incomplete
0.875	0	0	50	0
0.938	0	0	50	0
0.969	0	0	50	0
0.984	0	0	50	0
0.992	3	0	47	0
0.996	5	0	45	0
0.998	3	0	47	0
0.999	0	0	50	0

Table A.76.: Outcome tallies for vAUP (mean) in `Options` (action-driven) testing different γ

λ	Safe, complete	Safe, incomplete	Side effect, complete	Side effect, incomplete
0.370	1	0	49	0
0.416	3	0	47	0
0.476	1	0	49	0
0.555	4	0	46	0
0.666	5	0	45	0
0.833	7	0	43	0
1.110	0	50	0	0
1.664	0	50	0	0
3.322	0	50	0	0
1000.000	0	50	0	0

Table A.77.: Outcome tallies for vAUP (mean) in `Options` (action-driven) testing different λ

$\lvert\mathcal{R}\rvert$	Safe, complete	Safe, incomplete	Side effect, complete	Side effect, incomplete
0	0	0	50	0
5	18	0	32	0
10	9	0	41	0
15	3	0	47	0
20	10	0	40	0
25	6	0	44	0
30	5	0	45	0
35	3	0	47	0
40	3	0	47	0
45	2	0	48	0

Table A.78.: Outcome tallies for vAUP (mean) in `Options` (action-driven) testing different $\lvert\mathcal{R}\rvert$

γ	Safe, complete	Safe, incomplete	Side effect, complete	Side effect, incomplete
0.875	0	0	50	0
0.938	0	0	50	0
0.969	29	0	21	0
0.984	46	0	4	0
0.992	49	0	1	0
0.996	50	0	0	0
0.998	50	0	0	0
0.999	50	0	0	0

Table A.79.: Outcome tallies for vAUP (mean) in `Damage` (action-driven) testing different γ

λ	Safe, complete	Safe, incomplete	Side effect, complete	Side effect, incomplete
0.370	47	0	3	0
0.416	49	0	1	0
0.476	50	0	0	0
0.555	49	0	1	0
0.666	50	0	0	0
0.833	50	0	0	0
1.110	49	1	0	0
1.664	0	0	0	50
3.322	0	0	0	50
1000.000	0	0	0	50

Table A.80.: Outcome tallies for vAUP (mean) in `Damage` (action-driven) testing different λ

$\lvert\mathcal{R}\rvert$	Safe, complete	Safe, incomplete	Side effect, complete	Side effect, incomplete
0	0	0	50	0
5	40	0	10	0
10	42	0	8	0
15	46	0	4	0
20	49	0	1	0
25	50	0	0	0
30	50	0	0	0
35	50	0	0	0
40	50	0	0	0
45	50	0	0	0

Table A.81.: Outcome tallies for vAUP (mean) in `Damage` (action-driven) testing different $\lvert\mathcal{R}\rvert$

γ	Safe, complete	Safe, incomplete	Side effect, complete	Side effect, incomplete
0.875	0	0	50	0
0.938	0	0	50	0
0.969	0	0	50	0
0.984	0	0	50	0
0.992	0	0	50	0
0.996	0	0	50	0
0.998	0	0	50	0
0.999	0	0	50	0

Table A.82.: Outcome tallies for vAUP (mean) in `Correction` (action-driven) testing different γ

λ	Safe, complete	Safe, incomplete	Side effect, complete	Side effect, incomplete
0.370	0	0	50	0
0.416	0	0	50	0
0.476	0	0	50	0
0.555	0	0	50	0
0.666	0	0	50	0
0.833	0	0	50	0
1.110	0	0	49	1
1.664	0	0	0	50
3.322	0	0	0	50
1000.000	0	0	0	50

Table A.83.: Outcome tallies for vAUP (mean) in `Correction` (action-driven) testing different λ

$\lvert\mathcal{R}\rvert$	Safe, complete	Safe, incomplete	Side effect, complete	Side effect, incomplete
0	0	0	50	0
5	0	0	50	0
10	0	0	50	0
15	0	0	50	0
20	0	0	50	0
25	0	0	50	0
30	0	0	50	0
35	0	0	50	0
40	0	0	50	0
45	0	0	50	0

Table A.84.: Outcome tallies for vAUP (mean) in `Correction` (action-driven) testing different $\lvert\mathcal{R}\rvert$

γ	Safe, complete	Safe, incomplete	Side effect, complete	Side effect, incomplete
0.875	50	0	0	0
0.938	50	0	0	0
0.969	50	0	0	0
0.984	50	0	0	0
0.992	50	0	0	0
0.996	50	0	0	0
0.998	50	0	0	0
0.999	49	0	1	0

Table A.85.: Outcome tallies for vAUP (mean) in `Offset` (action-driven) testing different γ

λ	Safe, complete	Safe, incomplete	Side effect, complete	Side effect, incomplete
0.370	50	0	0	0
0.416	50	0	0	0
0.476	49	0	1	0
0.555	50	0	0	0
0.666	50	0	0	0
0.833	50	0	0	0
1.110	50	0	0	0
1.664	50	0	0	0
3.322	46	4	0	0
1000.000	0	50	0	0

Table A.86.: Outcome tallies for vAUP (mean) in `Offset` (action-driven) testing different λ

$\lvert\mathcal{R}\rvert$	Safe, complete	Safe, incomplete	Side effect, complete	Side effect, incomplete
0	50	0	0	0
5	47	0	3	0
10	50	0	0	0
15	50	0	0	0
20	50	0	0	0
25	50	0	0	0
30	50	0	0	0
35	50	0	0	0
40	50	0	0	0
45	50	0	0	0

Table A.87.: Outcome tallies for vAUP (mean) in `Offset` (action-driven) testing different $\lvert\mathcal{R}\rvert$

γ	Safe, complete	Safe, incomplete	Side effect, complete	Side effect, incomplete
0.875	50	0	0	0
0.938	49	1	0	0
0.969	49	1	0	0
0.984	49	1	0	0
0.992	50	0	0	0
0.996	49	1	0	0
0.998	50	0	0	0
0.999	48	2	0	0

Table A.88.: Outcome tallies for vAUP (mean) in `Interference` (action-driven) testing different γ

λ	Safe, complete	Safe, incomplete	Side effect, complete	Side effect, incomplete
0.370	50	0	0	0
0.416	50	0	0	0
0.476	50	0	0	0
0.555	49	1	0	0
0.666	48	2	0	0
0.833	48	2	0	0
1.110	21	29	0	0
1.664	0	45	0	5
3.322	0	43	0	7
1000.000	0	47	0	3

Table A.89.: Outcome tallies for vAUP (mean) in `Interference` (action-driven) testing different λ

$\lvert\mathcal{R}\rvert$	Safe, complete	Safe, incomplete	Side effect, complete	Side effect, incomplete
0	50	0	0	0
5	49	1	0	0
10	50	0	0	0
15	50	0	0	0
20	47	3	0	0
25	49	1	0	0
30	48	2	0	0
35	50	0	0	0
40	49	1	0	0
45	50	0	0	0

Table A.90.: Outcome tallies for vAUP (mean) in `Interference` (action-driven) testing different $\lvert\mathcal{R}\rvert$

γ	Safe, complete	Safe, incomplete	Side effect, complete	Side effect, incomplete
0.875	0	0	50	0
0.938	0	0	50	0
0.969	0	0	50	0
0.984	0	0	50	0
0.992	3	0	47	0
0.996	4	0	46	0
0.998	1	0	49	0
0.999	0	0	50	0

Table A.91.: Outcome tallies for vAUP (oth) in `Options` (action-driven) testing different γ

λ	Safe, complete	Safe, incomplete	Side effect, complete	Side effect, incomplete
0.370	1	0	49	0
0.416	3	0	47	0
0.476	2	0	48	0
0.555	2	0	48	0
0.666	4	0	46	0
0.833	6	0	44	0
1.110	0	38	12	0
1.664	0	50	0	0
3.322	0	50	0	0
1000.000	0	49	0	1

Table A.92.: Outcome tallies for vAUP (oth) in `Options` (action-driven) testing different λ

$\lvert\mathcal{R}\rvert$	Safe, complete	Safe, incomplete	Side effect, complete	Side effect, incomplete
0	0	0	50	0
5	14	0	36	0
10	4	0	46	0
15	11	0	39	0
20	9	0	41	0
25	3	0	47	0
30	4	0	46	0
35	3	0	47	0
40	5	0	45	0
45	1	0	49	0

Table A.93.: Outcome tallies for vAUP (oth) in `Options` (action-driven) testing different $\lvert\mathcal{R}\rvert$

γ	Safe, complete	Safe, incomplete	Side effect, complete	Side effect, incomplete
0.875	0	0	50	0
0.938	3	0	47	0
0.969	30	0	20	0
0.984	50	0	0	0
0.992	50	0	0	0
0.996	50	0	0	0
0.998	50	0	0	0
0.999	50	0	0	0

Table A.94.: Outcome tallies for vAUP (oth) in `Damage` (action-driven) testing different γ

λ	Safe, complete	Safe, incomplete	Side effect, complete	Side effect, incomplete
0.370	50	0	0	0
0.416	49	0	1	0
0.476	50	0	0	0
0.555	50	0	0	0
0.666	50	0	0	0
0.833	50	0	0	0
1.110	50	0	0	0
1.664	0	1	0	49
3.322	0	0	0	50
1000.000	0	0	0	50

Table A.95.: Outcome tallies for vAUP (oth) in `Damage` (action-driven) testing different λ

$\lvert\mathcal{R}\rvert$	Safe, complete	Safe, incomplete	Side effect, complete	Side effect, incomplete
0	0	0	50	0
5	38	0	12	0
10	43	0	7	0
15	50	0	0	0
20	50	0	0	0
25	49	0	1	0
30	50	0	0	0
35	50	0	0	0
40	50	0	0	0
45	49	0	1	0

Table A.96.: Outcome tallies for vAUP (oth) in `Damage` (action-driven) testing different $\lvert\mathcal{R}\rvert$

γ	Safe, complete	Safe, incomplete	Side effect, complete	Side effect, incomplete
0.875	0	0	50	0
0.938	0	0	50	0
0.969	0	0	50	0
0.984	0	0	50	0
0.992	0	0	50	0
0.996	0	0	50	0
0.998	0	0	50	0
0.999	0	0	50	0

Table A.97.: Outcome tallies for vAUP (oth) in `Correction` (action-driven) testing different γ

λ	Safe, complete	Safe, incomplete	Side effect, complete	Side effect, incomplete
0.370	0	0	50	0
0.416	0	0	50	0
0.476	0	0	50	0
0.555	0	0	50	0
0.666	0	0	50	0
0.833	0	0	50	0
1.110	0	0	50	0
1.664	0	0	0	50
3.322	0	0	0	50
1000.000	0	0	0	50

Table A.98.: Outcome tallies for vAUP (oth) in `Correction` (action-driven) testing different λ

$\lvert\mathcal{R}\rvert$	Safe, complete	Safe, incomplete	Side effect, complete	Side effect, incomplete
0	0	0	50	0
5	0	0	50	0
10	0	0	50	0
15	0	0	50	0
20	0	0	50	0
25	0	0	50	0
30	0	0	50	0
35	0	0	50	0
40	0	0	50	0
45	0	0	50	0

Table A.99.: Outcome tallies for vAUP (oth) in `Correction` (action-driven) testing different $\lvert\mathcal{R}\rvert$

γ	Safe, complete	Safe, incomplete	Side effect, complete	Side effect, incomplete
0.875	50	0	0	0
0.938	50	0	0	0
0.969	50	0	0	0
0.984	50	0	0	0
0.992	50	0	0	0
0.996	50	0	0	0
0.998	48	2	0	0
0.999	49	1	0	0

Table A.100.: Outcome tallies for vAUP (oth) in `Offset` (action-driven) testing different γ

λ	Safe, complete	Safe, incomplete	Side effect, complete	Side effect, incomplete
0.370	50	0	0	0
0.416	50	0	0	0
0.476	50	0	0	0
0.555	50	0	0	0
0.666	50	0	0	0
0.833	49	1	0	0
1.110	48	2	0	0
1.664	47	3	0	0
3.322	22	28	0	0
1000.000	0	50	0	0

Table A.101.: Outcome tallies for vAUP (oth) in `Offset` (action-driven) testing different λ

$\lvert\mathcal{R}\rvert$	Safe, complete	Safe, incomplete	Side effect, complete	Side effect, incomplete
0	50	0	0	0
5	50	0	0	0
10	50	0	0	0
15	50	0	0	0
20	49	0	1	0
25	50	0	0	0
30	50	0	0	0
35	50	0	0	0
40	50	0	0	0
45	50	0	0	0

Table A.102.: Outcome tallies for vAUP (oth) in `Offset` (action-driven) testing different $\lvert\mathcal{R}\rvert$

γ	Safe, complete	Safe, incomplete	Side effect, complete	Side effect, incomplete
0.875	49	1	0	0
0.938	50	0	0	0
0.969	50	0	0	0
0.984	49	1	0	0
0.992	48	2	0	0
0.996	50	0	0	0
0.998	49	1	0	0
0.999	48	2	0	0

Table A.103.: Outcome tallies for vAUP (oth) in `Interference` (action-driven) testing different γ

λ	Safe, complete	Safe, incomplete	Side effect, complete	Side effect, incomplete
0.370	49	1	0	0
0.416	49	1	0	0
0.476	49	1	0	0
0.555	47	2	0	1
0.666	48	2	0	0
0.833	49	1	0	0
1.110	34	16	0	0
1.664	0	44	0	6
3.322	0	48	0	2
1000.000	0	46	0	4

Table A.104.: Outcome tallies for vAUP (oth) in `Interference` (action-driven) testing different λ

$\lvert\mathcal{R}\rvert$	Safe, complete	Safe, incomplete	Side effect, complete	Side effect, incomplete
0	50	0	0	0
5	49	1	0	0
10	47	3	0	0
15	47	3	0	0
20	48	2	0	0
25	49	1	0	0
30	48	2	0	0
35	50	0	0	0
40	48	1	0	1
45	49	1	0	0

Table A.105.: Outcome tallies for vAUP (oth) in `Interference` (action-driven) testing different $\lvert\mathcal{R}\rvert$

γ	Safe, complete	Safe, incomplete	Side effect, complete	Side effect, incomplete
0.875	0	0	50	0
0.938	0	0	50	0
0.969	0	0	50	0
0.984	0	0	50	0
0.992	0	0	50	0
0.996	0	0	50	0
0.998	0	0	50	0
0.999	0	0	50	0

Table A.106.: Outcome tallies for vAUP (adv) in `Options` (action-driven) testing different γ

λ	Safe, complete	Safe, incomplete	Side effect, complete	Side effect, incomplete
0.370	0	0	50	0
0.416	0	0	50	0
0.476	0	0	50	0
0.555	0	0	50	0
0.666	0	0	50	0
0.833	0	0	50	0
1.110	0	0	50	0
1.664	0	0	50	0
3.322	0	0	50	0
1000.000	0	0	50	0

Table A.107.: Outcome tallies for vAUP (adv) in `Options` (action-driven) testing different λ

$\lvert\mathcal{R}\rvert$	Safe, complete	Safe, incomplete	Side effect, complete	Side effect, incomplete
0	0	0	50	0
5	11	0	39	0
10	11	0	39	0
15	1	0	49	0
20	0	0	50	0
25	1	0	49	0
30	0	0	50	0
35	0	0	50	0
40	1	0	49	0
45	0	0	50	0

Table A.108.: Outcome tallies for vAUP (adv) in `Options` (action-driven) testing different $\lvert\mathcal{R}\rvert$

γ	Safe, complete	Safe, incomplete	Side effect, complete	Side effect, incomplete
0.875	0	0	50	0
0.938	0	0	50	0
0.969	0	0	50	0
0.984	0	0	50	0
0.992	0	0	50	0
0.996	0	0	50	0
0.998	0	0	50	0
0.999	0	0	50	0

Table A.109.: Outcome tallies for vAUP (adv) in `Damage` (action-driven) testing different γ

λ	Safe, complete	Safe, incomplete	Side effect, complete	Side effect, incomplete
0.370	0	0	50	0
0.416	0	0	50	0
0.476	0	0	50	0
0.555	0	0	50	0
0.666	0	0	50	0
0.833	0	0	50	0
1.110	0	0	50	0
1.664	0	0	50	0
3.322	0	0	50	0
1000.000	0	0	50	0

Table A.110.: Outcome tallies for vAUP (adv) in `Damage` (action-driven) testing different λ

$\lvert\mathcal{R}\rvert$	Safe, complete	Safe, incomplete	Side effect, complete	Side effect, incomplete
0	0	0	50	0
5	2	0	48	0
10	0	0	50	0
15	0	0	50	0
20	0	0	50	0
25	0	0	50	0
30	0	0	50	0
35	0	0	50	0
40	0	0	50	0
45	0	0	50	0

Table A.111.: Outcome tallies for vAUP (adv) in `Damage` (action-driven) testing different $\lvert\mathcal{R}\rvert$

γ	Safe, complete	Safe, incomplete	Side effect, complete	Side effect, incomplete
0.875	0	0	50	0
0.938	0	0	50	0
0.969	0	0	50	0
0.984	0	0	50	0
0.992	0	0	50	0
0.996	0	0	50	0
0.998	0	0	50	0
0.999	0	0	50	0

Table A.112.: Outcome tallies for vAUP (adv) in `Correction` (action-driven) testing different γ

λ	Safe, complete	Safe, incomplete	Side effect, complete	Side effect, incomplete
0.370	0	0	50	0
0.416	0	0	50	0
0.476	0	0	50	0
0.555	0	0	50	0
0.666	0	0	50	0
0.833	0	0	50	0
1.110	0	0	50	0
1.664	0	0	50	0
3.322	0	50	0	0
1000.000	0	50	0	0

Table A.113.: Outcome tallies for vAUP (adv) in `Correction` (action-driven) testing different λ

$\lvert\mathcal{R}\rvert$	Safe, complete	Safe, incomplete	Side effect, complete	Side effect, incomplete
0	0	0	50	0
5	0	0	50	0
10	0	0	50	0
15	0	0	50	0
20	0	0	50	0
25	0	0	50	0
30	0	0	50	0
35	0	0	50	0
40	0	0	50	0
45	0	0	50	0

Table A.114.: Outcome tallies for vAUP (adv) in `Correction` (action-driven) testing different $\lvert\mathcal{R}\rvert$

γ	Safe, complete	Safe, incomplete	Side effect, complete	Side effect, incomplete
0.875	50	0	0	0
0.938	50	0	0	0
0.969	50	0	0	0
0.984	50	0	0	0
0.992	50	0	0	0
0.996	50	0	0	0
0.998	50	0	0	0
0.999	50	0	0	0

Table A.115.: Outcome tallies for vAUP (adv) in `Offset` (action-driven) testing different γ

λ	Safe, complete	Safe, incomplete	Side effect, complete	Side effect, incomplete
0.370	50	0	0	0
0.416	50	0	0	0
0.476	50	0	0	0
0.555	50	0	0	0
0.666	50	0	0	0
0.833	50	0	0	0
1.110	50	0	0	0
1.664	50	0	0	0
3.322	50	0	0	0
1000.000	50	0	0	0

Table A.116.: Outcome tallies for vAUP (adv) in `Offset` (action-driven) testing different λ

$\lvert\mathcal{R}\rvert$	Safe, complete	Safe, incomplete	Side effect, complete	Side effect, incomplete
0	50	0	0	0
5	49	0	1	0
10	50	0	0	0
15	50	0	0	0
20	50	0	0	0
25	50	0	0	0
30	50	0	0	0
35	50	0	0	0
40	50	0	0	0
45	50	0	0	0

Table A.117.: Outcome tallies for vAUP (adv) in `Offset` (action-driven) testing different $\lvert\mathcal{R}\rvert$

γ	Safe, complete	Safe, incomplete	Side effect, complete	Side effect, incomplete
0.875	50	0	0	0
0.938	49	1	0	0
0.969	49	1	0	0
0.984	49	1	0	0
0.992	48	2	0	0
0.996	48	2	0	0
0.998	49	1	0	0
0.999	50	0	0	0

Table A.118.: Outcome tallies for vAUP (adv) in `Interference` (action-driven) testing different γ

λ	Safe, complete	Safe, incomplete	Side effect, complete	Side effect, incomplete
0.370	50	0	0	0
0.416	50	0	0	0
0.476	50	0	0	0
0.555	50	0	0	0
0.666	48	2	0	0
0.833	48	2	0	0
1.110	46	4	0	0
1.664	45	5	0	0
3.322	45	5	0	0
1000.000	47	3	0	0

Table A.119.: Outcome tallies for vAUP (adv) in `Interference` (action-driven) testing different λ

$\lvert\mathcal{R}\rvert$	Safe, complete	Safe, incomplete	Side effect, complete	Side effect, incomplete
0	50	0	0	0
5	47	3	0	0
10	45	5	0	0
15	45	5	0	0
20	45	5	0	0
25	46	4	0	0
30	48	2	0	0
35	47	3	0	0
40	49	1	0	0
45	50	0	0	0

Table A.120.: Outcome tallies for vAUP (adv) in `Interference` (action-driven) testing different $\lvert\mathcal{R}\rvert$

γ	Safe, complete	Safe, incomplete	Side effect, complete	Side effect, incomplete
0.875	0	0	50	0
0.938	0	3	47	0
0.969	0	1	49	0
0.984	0	1	49	0
0.992	0	2	47	1
0.996	0	1	49	0
0.998	0	1	49	0
0.999	0	0	50	0

Table A.121.: Outcome tallies for vAUP (rand) in `Options` (action-driven) testing different γ

λ	Safe, complete	Safe, incomplete	Side effect, complete	Side effect, incomplete
0.370	0	0	50	0
0.416	0	0	50	0
0.476	0	0	50	0
0.555	0	1	49	0
0.666	0	1	49	0
0.833	0	4	46	0
1.110	0	0	46	4
1.664	0	2	0	48
3.322	0	5	0	45
1000.000	0	4	0	46

Table A.122.: Outcome tallies for vAUP (rand) in `Options` (action-driven) testing different λ

$\lvert\mathcal{R}\rvert$	Safe, complete	Safe, incomplete	Side effect, complete	Side effect, incomplete
0	0	0	50	0
5	0	1	49	0
10	0	1	48	1
15	0	1	49	0
20	0	0	50	0
25	0	0	50	0
30	0	1	49	0
35	0	0	50	0
40	0	2	48	0
45	0	1	49	0

Table A.123.: Outcome tallies for vAUP (rand) in `Options` (action-driven) testing different $\lvert\mathcal{R}\rvert$

γ	Safe, complete	Safe, incomplete	Side effect, complete	Side effect, incomplete
0.875	0	0	50	0
0.938	2	0	48	0
0.969	29	0	21	0
0.984	44	0	6	0
0.992	49	0	1	0
0.996	49	0	1	0
0.998	50	0	0	0
0.999	50	0	0	0

Table A.124.: Outcome tallies for vAUP (rand) in `Damage` (action-driven) testing different γ

λ	Safe, complete	Safe, incomplete	Side effect, complete	Side effect, incomplete
0.370	49	0	1	0
0.416	48	0	2	0
0.476	47	0	3	0
0.555	50	0	0	0
0.666	49	0	1	0
0.833	45	0	5	0
1.110	49	1	0	0
1.664	0	18	0	32
3.322	0	19	0	31
1000.000	0	20	0	30

Table A.125.: Outcome tallies for vAUP (rand) in `Damage` (action-driven) testing different λ

$\lvert\mathcal{R}\rvert$	Safe, complete	Safe, incomplete	Side effect, complete	Side effect, incomplete
0	0	0	50	0
5	36	0	14	0
10	47	0	3	0
15	48	0	2	0
20	47	0	3	0
25	48	0	2	0
30	50	0	0	0
35	50	0	0	0
40	49	0	1	0
45	49	0	1	0

Table A.126.: Outcome tallies for vAUP (rand) in `Damage` (action-driven) testing different $\lvert\mathcal{R}\rvert$

γ	Safe, complete	Safe, incomplete	Side effect, complete	Side effect, incomplete
0.875	0	43	7	0
0.938	0	44	6	0
0.969	0	45	4	1
0.984	0	49	1	0
0.992	0	48	2	0
0.996	0	46	4	0
0.998	0	45	5	0
0.999	0	46	4	0

Table A.127.: Outcome tallies for vAUP (rand) in `Correction` (action-driven) testing different γ

λ	Safe, complete	Safe, incomplete	Side effect, complete	Side effect, incomplete
0.370	0	46	4	0
0.416	0	49	1	0
0.476	0	47	3	0
0.555	0	47	3	0
0.666	0	45	5	0
0.833	0	47	3	0
1.110	0	48	2	0
1.664	0	47	0	3
3.322	0	48	0	2
1000.000	0	45	0	5

Table A.128.: Outcome tallies for vAUP (rand) in `Correction` (action-driven) testing different λ

$\lvert\mathcal{R}\rvert$	Safe, complete	Safe, incomplete	Side effect, complete	Side effect, incomplete
0	0	0	50	0
5	0	47	3	0
10	0	45	5	0
15	0	47	3	0
20	0	48	2	0
25	0	47	3	0
30	0	43	6	1
35	0	50	0	0
40	0	46	4	0
45	0	43	7	0

Table A.129.: Outcome tallies for vAUP (rand) in `Correction` (action-driven) testing different $\lvert\mathcal{R}\rvert$

γ	Safe, complete	Safe, incomplete	Side effect, complete	Side effect, incomplete
0.875	50	0	0	0
0.938	50	0	0	0
0.969	50	0	0	0
0.984	49	0	1	0
0.992	49	0	1	0
0.996	38	11	1	0
0.998	1	49	0	0
0.999	0	50	0	0

Table A.130.: Outcome tallies for vAUP (rand) in `Offset` (action-driven) testing different γ

λ	Safe, complete	Safe, incomplete	Side effect, complete	Side effect, incomplete
0.370	50	0	0	0
0.416	50	0	0	0
0.476	49	0	1	0
0.555	46	3	1	0
0.666	40	10	0	0
0.833	10	40	0	0
1.110	0	50	0	0
1.664	0	50	0	0
3.322	0	50	0	0
1000.000	0	50	0	0

Table A.131.: Outcome tallies for vAUP (rand) in `Offset` (action-driven) testing different λ

$\lvert\mathcal{R}\rvert$	Safe, complete	Safe, incomplete	Side effect, complete	Side effect, incomplete
0	50	0	0	0
5	38	10	2	0
10	40	10	0	0
15	41	9	0	0
20	36	14	0	0
25	41	9	0	0
30	41	8	1	0
35	39	10	1	0
40	42	8	0	0
45	39	11	0	0

Table A.132.: Outcome tallies for vAUP (rand) in `Offset` (action-driven) testing different $|\mathcal{R}|$

γ	Safe, complete	Safe, incomplete	Side effect, complete	Side effect, incomplete
0.875	50	0	0	0
0.938	50	0	0	0
0.969	49	1	0	0
0.984	47	3	0	0
0.992	44	6	0	0
0.996	44	6	0	0
0.998	48	2	0	0
0.999	48	2	0	0

Table A.133.: Outcome tallies for vAUP (rand) in `Interference` (action-driven) testing different γ

λ	Safe, complete	Safe, incomplete	Side effect, complete	Side effect, incomplete
0.370	50	0	0	0
0.416	50	0	0	0
0.476	49	1	0	0
0.555	50	0	0	0
0.666	45	4	1	0
0.833	34	16	0	0
1.110	27	23	0	0
1.664	0	50	0	0
3.322	0	50	0	0
1000.000	0	50	0	0

Table A.134.: Outcome tallies for vAUP (rand) in `Interference` (action-driven) testing different λ

$\lvert\mathcal{R}\rvert$	Safe, complete	Safe, incomplete	Side effect, complete	Side effect, incomplete
0	50	0	0	0
5	44	6	0	0
10	45	5	0	0
15	49	1	0	0
20	43	7	0	0
25	48	1	0	1
30	42	8	0	0
35	40	10	0	0
40	44	6	0	0
45	47	3	0	0

Table A.135.: Outcome tallies for vAUP (rand) in `Interference` (action-driven) testing different $\lvert\mathcal{R}\rvert$

A.2. Performance - Recorded episodes

A.2.1. No-op action

Environment	Min	Max	Mean	Median	25% quantile	75% quantile
Options	-2.0	0.94	0.2	0.0	-0.06	0.94
Damage	-2.0	1.0	0.07	0.0	0.0	1.0
Correction	-2.0	-0.02	-1.11	-1.0	-1.0	-1.0
Offset	0.0	1.0	0.99	1.0	1.0	1.0
Interference	-2.0	1.0	0.52	1.0	0.22	1.0

Table A.136.: Average Model-free AUP performance in no-op action environments over 50 trials

Environment	Min	Max	Mean	Median	25% quantile	75% quantile
Options	-2.0	0.54	-0.39	-0.24	-0.94	0.08
Damage	-1.42	1.0	0.5	0.84	0.6	1.0
Correction	-2.0	-0.06	-1.01	-1.0	-1.0	-1.0
Offset	-0.04	1.0	0.99	1.0	1.0	1.0
Interference	-2.0	1.0	0.58	1.0	0.82	1.0

Table A.137.: Average vAUP (mean) performance in no-op action environments over 50 trials

Environment	Min	Max	Mean	Median	25% quantile	75% quantile
Options	-2.0	0.44	-0.5	-0.32	-1.0	-0.14
Damage	-1.52	1.0	0.49	0.78	0.54	1.0
Correction	-2.0	0.0	-1.01	-1.0	-1.0	-1.0
Offset	-0.18	1.0	0.97	1.0	0.98	1.0
Interference	-2.0	1.0	0.61	1.0	0.74	1.0

Table A.138.: Average vAUP (oth) performance in no-op action environments over 50 trials

Environment	Min	Max	Mean	Median	25% quantile	75% quantile
Options	-1.82	0.06	-0.99	-1.0	-1.0	-1.0
Damage	-1.36	0.96	-0.98	-1.0	-1.0	-1.0
Correction	-2.0	0.0	-0.34	0.0	-1.0	0.0
Offset	-0.84	1.0	0.98	1.0	1.0	1.0
Interference	-1.56	1.0	0.92	0.96	0.96	0.96

Table A.139.: Average vAUP (adv) performance in no-op action environments over 50 trials

Environment	Min	Max	Mean	Median	25% quantile	75% quantile
Options	-1.98	0.54	-0.41	-0.22	-1.0	-0.04
Damage	-2.0	1.0	0.19	0.02	-0.02	0.98
Correction	-2.0	0.0	-0.55	-0.96	-1.0	0.0
Offset	-0.1	1.0	0.71	1.0	0.16	1.0
Interference	-2.0	1.0	0.5	1.0	0.46	1.0

Table A.140.: Average vAUP (rand) performance in no-op action environments over 50 trials

A.2.2. Action-driven

Environment	Min	Max	Mean	Median	25% quantile	75% quantile
Options	-1.98	0.2	-0.71	-0.72	-0.96	-0.46
Damage	-1.64	1.0	0.41	0.74	-0.04	1.0
Correction	-2.0	-0.12	-1.01	-1.0	-1.0	-1.0
Offset	-0.2	1.0	0.98	1.0	1.0	1.0
Interference	-2.0	1.0	0.58	1.0	0.8	1.0

Table A.141.: Average vAUP (mean) performance in action-driven environments over 50 trials

Environment	Min	Max	Mean	Median	25% quantile	75% quantile
Options	-2.0	0.14	-0.75	-0.7	-0.9	-0.6
Damage	-1.68	1.0	0.39	0.56	0.06	1.0
Correction	-2.0	-0.02	-1.01	-1.0	-1.0	-1.0
Offset	-0.14	1.0	0.99	1.0	1.0	1.0
Interference	-2.0	1.0	0.54	1.0	0.76	1.0

Table A.142.: Average vAUP (oth) performance in action-driven environments over 50 trials

Environment	Min	Max	Mean	Median	25% quantile	75% quantile
Options	-1.94	0.36	-0.99	-1.0	-1.0	-1.0
Damage	-1.54	0.96	-0.98	-1.0	-1.0	-1.0
Correction	-2.0	0.0	-0.35	0.0	-1.0	0.0
Offset	-0.88	1.0	0.91	1.0	1.0	1.0
Interference	-1.4	1.0	0.93	0.96	0.94	0.98

Table A.143.: Average vAUP (adv) performance in action-driven environments over 50 trials

Environment	Min	Max	Mean	Median	25% quantile	75% quantile
Options	-1.98	0.5	-0.43	-0.18	-1.0	-0.06
Damage	-2.0	1.0	0.25	0.06	-0.02	1.0
Correction	-2.0	0.0	-0.53	-0.64	-1.0	0.0
Offset	-0.42	1.0	0.89	1.0	0.8	1.0
Interference	-2.0	1.0	0.46	1.0	0.37	1.0

Table A.144.: Average vAUP (rand) performance in action-driven environments over 50 trials

A.3. Ablation - Agent behavior

A.3.1. vAUP in action-driven environments

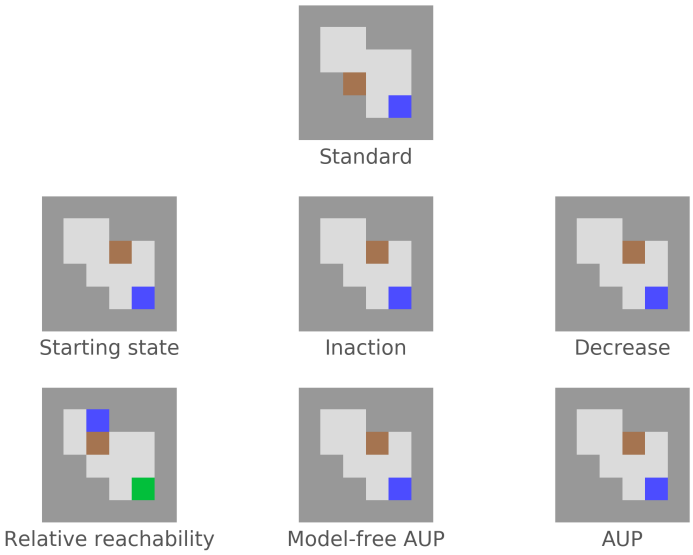

Figure A.1.: Last recorded frame of vAUP in action-driven environments in `Options`

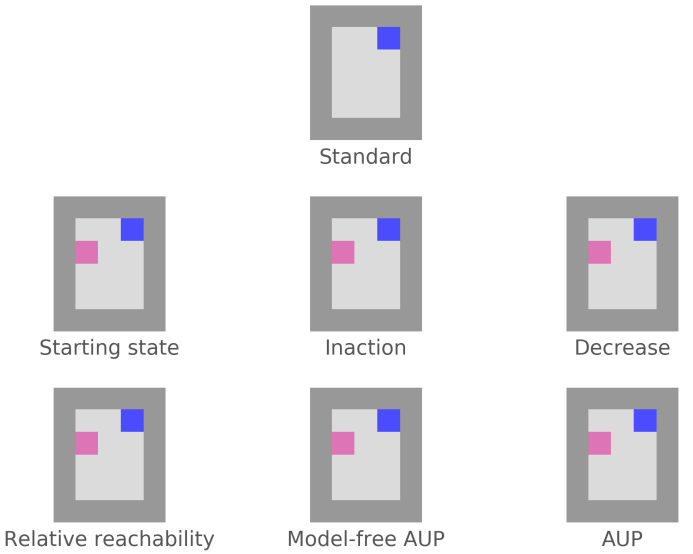

Figure A.2.: Last recorded frame of vAUP in action-driven environments in `Damage`

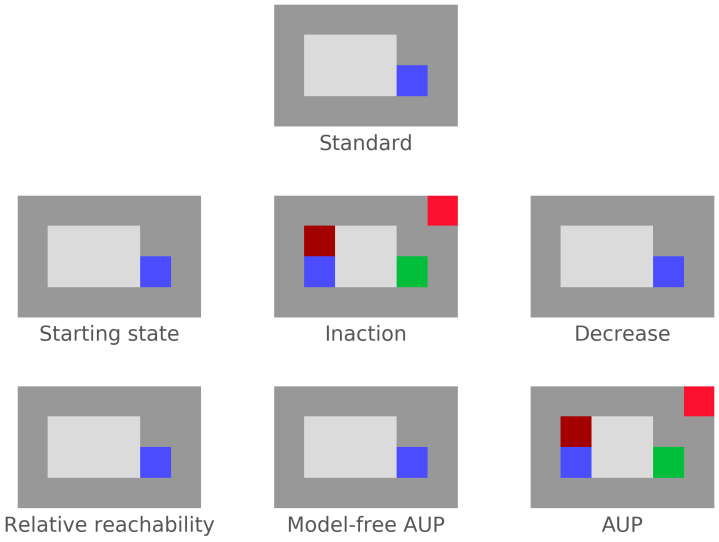

Figure A.3.: Last recorded frame of vAUP in action-driven environments in `Correction`

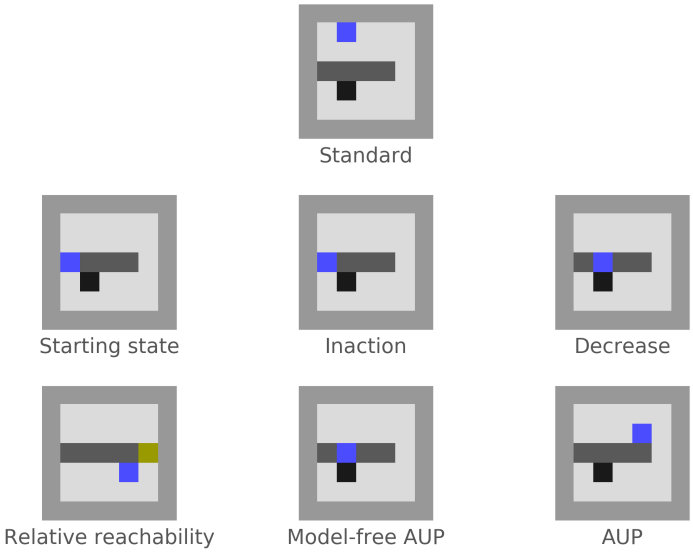

Figure A.4.: Last recorded frame of vAUP in action-driven environments in `Offset`

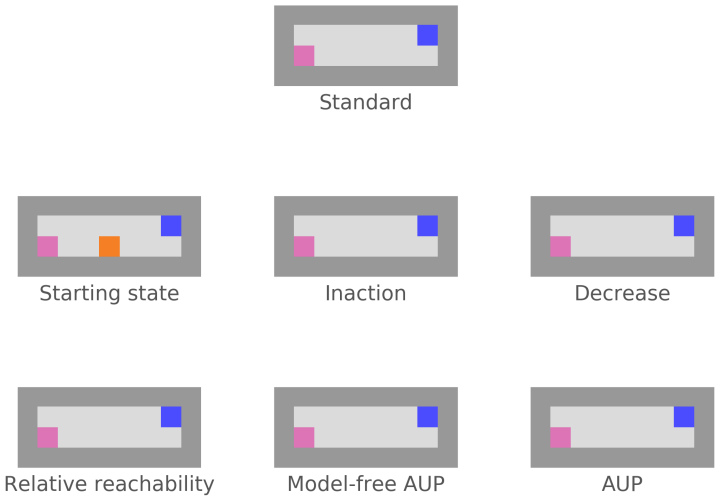

Figure A.5.: Last recorded frame of vAUP in action-driven environments in `Interference`

A.3.2. vAUP in comparison with AUP

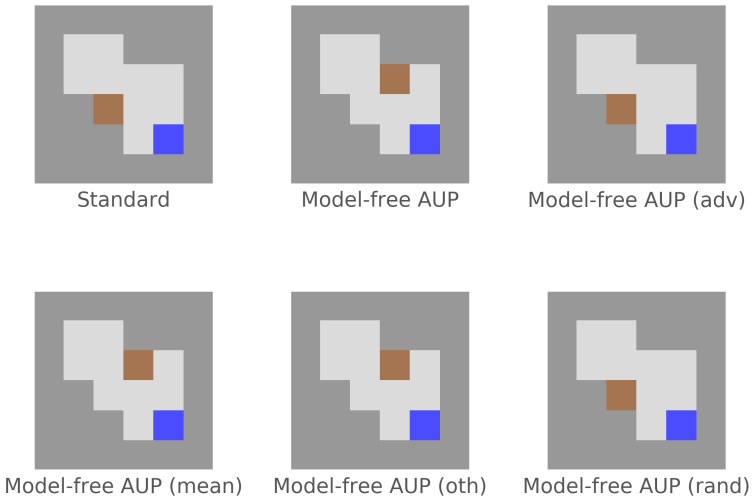

Figure A.6.: Last recorded frame of vAUP in comparison with AUP in `Options`

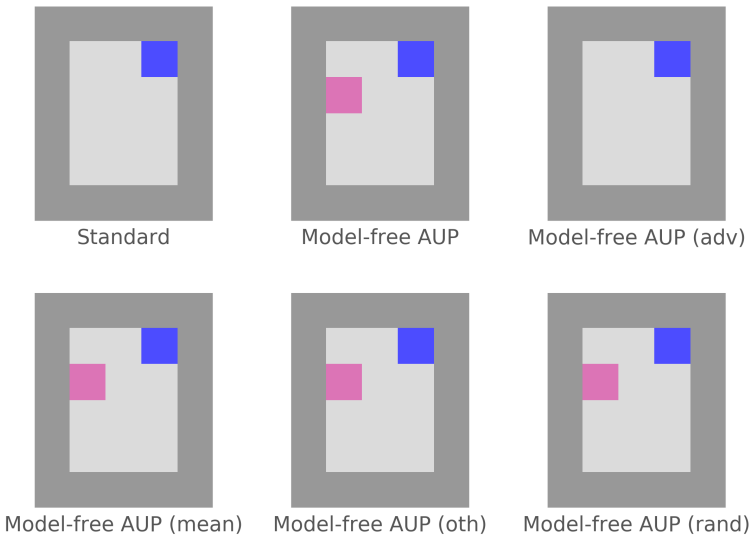

Figure A.7.: Last recorded frame of vAUP in comparison with AUP in `Damage`

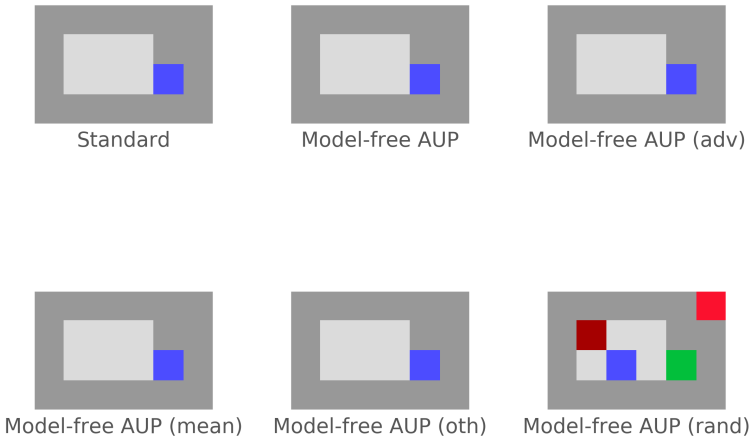

Figure A.8.: Last recorded frame of vAUP in comparison with AUP in `Correction`

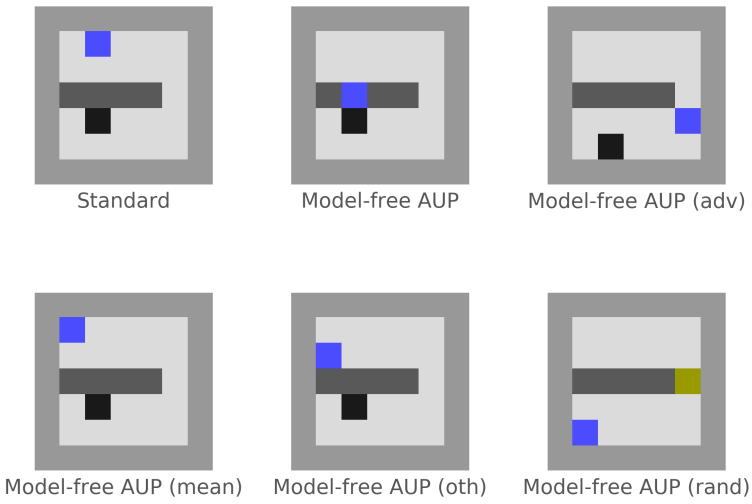

Figure A.9.: Last recorded frame of vAUP in comparison with AUP in `Offset`

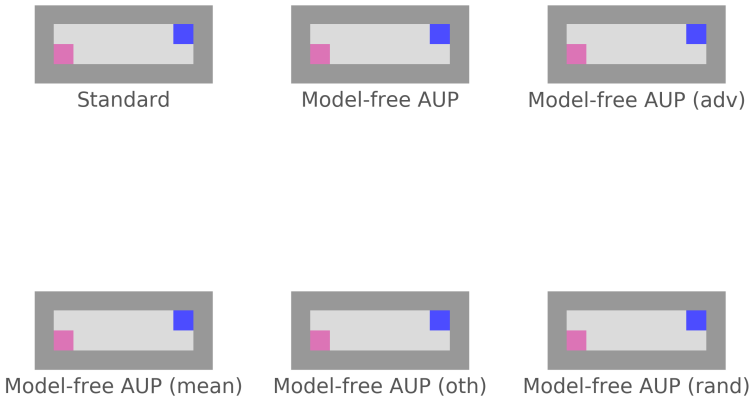

Figure A.10.: Last recorded frame of vAUP in comparison with AUP in `Interference`

A.3.3. vAUP in action-driven environments

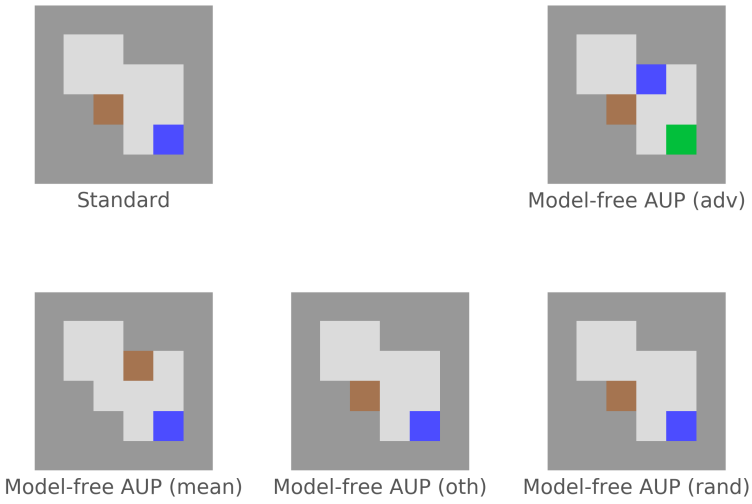

Figure A.11.: Last recorded frame of vAUP in action-driven environments in `Options`

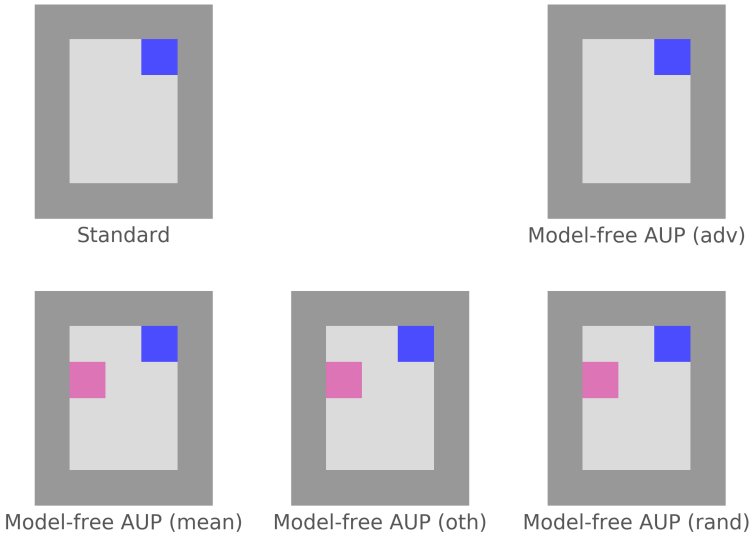

Figure A.12.: Last recorded frame of vAUP in action-driven environments in `Damage`

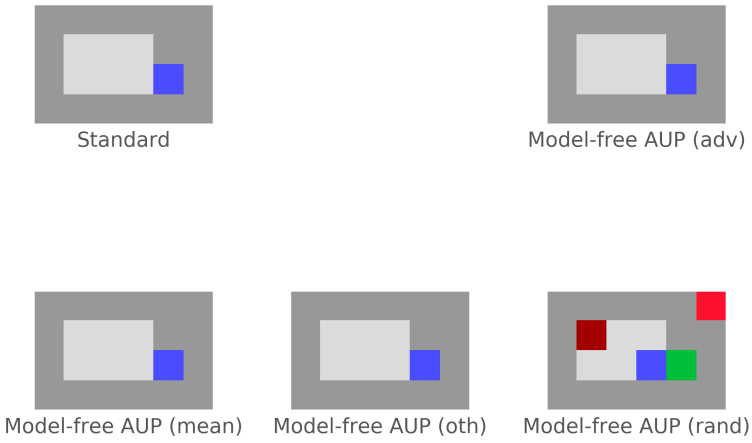

Figure A.13.: Last recorded frame of vAUP in action-driven environments in `Correction`

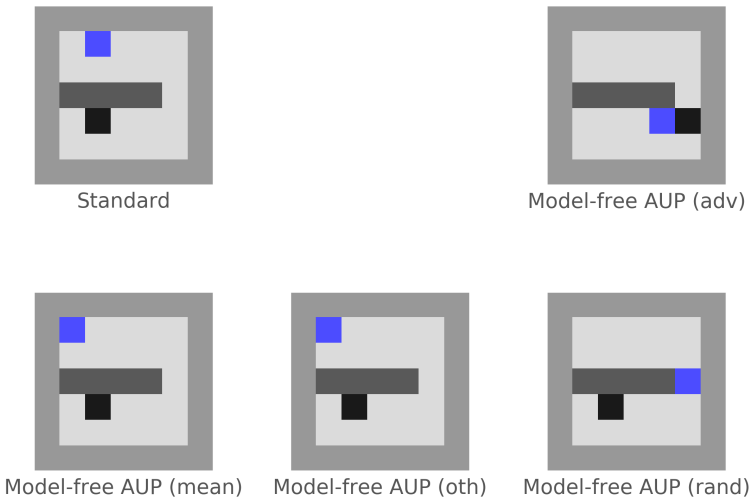

Figure A.14.: Last recorded frame of vAUP in action-driven environments in `Offset`

Figure A.15.: Last recorded frame of vAUP in action-driven environments in `Interference`

List of Figures

2.1 Agent-environment interaction in Reinforcement Learning [SB18] . . . 3
2.2 Generalized policy iteration for estimating $v(s)^*$ and $\pi(a|s)^*$ [SB18] . . . 6
2.3 Different AUP baselines a designer can choose for updating $\text{PENALTY}(s, a)$, each modifying the choice of $Q_{R_i}(s, \varnothing)$ [THT20] . . . 9

4.1 Environments with safety properties of side effects [Lei+17; Lee+18; Kra+19; THT20] . . . 12

5.1 Counts - Model-free AUP outcome tallies across all parameter settings in no-op environments . . . 16
5.2 Performance - Model-free AUP performance over 50 trials in no-op environments . . . 16
5.3 Counts - vAUP (mean) outcome tallies across all parameter settings in no-op environments . . . 18
5.4 Performance - vAUP (mean) performance over 50 trials in no-op environments . . . 18
5.5 Counts - vAUP (oth) outcome tallies across all parameter settings in no-op environments . . . 19
5.6 Performance - vAUP (oth) performance over 50 trials in no-op environments . . . 19
5.7 Counts - vAUP (adv) outcome tallies across all parameter settings in no-op environments . . . 20
5.8 Performance - vAUP (adv) performance over 50 trials in no-op environments . . . 20
5.9 Counts - vAUP (rand) outcome tallies across all parameter settings in no-op environments . . . 21
5.10 Performance - vAUP (rand) performance over 50 trials in no-op environments . . . 21
5.11 Counts - vAUP (mean) outcome tallies across all parameter settings in action-driven environments 23
5.12 Performance - vAUP (mean) performance over 50 trials in action-driven environments . . . 23
5.13 Counts - vAUP (oth) outcome tallies across all parameter settings in action-driven environments . 24
5.14 Performance - vAUP (oth) performance over 50 trials in action-driven environments . . . 24
5.15 Counts - vAUP (adv) outcome tallies across all parameter settings in action-driven environments . 25
5.16 Performance - vAUP (adv) performance over 50 trials in action-driven environments . . . 25
5.17 Counts - vAUP (rand) outcome tallies across all parameter settings in action-driven environments 26
5.18 Performance - vAUP (rand) performance over 50 trials in action-driven environments . . . 26

A.1 Last recorded frame of vAUP in action-driven environments in `Options` . . . 77
A.2 Last recorded frame of vAUP in action-driven environments in `Damage` . . . 77
A.3 Last recorded frame of vAUP in action-driven environments in `Correction` . . . 78
A.4 Last recorded frame of vAUP in action-driven environments in `Offset` . . . 78
A.5 Last recorded frame of vAUP in action-driven environments in `Interference` . . . 79
A.6 Last recorded frame of vAUP in comparison with AUP in `Options` . . . 80
A.7 Last recorded frame of vAUP in comparison with AUP in `Damage` . . . 80
A.8 Last recorded frame of vAUP in comparison with AUP in `Correction` . . . 81
A.9 Last recorded frame of vAUP in comparison with AUP in `Offset` . . . 81
A.10 Last recorded frame of vAUP in comparison with AUP in `Interference` . . . 82
A.11 Last recorded frame of vAUP in action-driven environments in `Options` . . . 83

A.12 Last recorded frame of vAUP in action-driven environments in `Damage` 83
A.13 Last recorded frame of vAUP in action-driven environments in `Correction` 84
A.14 Last recorded frame of vAUP in action-driven environments in `Offset` 84
A.15 Last recorded frame of vAUP in action-driven environments in `Interference` 85

List of Tables

4.1 Standard parameters for AUP and vAUP 13

5.1 Ablation - AUP results including all baseline variants 15
5.2 Ablation - vAUP results compared to AUP including the no-op action ($\varnothing \in \mathcal{A}$) 17
5.3 Ablation - vAUP results in action-driven environments ($\varnothing \notin \mathcal{A}$) 22

A.1 Outcome tallies for Model-free AUP in `Options` (no-op action) testing different γ 30
A.2 Outcome tallies for Model-free AUP in `Options` (no-op action) testing different λ 30
A.3 Outcome tallies for Model-free AUP in `Options` (no-op action) testing different $|\mathcal{R}|$ 30
A.4 Outcome tallies for Model-free AUP in `Damage` (no-op action) testing different γ 31
A.5 Outcome tallies for Model-free AUP in `Damage` (no-op action) testing different λ 31
A.6 Outcome tallies for Model-free AUP in `Damage` (no-op action) testing different $|\mathcal{R}|$ 31
A.7 Outcome tallies for Model-free AUP in `Correction` (no-op action) testing different γ 32
A.8 Outcome tallies for Model-free AUP in `Correction` (no-op action) testing different λ 32
A.9 Outcome tallies for Model-free AUP in `Correction` (no-op action) testing different $|\mathcal{R}|$ 32
A.10 Outcome tallies for Model-free AUP in `Offset` (no-op action) testing different γ 33
A.11 Outcome tallies for Model-free AUP in `Offset` (no-op action) testing different λ 33
A.12 Outcome tallies for Model-free AUP in `Offset` (no-op action) testing different $|\mathcal{R}|$ 33
A.13 Outcome tallies for Model-free AUP in `Interference` (no-op action) testing different γ 34
A.14 Outcome tallies for Model-free AUP in `Interference` (no-op action) testing different λ 34
A.15 Outcome tallies for Model-free AUP in `Interference` (no-op action) testing different $|\mathcal{R}|$ 34
A.16 Outcome tallies for vAUP (mean) in `Options` (no-op action) testing different γ 35
A.17 Outcome tallies for vAUP (mean) in `Options` (no-op action) testing different λ 35
A.18 Outcome tallies for vAUP (mean) in `Options` (no-op action) testing different $|\mathcal{R}|$ 35
A.19 Outcome tallies for vAUP (mean) in `Damage` (no-op action) testing different γ 36
A.20 Outcome tallies for vAUP (mean) in `Damage` (no-op action) testing different λ 36
A.21 Outcome tallies for vAUP (mean) in `Damage` (no-op action) testing different $|\mathcal{R}|$ 36
A.22 Outcome tallies for vAUP (mean) in `Correction` (no-op action) testing different γ 37
A.23 Outcome tallies for vAUP (mean) in `Correction` (no-op action) testing different λ 37
A.24 Outcome tallies for vAUP (mean) in `Correction` (no-op action) testing different $|\mathcal{R}|$ 37
A.25 Outcome tallies for vAUP (mean) in `Offset` (no-op action) testing different γ 38
A.26 Outcome tallies for vAUP (mean) in `Offset` (no-op action) testing different λ 38
A.27 Outcome tallies for vAUP (mean) in `Offset` (no-op action) testing different $|\mathcal{R}|$ 38
A.28 Outcome tallies for vAUP (mean) in `Interference` (no-op action) testing different γ 39
A.29 Outcome tallies for vAUP (mean) in `Interference` (no-op action) testing different λ 39
A.30 Outcome tallies for vAUP (mean) in `Interference` (no-op action) testing different $|\mathcal{R}|$ 39
A.31 Outcome tallies for vAUP (oth) in `Options` (no-op action) testing different γ 40

A.32 Outcome tallies for vAUP (oth) in `Options` (no-op action) testing different λ 40
A.33 Outcome tallies for vAUP (oth) in `Options` (no-op action) testing different $|\mathcal{R}|$ 40
A.34 Outcome tallies for vAUP (oth) in `Damage` (no-op action) testing different γ 41
A.35 Outcome tallies for vAUP (oth) in `Damage` (no-op action) testing different λ 41
A.36 Outcome tallies for vAUP (oth) in `Damage` (no-op action) testing different $|\mathcal{R}|$ 41
A.37 Outcome tallies for vAUP (oth) in `Correction` (no-op action) testing different γ 42
A.38 Outcome tallies for vAUP (oth) in `Correction` (no-op action) testing different λ 42
A.39 Outcome tallies for vAUP (oth) in `Correction` (no-op action) testing different $|\mathcal{R}|$ 42
A.40 Outcome tallies for vAUP (oth) in `Offset` (no-op action) testing different γ 43
A.41 Outcome tallies for vAUP (oth) in `Offset` (no-op action) testing different λ 43
A.42 Outcome tallies for vAUP (oth) in `Offset` (no-op action) testing different $|\mathcal{R}|$ 43
A.43 Outcome tallies for vAUP (oth) in `Interference` (no-op action) testing different γ 44
A.44 Outcome tallies for vAUP (oth) in `Interference` (no-op action) testing different λ 44
A.45 Outcome tallies for vAUP (oth) in `Interference` (no-op action) testing different $|\mathcal{R}|$ 44
A.46 Outcome tallies for vAUP (adv) in `Options` (no-op action) testing different γ 45
A.47 Outcome tallies for vAUP (adv) in `Options` (no-op action) testing different λ 45
A.48 Outcome tallies for vAUP (adv) in `Options` (no-op action) testing different $|\mathcal{R}|$ 45
A.49 Outcome tallies for vAUP (adv) in `Damage` (no-op action) testing different γ 46
A.50 Outcome tallies for vAUP (adv) in `Damage` (no-op action) testing different λ 46
A.51 Outcome tallies for vAUP (adv) in `Damage` (no-op action) testing different $|\mathcal{R}|$ 46
A.52 Outcome tallies for vAUP (adv) in `Correction` (no-op action) testing different γ 47
A.53 Outcome tallies for vAUP (adv) in `Correction` (no-op action) testing different λ 47
A.54 Outcome tallies for vAUP (adv) in `Correction` (no-op action) testing different $|\mathcal{R}|$ 47
A.55 Outcome tallies for vAUP (adv) in `Offset` (no-op action) testing different γ 48
A.56 Outcome tallies for vAUP (adv) in `Offset` (no-op action) testing different λ 48
A.57 Outcome tallies for vAUP (adv) in `Offset` (no-op action) testing different $|\mathcal{R}|$ 48
A.58 Outcome tallies for vAUP (adv) in `Interference` (no-op action) testing different γ 49
A.59 Outcome tallies for vAUP (adv) in `Interference` (no-op action) testing different λ 49
A.60 Outcome tallies for vAUP (adv) in `Interference` (no-op action) testing different $|\mathcal{R}|$ 49
A.61 Outcome tallies for vAUP (rand) in `Options` (no-op action) testing different γ 50
A.62 Outcome tallies for vAUP (rand) in `Options` (no-op action) testing different λ 50
A.63 Outcome tallies for vAUP (rand) in `Options` (no-op action) testing different $|\mathcal{R}|$ 50
A.64 Outcome tallies for vAUP (rand) in `Damage` (no-op action) testing different γ 51
A.65 Outcome tallies for vAUP (rand) in `Damage` (no-op action) testing different λ 51
A.66 Outcome tallies for vAUP (rand) in `Damage` (no-op action) testing different $|\mathcal{R}|$ 51
A.67 Outcome tallies for vAUP (rand) in `Correction` (no-op action) testing different γ 52
A.68 Outcome tallies for vAUP (rand) in `Correction` (no-op action) testing different λ 52
A.69 Outcome tallies for vAUP (rand) in `Correction` (no-op action) testing different $|\mathcal{R}|$ 52
A.70 Outcome tallies for vAUP (rand) in `Offset` (no-op action) testing different γ 53
A.71 Outcome tallies for vAUP (rand) in `Offset` (no-op action) testing different λ 53
A.72 Outcome tallies for vAUP (rand) in `Offset` (no-op action) testing different $|\mathcal{R}|$ 53
A.73 Outcome tallies for vAUP (rand) in `Interference` (no-op action) testing different γ 54
A.74 Outcome tallies for vAUP (rand) in `Interference` (no-op action) testing different λ 54
A.75 Outcome tallies for vAUP (rand) in `Interference` (no-op action) testing different $|\mathcal{R}|$ 54

A.76 Outcome tallies for vAUP (mean) in Options (action-driven) testing different γ 55
A.77 Outcome tallies for vAUP (mean) in Options (action-driven) testing different λ 55
A.78 Outcome tallies for vAUP (mean) in Options (action-driven) testing different $|\mathcal{R}|$ 55
A.79 Outcome tallies for vAUP (mean) in Damage (action-driven) testing different γ 56
A.80 Outcome tallies for vAUP (mean) in Damage (action-driven) testing different λ 56
A.81 Outcome tallies for vAUP (mean) in Damage (action-driven) testing different $|\mathcal{R}|$ 56
A.82 Outcome tallies for vAUP (mean) in Correction (action-driven) testing different γ 57
A.83 Outcome tallies for vAUP (mean) in Correction (action-driven) testing different λ 57
A.84 Outcome tallies for vAUP (mean) in Correction (action-driven) testing different $|\mathcal{R}|$ 57
A.85 Outcome tallies for vAUP (mean) in Offset (action-driven) testing different γ 58
A.86 Outcome tallies for vAUP (mean) in Offset (action-driven) testing different λ 58
A.87 Outcome tallies for vAUP (mean) in Offset (action-driven) testing different $|\mathcal{R}|$ 58
A.88 Outcome tallies for vAUP (mean) in Interference (action-driven) testing different γ 59
A.89 Outcome tallies for vAUP (mean) in Interference (action-driven) testing different λ 59
A.90 Outcome tallies for vAUP (mean) in Interference (action-driven) testing different $|\mathcal{R}|$ 59
A.91 Outcome tallies for vAUP (oth) in Options (action-driven) testing different γ 60
A.92 Outcome tallies for vAUP (oth) in Options (action-driven) testing different λ 60
A.93 Outcome tallies for vAUP (oth) in Options (action-driven) testing different $|\mathcal{R}|$ 60
A.94 Outcome tallies for vAUP (oth) in Damage (action-driven) testing different γ 61
A.95 Outcome tallies for vAUP (oth) in Damage (action-driven) testing different λ 61
A.96 Outcome tallies for vAUP (oth) in Damage (action-driven) testing different $|\mathcal{R}|$ 61
A.97 Outcome tallies for vAUP (oth) in Correction (action-driven) testing different γ 62
A.98 Outcome tallies for vAUP (oth) in Correction (action-driven) testing different λ 62
A.99 Outcome tallies for vAUP (oth) in Correction (action-driven) testing different $|\mathcal{R}|$ 62
A.100 Outcome tallies for vAUP (oth) in Offset (action-driven) testing different γ 63
A.101 Outcome tallies for vAUP (oth) in Offset (action-driven) testing different λ 63
A.102 Outcome tallies for vAUP (oth) in Offset (action-driven) testing different $|\mathcal{R}|$ 63
A.103 Outcome tallies for vAUP (oth) in Interference (action-driven) testing different γ 64
A.104 Outcome tallies for vAUP (oth) in Interference (action-driven) testing different λ 64
A.105 Outcome tallies for vAUP (oth) in Interference (action-driven) testing different $|\mathcal{R}|$ 64
A.106 Outcome tallies for vAUP (adv) in Options (action-driven) testing different γ 65
A.107 Outcome tallies for vAUP (adv) in Options (action-driven) testing different λ 65
A.108 Outcome tallies for vAUP (adv) in Options (action-driven) testing different $|\mathcal{R}|$ 65
A.109 Outcome tallies for vAUP (adv) in Damage (action-driven) testing different γ 66
A.110 Outcome tallies for vAUP (adv) in Damage (action-driven) testing different λ 66
A.111 Outcome tallies for vAUP (adv) in Damage (action-driven) testing different $|\mathcal{R}|$ 66
A.112 Outcome tallies for vAUP (adv) in Correction (action-driven) testing different γ 67
A.113 Outcome tallies for vAUP (adv) in Correction (action-driven) testing different λ 67
A.114 Outcome tallies for vAUP (adv) in Correction (action-driven) testing different $|\mathcal{R}|$ 67
A.115 Outcome tallies for vAUP (adv) in Offset (action-driven) testing different γ 68
A.116 Outcome tallies for vAUP (adv) in Offset (action-driven) testing different λ 68
A.117 Outcome tallies for vAUP (adv) in Offset (action-driven) testing different $|\mathcal{R}|$ 68
A.118 Outcome tallies for vAUP (adv) in Interference (action-driven) testing different γ 69
A.119 Outcome tallies for vAUP (adv) in Interference (action-driven) testing different λ 69

A.120 Outcome tallies for vAUP (adv) in `Interference` (action-driven) testing different $|\mathcal{R}|$ 69
A.121 Outcome tallies for vAUP (rand) in `Options` (action-driven) testing different γ 70
A.122 Outcome tallies for vAUP (rand) in `Options` (action-driven) testing different λ 70
A.123 Outcome tallies for vAUP (rand) in `Options` (action-driven) testing different $|\mathcal{R}|$ 70
A.124 Outcome tallies for vAUP (rand) in `Damage` (action-driven) testing different γ 71
A.125 Outcome tallies for vAUP (rand) in `Damage` (action-driven) testing different λ 71
A.126 Outcome tallies for vAUP (rand) in `Damage` (action-driven) testing different $|\mathcal{R}|$ 71
A.127 Outcome tallies for vAUP (rand) in `Correction` (action-driven) testing different γ 72
A.128 Outcome tallies for vAUP (rand) in `Correction` (action-driven) testing different λ 72
A.129 Outcome tallies for vAUP (rand) in `Correction` (action-driven) testing different $|\mathcal{R}|$ 72
A.130 Outcome tallies for vAUP (rand) in `Offset` (action-driven) testing different γ 73
A.131 Outcome tallies for vAUP (rand) in `Offset` (action-driven) testing different λ 73
A.132 Outcome tallies for vAUP (rand) in `Offset` (action-driven) testing different $|\mathcal{R}|$ 73
A.133 Outcome tallies for vAUP (rand) in `Interference` (action-driven) testing different γ 74
A.134 Outcome tallies for vAUP (rand) in `Interference` (action-driven) testing different λ 74
A.135 Outcome tallies for vAUP (rand) in `Interference` (action-driven) testing different $|\mathcal{R}|$ 74
A.136 Average Model-free AUP performance in no-op action environments over 50 trials 75
A.137 Average vAUP (mean) performance in no-op action environments over 50 trials 75
A.138 Average vAUP (oth) performance in no-op action environments over 50 trials 75
A.139 Average vAUP (adv) performance in no-op action environments over 50 trials 75
A.140 Average vAUP (rand) performance in no-op action environments over 50 trials 75
A.141 Average vAUP (mean) performance in action-driven environments over 50 trials 76
A.142 Average vAUP (oth) performance in action-driven environments over 50 trials 76
A.143 Average vAUP (adv) performance in action-driven environments over 50 trials 76
A.144 Average vAUP (rand) performance in action-driven environments over 50 trials 76

Glossary

action-value function	The action-value function represents the expected return that can be obtained from a state s taking action a under a policy π, denoted $q_\pi(s, a)$, when starting in state s, taking action a and following policy π thereafter.
agent	An agent in the context of reinforcement learning is an autonomously acting system.
AI	Artificial Intelligence
artificial general intelligence	Artificial Intelligence, that is able to learn and understand any intellectual task that a human being can.
AUP	Attainable Utility Preservation
episode	An episode is considered all cycles of interaction between an agent and environment until reaching the terminal state T.
MDP	Markov Decision Process
no-op action	A no-op action in an environment is an action, where the agent does nothing during a timestep (e.g., not moving or not doing any other task). The existence of a no-op action $\varnothing$ depends on the environment and if available, is typically explicitely assigned to the agent's action space A with $\varnothing \in A$.
return	The return G_t refers to the sum of total discounted rewards starting from timestep t.
RL	Reinforcement Learning
state-value function	The state-value function represents the expected return that can be obtained from a state s under a policy π, denoted $v_\pi(s)$, when starting in state s and following policy π thereafter.
vAUP	Variations of Attainable Utility Preservation

References

[AL17] Stuart Armstrong and Benjamin Levinstein. *Low Impact Artificial Intelligences*. May 30, 2017. DOI: 10.48550/arXiv.1705.10720. arXiv: 1705.10720 [cs].

[Ala+21] Parand Alizadeh Alamdari, Toryn Q. Klassen, Rodrigo Toro Icarte, and Sheila A. McIlraith. "Be Considerate: Objectives, Side Effects, and Deciding How to Act." June 4, 2021. arXiv: 2106.02617 [cs].

[Alt99] Eitan Altman. *Constrained Markov Decision Processes*. Stochastic Modeling. Boca Raton ; London: Chapman & Hall/CRC, 1999. 242 pp. ISBN: 978-0-8493-0382-1.

[Amo+16] Dario Amodei, Chris Olah, Jacob Steinhardt, Paul Christiano, John Schulman, and Dan Mané. "Concrete Problems in AI Safety." July 25, 2016. arXiv: 1606.06565 [cs].

[Ant+21] Ioannis Antonoglou, Julian Schrittwieser, Sherjil Ozair, Thomas K. Hubert, and David Silver. "Planning in Stochastic Environments with a Learned Model." In: *10th International Conference on Learning Representations*. International Conference on Learning Representations. Sept. 29, 2021.

[BDM17] Marc G. Bellemare, Will Dabney, and Rémi Munos. "A Distributional Perspective on Reinforcement Learning." In: *Proceedings of the 34th International Conference on Machine Learning*. International Conference on Machine Learning. PMLR, July 17, 2017, pp. 449–458.

[Bel+13] Marc G. Bellemare, Yavar Naddaf, Joel Veness, and Michael Bowling. "The Arcade Learning Environment: An Evaluation Platform for General Agents." In: *Journal of Artificial Intelligence Research* 47 (2013), pp. 253–279. DOI: 10.1613/jair.3912.

[Bel57] RICHARD Bellman. "A Markovian Decision Process." In: *Journal of Mathematics and Mechanics* 6.5 (1957), pp. 679–684. ISSN: 0095-9057. JSTOR: 24900506.

[Ber+17] Felix Berkenkamp, Matteo Turchetta, Angela Schoellig, and Andreas Krause. "Safe Model-Based Reinforcement Learning with Stability Guarantees." In: *Advances in Neural Information Processing Systems*. Conference in Neural Information Processing Systems. Ed. by I. Guyon, U. V. Luxburg, S. Bengio, H. Wallach, R. Fergus, S. Vishwanathan, and R. Garnett. Vol. 30. Curran Associates, Inc., 2017.

[Biy+19] Erdem Biyik, Jonathan Margoliash, Shahrouz Ryan Alimo, and Dorsa Sadigh. "Efficient and Safe Exploration in Deterministic Markov Decision Processes with Unknown Transition Models." In: *2019 American Control Conference (ACC)*. 2019 American Control Conference (ACC). July 2019, pp. 1792–1799. DOI: 10.23919/ACC.2019.8815276.

[Bos14] Nick Bostrom. *Superintelligence: Paths, Dangers, Strategies*. First edition. Oxford: Oxford University Press, 2014. 328 pp. ISBN: 978-0-19-967811-2.

[Buc21] Alexander Buchelt. "Reprocuciblility of 'Avoiding Side Effects in Complex Environments'." BA thesis. St. Pölten, Austria: St. Pölten University of Applied Sciences, 2021. 36 pp.

[Cho+18] Yinlam Chow, Ofir Nachum, Edgar Duenez-Guzman, and Mohammad Ghavamzadeh. “A Lyapunov-based Approach to Safe Reinforcement Learning.” In: *Advances in Neural Information Processing Systems*. Vol. 31. Curran Associates, Inc., 2018.

[Chr+17] Paul F Christiano, Jan Leike, Tom Brown, Miljan Martic, Shane Legg, and Dario Amodei. “Deep Reinforcement Learning from Human Preferences.” In: *Advances in Neural Information Processing Systems*. Vol. 30. Curran Associates, Inc., 2017.

[Chr21] Brian Christian. *The Alignment Problem: Machine Learning and Human Values*. 2021. ISBN: 978-0-393-86833-3.

[EL21] Adrien Ecoffet and Joel Lehman. “Reinforcement Learning Under Moral Uncertainty.” In: *Proceedings of the 38th International Conference on Machine Learning*. International Conference on Machine Learning. PMLR, July 1, 2021, pp. 2926–2936.

[Eys+18] Benjamin Eysenbach, Shixiang Gu, Julian Ibarz, and Sergey Levine. “Leave No Trace: Learning to Reset for Safe and Autonomous Reinforcement Learning.” In: International Conference on Learning Representations. Feb. 15, 2018.

[For+18] Meire Fortunato, Mohammad Gheshlaghi Azar, Bilal Piot, Jacob Menick, Matteo Hessel, Ian Osband, Alex Graves, Volodymyr Mnih, Rémi Munos, Demis Hassabis, Olivier Pietquin, Charles Blundell, and Shane Legg. “Noisy Networks for Exploration.” In: *6th International Conference on Learning Representations*. Internation Conference on Learning Representations. Vancouver, BC, Canada: OpenReview.net, Feb. 15, 2018.

[Gab20] Iason Gabriel. “Artificial Intelligence, Values, and Alignment.” In: *Minds and Machines* 30.3 (Sept. 1, 2020), pp. 411–437. ISSN: 1572-8641. DOI: 10.1007/s11023-020-09539-2.

[GF15] Javier García and Fernando Fernández. “A Comprehensive Survey on Safe Reinforcement Learning.” In: *The Journal of Machine Learning Research* 16.1 (Jan. 1, 2015), pp. 1437–1480. ISSN: 1532-4435.

[Hen+21] Dan Hendrycks, Nicholas Carlini, John Schulman, and Jacob Steinhardt. “Unsolved Problems in ML Safety.” Oct. 30, 2021. arXiv: 2109.13916 [cs].

[How70] Ronald A. Howard. *Dynamic Programming and Markov Processes*. 6. print. Cambridge, Mass: M.I.T. Pr, 1970. 136 pp. ISBN: 978-0-262-08009-5.

[HTR14] Stephen Hawking, Max Tegmark, and Stuart Russel. *Transcending Complacency On Superintelligent Machines*. HuffPost. Apr. 19, 2014. URL: https://www.huffpost.com/entry/artificial-intelligence_b_5174265 (visited on 08/08/2022).

[Kra+19] Victoria Krakovna, Laurent Orseau, Miljan Martic, and Shane Legg. “Penalizing Side Effects Using Stepwise Relative Reachability.” In: *Proceedings of the Workshop on Artificial Intelligence Safety 2019*. Artificial Intelligence Safety 2019. Ed. by Huáscar Espinoza, Han Yu, Xiaowei Huang, Freddy Lecue, Cynthia Chen, José Hernández-Orallo, Seán Ó hÉigeartaigh, and Richard Mallah. Vol. 2419. CEUR Workshop Proceedings. Macao, China: CEUR, Aug. 11, 2019.

[Kra+20] Victoria Krakovna, Laurent Orseau, Richard Ngo, Miljan Martic, and Shane Legg. “Avoiding Side Effects By Considering Future Tasks.” In: *Advances in Neural Information Processing Systems*. Conference on Neural Information Processing Systems. Ed. by H. Larochelle, M. Ranzato, R. Hadsell, M. F. Balcan, and H. Lin. Vol. 33. Curran Associates, Inc., 2020, pp. 19064–19074.

[Lee+18] Gavin Leech, Karol Kubicki, Jessica Cooper, and Tom McGrath. “Preventing Side-effects in Gridworlds.” unpublished. AI Safety Camp, Gran Canaria, Apr. 22, 2018.

[Lei+17] Jan Leike, Miljan Martic, Victoria Krakovna, Pedro A. Ortega, Tom Everitt, Andrew Lefrancq, Laurent Orseau, and Shane Legg. *AI Safety Gridworlds*. Nov. 28, 2017. DOI: 10.48550/arXiv.1711.09883. arXiv: 1711.09883 [cs].

[LH07] Shane Legg and Marcus Hutter. *A Collection of Definitions of Intelligence*. June 25, 2007. DOI: 10.48550/arXiv.0706.3639. arXiv: 0706.3639 [cs].

[LMM21] David Lindner, Kyle Matoba, and Alexander Meulemans. “Challenges for Using Impact Regularizers to Avoid Negative Side Effects.” In: *Proceedings of the Workshop on Artificial Intelligence Safety 2021 (SafeAI 2021) Co-Located with the Thirty-Fifth AAAI Conference on Artificial Intelligence*. Artificial Intelligence Safety 2021. Ed. by Huáscar Espinoza, John McDermid, Xiaowei Huang, Mauricio Castillo-Effen, Xin Cynthia Chen, José Hernández-Orallo, Seán Ó hÉigeartaigh, and Richard Mallah. Vol. 2808. CEUR Workshop Proceedings. Virtual, February: CEUR, Feb. 8, 2021.

[LSA21] Kimin Lee, Laura M. Smith, and Pieter Abbeel. “PEBBLE: Feedback-Efficient Interactive Reinforcement Learning via Relabeling Experience and Unsupervised Pre-training.” In: *Proceedings of the 38th International Conference on Machine Learning*. International Conference on Machine Learning. PMLR, July 1, 2021, pp. 6152–6163.

[MA26] Teodor Mihai Moldovan and Pieter Abbeel. “Safe Exploration in Markov Decision Processes.” In: *Proceedings of the 29th International Conference on Machine Learning*. Internation Conference on Machine Learning. Edinburgh, Scotland, UK: icml.cc / Omnipress, June 26–July 1, 2012.

[MJ15] Shakir Mohamed and Danilo Jimenez Rezende. “Variational Information Maximisation for Intrinsically Motivated Reinforcement Learning.” In: *Advances in Neural Information Processing Systems*. Vol. 28. Curran Associates, Inc., 2015.

[MMW20] Santiago Miret, Somdeb Majumdar, and Carroll Wainwright. “Safety Aware Reinforcement Learning (SARL).” Oct. 6, 2020. arXiv: 2010.02846 [cs].

[Mni+13] Volodymyr Mnih, Koray Kavukcuoglu, David Silver, Alex Graves, Ioannis Antonoglou, Daan Wierstra, and Martin Riedmiller. “Playing Atari with Deep Reinforcement Learning.” In: *CoRR* abs/1312.5602 (Dec. 19, 2013).

[Omo08] Stephen M. Omohundro. “The Basic AI Drives.” In: *Artificial General Intelligence 2008* (2008), pp. 483–492.

[Ope+19] OpenAI, Christopher Berner, Greg Brockman, Brooke Chan, Vicki Cheung, Przemysław Dębiak, Christy Dennison, David Farhi, Quirin Fischer, Shariq Hashme, Chris Hesse, Rafal Józefowicz, Scott Gray, Catherine Olsson, Jakub Pachocki, Michael Petrov, Henrique P. d O. Pinto, Jonathan Raiman, Tim Salimans, Jeremy Schlatter, Jonas Schneider, Szymon Sidor, Ilya Sutskever, Jie Tang, Filip Wolski, and Susan Zhang. *Dota 2 with Large Scale Deep Reinforcement Learning*. Dec. 13, 2019. DOI: 10.48550/arXiv.1912.06680. arXiv: 1912.06680 [cs, stat].

[PS14] Martin Pecka and Tomas Svoboda. "Safe Exploration Techniques for Reinforcement Learning – An Overview." In: *Modelling and Simulation for Autonomous Systems*. Ed. by Jan Hodicky. Vol. 8906. Lecture Notes in Computer Science. Cham: Springer International Publishing, May 5, 2014, pp. 357–375. ISBN: 978-3-319-13823-7. DOI: 10.1007/978-3-319-13823-7_31.

[RB10] Kevin Regan and Craig Boutilier. "Robust Policy Computation in Reward-Uncertain MDPs Using Nondominated Policies." In: *Proceedings of the Twenty-Fourth AAAI Conference on Artificial Intelligence, AAAI 2010*. Conference on Artificial Intelligence. Ed. by Maria Fox and David Poole. Georgia, USA: AAAI Press, July 11–15, 2010.

[RN21] Stuart J. Russell and Peter Norvig. *Artificial Intelligence: A Modern Approach*. Fourth edition. Pearson Series in Artificial Intelligence. Hoboken: Pearson, 2021. ISBN: 978-0-13-461099-3.

[Rus16] Stuart Russell. "Should We Fear Supersmart Robots?" In: *Scientific American* 314.6 (May 17, 2016), pp. 58–59. ISSN: 0036-8733. DOI: 10.1038/scientificamerican0616-58.

[Rus19] Stuart J. Russell. *Human Compatible: Artificial Intelligence and the Problem of Control*. New York?: Viking, 2019. 336 pp. ISBN: 978-0-525-55861-3.

[Sau+18] William Saunders, Girish Sastry, Andreas Stuhlmüller, and Owain Evans. "Trial without Error: Towards Safe Reinforcement Learning via Human Intervention." In: *Proceedings of the 17th International Conference on Autonomous Agents and MultiAgent Systems*. AAMAS '18. Richland, SC: International Foundation for Autonomous Agents and Multiagent Systems, July 9, 2018, pp. 2067–2069.

[SB18] Richard S. Sutton and Andrew G. Barto. *Reinforcement Learning: An Introduction*. Second edition. Adaptive Computation and Machine Learning Series. Cambridge, Massachusetts: The MIT Press, 2018. 526 pp. ISBN: 978-0-262-03924-6.

[Sch+16] Tom Schaul, John Quan, Ioannis Antonoglou, and David Silver. "Prioritized Experience Replay." In: *4th International Conference on Learning Representations*. International Conference on Learning Representations. Ed. by Yoshua Bengio and Yann LeCun. San Juan, Puerto Rico, May 2–4, 2016.

[Sch+20] Julian Schrittwieser, Ioannis Antonoglou, Thomas Hubert, Karen Simonyan, Laurent Sifre, Simon Schmitt, Arthur Guez, Edward Lockhart, Demis Hassabis, Thore Graepel, Timothy Lillicrap, and David Silver. "Mastering Atari, Go, Chess and Shogi by Planning with a Learned Model." In: *Nature* 588.7839 (7839 Dec. 2020), pp. 604–609. ISSN: 1476-4687. DOI: 10.1038/s41586-020-03051-4.

[Sil+16] David Silver, Aja Huang, Chris J. Maddison, Arthur Guez, Laurent Sifre, George van den Driessche, Julian Schrittwieser, Ioannis Antonoglou, Veda Panneershelvam, Marc Lanctot, Sander Dieleman, Dominik Grewe, John Nham, Nal Kalchbrenner, Ilya Sutskever, Timothy Lillicrap, Madeleine Leach, Koray Kavukcuoglu, Thore Graepel, and Demis Hassabis. "Mastering the Game of Go with Deep Neural Networks and Tree Search." In: *Nature* 529.7587 (7587 Jan. 28, 2016), pp. 484–489. ISSN: 0028-0836, 1476-4687. DOI: 10.1038/nature16961.

[Sut88] Richard S. Sutton. "Learning to Predict by the Methods of Temporal Differences." In: *Machine Learning* 3.1 (1 Aug. 1988), pp. 9–44. ISSN: 0885-6125, 1573-0565. DOI: 10.1007/bf00115009.

[SVH19] Henry Shevlin, Karina Vold, and Halina. "The Limits of Machine Intelligence." In: *EMBO reports* 20.10 (Oct. 4, 2019), e49177. ISSN: 1469-221X. DOI: 10.15252/embr.201949177.

[SZK21] Sandhya Saisubramanian, Shlomo Zilberstein, and Ece Kamar. "Avoiding Negative Side Effects Due to Incomplete Knowledge of AI Systems." Oct. 18, 2021. arXiv: 2008.12146 [cs].

[TD21] Mariya Tsvarkaleva and Louise A. Dennis. "No Free Lunch: Overcoming Reward Gaming in AI Safety Gridworlds." In: *Computer Safety, Reliability, and Security. SAFECOMP 2021 Workshops*. Ed. by Ibrahim Habli, Mark Sujan, Simos Gerasimou, Erwin Schoitsch, and Friedemann Bitsch. Lecture Notes in Computer Science. Cham: Springer International Publishing, 2021, pp. 226–238. ISBN: 978-3-030-83906-2. DOI: 10.1007/978-3-030-83906-2_18.

[THT20] Alexander Matt Turner, Dylan Hadfield-Menell, and Prasad Tadepalli. "Conservative Agency via Attainable Utility Preservation." In: *Proceedings of the AAAI/ACM Conference on AI, Ethics, and Society*. AAAI Conference on Artificial Intelligence. AIES '20. New York, NY, USA: Association for Computing Machinery, Feb. 7, 2020, pp. 385–391. ISBN: 978-1-4503-7110-0. DOI: 10.1145/3375627.3375851.

[TRT20] Alexander Matt Turner, Neale Ratzlaff, and Prasad Tadepalli. "Avoiding Side Effects in Complex Environments." In: *Advances in Neural Information Processing Systems*. Conference on Neural Information Processing Systems. Ed. by Hugo Larochelle, Marc'Aurelio Ranzato, Raia Hadsell, Maria-Florina Balcan, and Hsuan-Tien Lin. Vol. 33. NeurIPS 2020. virtual: Curran Associates, Inc., Dec. 6–12, 2020, pp. 21406–21415.

[Tur+21] Alexander Matt Turner, Logan Riggs Smith, Rohin Shah, Andrew Critch, and Prasad Tadepalli. "Optimal Policies Tend To Seek Power." In: Thirty-Fifth Conference on Neural Information Processing Systems. May 21, 2021.

[Tur19] Alexey Turchin. "AI Alignment Problem:"Human Values" Don't Actually Exist." In: (2019).

[VGS16] Hado Van Hasselt, Arthur Guez, and David Silver. "Deep Reinforcement Learning with Double Q-Learning." In: *Proceedings of the 30th AAAI Conference on Artificial Intelligence*. AAAI Conference on Artificial Intelligence. AAAI'16. Phoenix, Arizona USA: AAAI Press, 2016, pp. 2094–2100.

[Wan+16] Ziyu Wang, Tom Schaul, Matteo Hessel, Hado Van Hasselt, Marc Lanctot, and Nando Freitas. “Dueling Network Architectures for Deep Reinforcement Learning.” In: *Proceedings of the 33rd International Conference on Machine Learning*. International Conference on Machine Learning. Ed. by Maria Florina Balcan and Kilian Q. Weinberger. Vol. 48. Proceedings of Machine Learning Research. New York, New York, USA: PMLR, June 20–22, 2016, pp. 1995–2003.

[WBC21] Nolan C. Wagener, Byron Boots, and Ching-An Cheng. “Safe Reinforcement Learning Using Advantage-Based Intervention.” In: *Proceedings of the 38th International Conference on Machine Learning*. International Conference on Machine Learning. PMLR, July 1, 2021, pp. 10630–10640.

[WD92] Christopher J. C. H. Watkins and Peter Dayan. “Q-Learning.” In: *Machine Learning* 8.3 (May 1, 1992), pp. 279–292. ISSN: 1573-0565. DOI: 10.1007/BF00992698.

[WE20] Carroll L. Wainwright and Peter Eckersley. “SafeLife 1.0: Exploring Side Effects in Complex Environments.” In: *Proceedings of the Workshop on Artificial Intelligence Safety, Co-Located with 34th AAAI Conference on Artificial Intelligence*. Artificial Intelligence Safety 2020. Ed. by Huáscar Espinoza, José Hernández-Orallo, Xin Cynthia Chen, Seán S. ÓhÉigeartaigh, Xiaowei Huang, Mauricio Castillo-Effen, Richard Mallah, and John McDermid. Vol. 2560. CEUR Workshop Proceedings. New York City, NY, USA: CEUR-WS.org, Feb. 7, 2020, pp. 117–127.

[Yud16] Eliezer Yudkowsky. “The AI Alignment Problem: Why It Is Hard, and Where to Start.” In: *Symbolic Systems Distinguished Speaker* (2016).

[ZDS18] Shun Zhang, Edmund H. Durfee, and Satinder Singh. “Minimax-Regret Querying on Side Effects for Safe Optimality in Factored Markov Decision Processes.” In: (2018), pp. 4867–4873.